

LLM Governance Conflation

The Cross-Cutting Anti-Pattern Where LLM Outputs Are Treated as Governance Decisions Rather Than Labor Outputs, Violating A1.01, A1.12, and B1.01 by Allowing the Instinct Layer to Effectively Govern the Deployment While Appearing Human-Governed

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 12, 2026

License: Creative Commons Attribution 4.0 International (CC BY 4.0)

Attribution

This work derives from and formalizes the instinct/reasoning separation pattern introduced in "The Instinct/Reasoning Separation Outside the Model" (Li, April 2026), the second paper in the CKS theory series following "Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems" (Li, April 2026).

It does not introduce new axioms. Its contribution is to identify and formalize a cross-cutting anti-pattern that emerges when the architecture's labor/governance distinction — foundational to both source papers — is violated consistently and systematically across a deployment.

Abstract

LLM Governance Conflation is a cross-cutting anti-pattern in which LLM outputs are treated as governance decisions rather than labor outputs, causing a deployment that appears human-governed (A1.01) to be effectively governed by the instinct layer (B1.01). The anti-pattern violates five architectural commitments simultaneously: A1.01 (governance is human authority), A1.12 (LLM is labor, not governance), B1.01 (instinct/reasoning separation), A2.04 (rule authoring is human), and B2.41 (birth governance is distinct from birth labor). It manifests in three recognizable forms — LLM-as-DNA-author, LLM-as-governance-advisor, and LLM-as-compliance-auditor — each of which substitutes LLM judgment for human governance judgment at a different point in the deployment lifecycle. The anti-pattern differs from Instinct-Reasoning Collapse (B3.02) in scope: B3.02 concerns cell-level operational separation; LLM Governance Conflation concerns the use of LLM outputs as governance acts across the entire deployment architecture. Detection proceeds through a governance reasoning test that asks human authorizers to explain their reasoning in their own words, independently of what any LLM said. Remediation requires governance culture change alongside technical correction.

1. Pattern name and cross-cutting character

Pattern name: LLM Governance Conflation

Series position: B3.25 (cross-cutting anti-pattern)

Commitments violated: A1.01 (human governance — authority, not labor), A1.12 (labor allocation — LLM is labor, not governance), B1.01 (instinct/reasoning separation — LLM is instinct, not the reasoning layer that governs), A2.04 (rule authoring is human), B2.41 (birth governance is distinct from birth labor)

LLM Governance Conflation is cross-cutting in the precise sense that it is not a failure of one architectural commitment localized to one layer or one lifecycle event. It is a systematic misassignment — treating LLM outputs as governance acts — that corrupts every commitment that depends on the labor/governance distinction. When the conflation is present, A1.01’s human governance is nominally preserved (a human signs every governance record) while being substantively absent (the reasoning behind every governance record is the LLM’s, not the human’s). The commitments violated are those whose meaning depends on human governance being genuine rather than ceremonial.

Distinction from B3.02 (Instinct-Reasoning Collapse). B3.02 names the cell-level anti-pattern where the instinct and reasoning layers are not architecturally separated within a cell’s operational cycle — where the LLM performs both fast-pattern response and deliberate reasoning in one undifferentiated pass. LLM Governance Conflation operates at a different level: it names the deployment-level anti-pattern where the LLM performs governance acts — acts of authority, not acts of labor — across the entire deployment. A deployment can maintain clean instinct/reasoning separation within every cell (B3.02 is absent) while allowing LLM outputs to drive governance decisions across all cells (B3.25 is present). The two anti-patterns are independent.

2. The labor/governance distinction this anti-pattern violates

The architectural foundation against which LLM Governance Conflation is measured is the labor/governance distinction that both source papers defend and that A1.12 and B1.01 make explicit.

Under A1.12, the labor allocation framework designates LLMs as one of three labor modes: LLMs perform operations under orchestration rules, alongside direct human labor and stable-cell automation. Labor and governance authority are architecturally distinct: labor is what produces or processes substrate content; governance is the human authority to author the rules under which labor operates, to approve what labor produces before it becomes authoritative, and to override any operation or output at any time. The distinction is not merely procedural — it is what makes the CKS architecture scalable (governance cost does not grow with labor volume) and human-governed in the sense A1.01 requires.

Under B1.01, the instinct layer is the LLM — fast-pattern, stateless-within-execution, capable of recognizing and generating — and the reasoning layer is the CKS substrate: deliberate, traceable, human-governed. The reasoning layer governs; the instinct layer labors under the reasoning layer’s rules. An architecture that allows the instinct layer to determine what the reasoning layer contains, or to decide what governance acts the reasoning layer records, inverts this relationship. The instinct layer is governing while appearing to serve the reasoning layer.

The operational test for conflation: Could the human who signed a governance record explain the governance reasoning behind that record in their own words, without consulting or deferring to what any LLM said? If yes, the governance act is genuine. If no — if the human’s explanation amounts to “the LLM recommended it” — the governance act is conflated. The human signature is present; the human governance judgment is absent.

3. Recognizable forms

LLM Governance Conflation manifests in three recognizable sub-forms, each locating the conflation at a different governance moment.

Form 1: LLM-as-DNA-author

In this form, LLMs draft DNA-layer rules — orchestration substrates, governing specifications, behavioral rules for cells — and those drafts are added to the substrate without substantive human governance review. The LLM effectively determines what governs cell behavior; the human governance act is rubber-stamp approval of LLM-generated specifications the human authorizer did not independently develop or deeply understand.

Recognition signals: DNA rule content reads as LLM output — natural-language governance specifications with the fluency and breadth of generated text — rather than precisely authored governance rules reflecting deliberate human design decisions. Governance records per A2.04 exist for every rule, but review durations and approval patterns suggest the review was nominal. When a cell’s behavior needs to be explained, the explanation cannot be grounded in governance intent because governance intent was the LLM’s, not the authorizing human’s. The A5.09 accountability questions — particularly “under what human reasoning was this authorized?” — cannot be answered coherently; the answer is “the LLM drafted it and a human approved it.”

A2.04 requires that rules governing cell behavior be authored by humans. LLM drafting that becomes authoritative through human approval is admissible under A1.12 as a labor allocation — LLM drafts, human governs. LLM drafting that becomes authoritative through nominal approval by a human who did not substantively govern violates A2.04. The distinction is whether the human’s approval reflects a governance judgment or merely records that the approval happened.

Form 2: LLM-as-governance-advisor

In this form, governance decisions are made by asking LLMs — “should we approve this DNA change?”, “is this cell birth specification appropriate?”, “does this proposal satisfy the architectural commitments?” — and treating LLM answers as the primary governance reasoning. The human governance act is asking the LLM rather than applying independent human judgment. The LLM answer is not used as one input among many; it is the governance reasoning, with a human signature appended.

Recognition signals: governance decision records per A2.40 contain LLM consultation records as the primary or only reasoning trace. Stage 2 approval per B2.75 is effectively the LLM’s recommendation

with a human cosignature. Human authorizers, when asked to explain a governance decision, reproduce the LLM’s reasoning without being able to say what they would have decided absent the LLM’s input. When governance teams disagree among themselves, the resolution mechanism is asking the LLM again rather than applying human judgment to resolve the disagreement.

This form is the most difficult to detect because it looks like governance. There are human authorizers; there are governance records; there is a traceable chain. What is missing is the substance: the human judgment that A1.01 requires is not present behind the signatures.

Form 3: LLM-as-compliance-auditor

In this form, LLMs are used to verify governance compliance — “does this DNA specification satisfy the Paper 1 and Paper 2 commitments?”, “does this substrate structure meet the B2.14/B2.19/B2.24 verification requirements?” — and the LLM’s compliance assessment is treated as authoritative verification. The operational tests that B2.109 specifies as human-governed activities are run by LLM as an automated pipeline, without human governance of the verification process itself.

Recognition signals: verification records for B2.14/B2.19/B2.24 milestones are produced by LLM as the primary verifier. The A5.01–A5.16 operational test suite is executed as an automated LLM-driven process without human governance of what the tests mean or what their results require. When compliance questions arise, the answer is “we ran it through the LLM compliance checker and it passed” rather than “here is what the human-governed operational tests showed.”

This form is particularly consequential because it creates a circular compliance structure: a deployment governed by LLMs is verified as compliant by the same class of system whose governance role it is conflating.

4. Emergence conditions

LLM Governance Conflation emerges under three conditions, typically in combination.

Capability trust. LLM capabilities in governance-adjacent domains — specification drafting, pattern recognition, compliance analysis, structured reasoning — are genuinely impressive. Governance teams observe LLM performance on governance tasks and conclude that LLM judgment is more reliable than their own. This is not an irrational response to the observed evidence; it is an error about what governance is. Governance is not a performance task where the best output wins; it is an authority assignment where the right agent holds authority. A more capable advisor does not become the decision-maker by virtue of capability.

Expertise gap. Governance teams frequently lack confidence in their own governance judgment, particularly for technically complex DNA specifications or cross-cutting architectural compliance questions. The LLM appears to understand the architecture better than the human authorizers do. Rather than investing in governance capability (understanding the architecture sufficiently to make genuine governance judgments), governance teams substitute LLM consultation for governance judgment. The LLM fills the expertise gap — but in doing so, it displaces governance rather than

supporting it.

Efficiency shortcut. Genuine governance review is time-consuming. Understanding a DNA specification well enough to make an independent judgment about it requires domain knowledge, architectural knowledge, and time. LLM consultation is faster and produces confident-sounding answers. When deployment timelines are compressed and governance is treated as an approval bottleneck rather than an architectural requirement, the efficiency shortcut is taken consistently, and LLM Governance Conflation becomes the default operating mode.

5. Operational consequences

AI-governed deployment appearing human-governed. The most fundamental consequence is that the deployment's actual governance has migrated to the instinct layer, while the formal record represents the deployment as human-governed. Governance integrity — the property that human authority is genuinely exercised, not merely formally recorded — is absent. This is not a compliance gap in one area; it is a system-wide inversion of the deployment's governance architecture.

Unknown governance reasoning. When LLMs govern, the reasoning behind governance decisions is the LLM's internal processing, which is not available for inspection in the form that A1.07 path retraceability requires. A1.07 requires that the reasoning behind substrate decisions be traceable — that a human inspecting the substrate can determine why a governance decision was made. When the governance reasoning is “the LLM recommended it,” that reasoning is the LLM's inference at a specific moment, which cannot be recovered or explained independently of repeating the LLM query under potentially different conditions. Retraceability is broken at the governance reasoning level.

Systemic LLM dependence. When governance depends on LLM judgment, the stability of governance depends on LLM stability. If the LLM changes — through instinct evolution (B1.10), through infrastructure upgrade, through model mutation per B1.13 — the governance judgments it would produce for the same inputs change unpredictably. A deployment that appeared to have coherent governance architecture under one LLM version may produce different governance outputs under a subsequent version, with no human-authored governance record that explains what the intended governance position is. Governance becomes fragile in proportion to its LLM dependence.

Compliance invalidity. In contexts where human governance is required by regulation, audit standard, or contractual commitment, LLM Governance Conflation produces compliance violations regardless of how complete the formal record appears. A record of LLM consultation and human signature does not satisfy a requirement for human governance judgment. This consequence is distinct from the architectural consequences above: it is external, not internal, and may have legal or contractual force.

6. Detection

Governance reasoning test. The primary detection instrument is asking human authorizers — at any point in the deployment lifecycle — to explain the reasoning behind a governance decision in their own words, without consulting any LLM. The question is not “what did the governance record say?” but

“what was your governance reasoning?” A human who made a genuine governance judgment can reproduce the core of that reasoning independently of any LLM input. A human who is conflating LLM advice with governance judgment will either reproduce LLM output or acknowledge that their reasoning was primarily LLM-derived. The test can be applied spot-check style across a sample of governance records without requiring comprehensive review.

A5.09 accountability questions. The four accountability questions — what was decided, by whom, under what authority, and under what reasoning — are applied to a sample of governance records. The “under what reasoning?” question is the diagnostic. If the answer consistently references LLM consultation as the primary or only reasoning, governance conflation is present. This test is applicable to DNA rule authoring records (Form 1), governance decision records (Form 2), and compliance verification records (Form 3).

DNA content quality analysis. DNA rule content can be analyzed for whether it reads as genuinely authored governance specification or as generated specification content. Precisely authored governance rules reflect deliberate human design decisions — they have the specificity and idiosyncrasy of human intent. LLM-authored governance rules tend toward fluent completeness, natural-language breadth, and a kind of organized comprehensiveness that reflects LLM generation rather than human governance design. This signal is secondary to the governance reasoning test, but it can flag candidates for closer review.

Verification record analysis. Compliance verification records (Form 3) can be analyzed for whether they reflect human-governed operational test execution or automated LLM pipeline output. Human-governed verification produces records that reference specific operational observations, human reviewers, and test conditions. LLM-driven verification produces records that reference LLM assessment of specification text. The distinction is visible in the record structure.

7. Remediation

Remediation for LLM Governance Conflation requires two parallel tracks: technical correction and governance culture change. Technical correction alone does not remediate the anti-pattern; governance culture change is required because the conflation is a behavior pattern, not a system configuration.

Technical correction. Governance records for DNA rule authoring, governance decisions, and compliance verification are audited against the governance reasoning test. Records that cannot be supported by human governance reasoning are flagged as conflated. Affected rules, decisions, and verifications are re-executed under genuine human governance — not by re-running the same process with more thorough LLM consultation, but by having human authorizers develop and record their own governance positions. Verification substrates and operational tests (B2.109, A5.01–A5.16) are restructured as human-governed activities with human reviewers examining test results and drawing conclusions, rather than as LLM-automated pipelines.

Governance culture change. The governance teams operating the deployment must relearn what genuine human governance means, and specifically what it requires of them. The core relearning is that

governance is an authority assignment — the human’s obligation is to hold and exercise genuine judgment, not to aggregate inputs and surface the best answer. LLMs may continue to perform labor: drafting DNA rule specifications for human governance review, analyzing proposed changes for pattern recognition, surfacing relevant prior governance decisions for human consideration. The distinction that must become operational culture is: LLM drafting plus human governance judgment equals appropriate labor allocation under A1.12; LLM judgment plus human signature equals conflation that this note formalizes. The governance teams must be able to explain, in their own words and independent of LLM input, why they authorized what they authorized.

Ongoing monitoring. The governance reasoning test should be applied on a regular schedule to a random sample of governance records, not only during remediation. This provides an ongoing signal about whether governance culture change has been sustained. A deployment that passes spot-check governance reasoning tests consistently over time has addressed the anti-pattern. A deployment that passes initial remediation but fails subsequent spot checks has not.

8. Conclusion

LLM Governance Conflation names a systematic failure of the labor/governance distinction that both source papers defend as foundational. The failure is not located in one cell, one lifecycle event, or one architectural layer — it is present wherever LLM outputs have been substituted for human governance judgment across the deployment. The architectural consequence is a deployment that is formally human-governed and substantively LLM-governed, with the human governance record serving as a façade over an instinct-layer governance structure. Detection is possible through the governance reasoning test, which asks whether the humans who signed governance records can reproduce the governance reasoning without LLM support. Remediation requires governance culture change: human authorizers must develop and hold genuine governance judgment, with LLMs allocated to labor roles under the governance framework the humans own.

The labor/governance distinction is not an administrative requirement. It is the property that makes the CKS architecture coherent — the property that allows governance cost to remain bounded, allows retraceability to hold, and allows human authority to mean something more than the right to sign off on what the instinct layer has decided.

Source papers

Li, W. (2026). *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems*. April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *The Instinct/Reasoning Separation Outside the Model: Extending the Coordination Knowledge Substrate Pattern to AI Selves under Human Governance*. April 2026. ORCID: 0009-0004-8065-3235.

How to cite this note

Li, W. (2026). *LLM Governance Conflation — The Cross-Cutting Anti-Pattern Where LLM Outputs Are Treated as Governance Decisions Rather Than Labor Outputs, Violating A1.01, A1.12, and B1.01 by Allowing the Instinct Layer to Effectively Govern the Deployment While Appearing Human-Governed*. May 12, 2026. ORCID: 0009-0004-8065-3235.