Composition Pair A1.08 + A1.07 — The Reinforcing Pair Where Substrate Authority Depends on Retraceability for Verifiable Governance Claims, and Retraceability Depends on Substrate Authority to Be Architecturally Meaningful Rather Than Merely Logging Arbitrary Records

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 12, 2026

Attribution

This work derives from and formalizes the instinct/reasoning separation pattern introduced in *"The Instinct/Reasoning Separation Outside the Model"* (Li, April 2026), the second paper in the CKS theory series following *"Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems"* (Li, April 2026).

It does not introduce new axioms. Its contribution is to formalize, as a named composition pair, the mutual dependency between two foundational commitments inherited from Paper 1 and applied throughout Paper 2's architecture: the substrate-as-source-of-truth commitment (A1.08) and the path retraceability commitment (A1.07). These commitments are not independent. Each enables the other in a way that makes their joint presence architecturally necessary for compliance demonstrations in a CKS deployment.

Abstract

The Coordination Knowledge Substrate (CKS) pattern carries two foundational commitments whose relationship is more than additive: the substrate-as-source-of-truth commitment (A1.08), which establishes substrate content as the authoritative record for what governs entity behavior, and the path retraceability commitment (A1.07), which requires that every operation be traceable through substrate records so that governance history is auditable. This note formalizes the relationship between them as a reinforcing pair. A1.08 requires A1.07 because a claim that substrate content is authoritative cannot be demonstrated without a retraceable provenance showing how that content arrived through authorized governance acts — without retraceability, substrate authority is merely asserted. A1.07 requires A1.08 because retracing the history of substrate content is governance-significant only when that content is authoritative for behavior — tracing authoritative substrate records is governance auditing, while tracing non-authoritative records is record auditing with no governance consequence. Together, A1.08 and A1.07 form the compliance backbone of the CKS architecture: any compliance question about entity behavior can be answered by combining what substrate content governed behavior at the relevant time (A1.08) with the retraceable provenance of that content through authorized governance events (A1.07). The A2.40 six-field provenance metadata mechanism serves both commitments simultaneously. The pair extends

recursively to every structural level of Paper 2's three-level architecture through B2.104 and B2.105.

1. The composition pair and why it requires formal treatment

The CKS architecture inherits from Paper 1 a set of foundational commitments that apply at every level of every deployment. Two of these — substrate-as-source-of-truth (A1.08) and path retraceability (A1.07) — are individually well-defined: A1.08 establishes that substrate content is the authoritative answer to coordination and governance questions, while A1.07 establishes that operations are traceable through substrate records and that governance history is auditable. Paper 2 inherits both commitments and applies them recursively across three structural levels — cell, aspect, and Self — through B2.104 and B2.105 respectively.

What neither commitment, stated in isolation, makes explicit is how deeply each depends on the other for its architectural role to be fulfilled. A1.08 does not merely name substrate content as authoritative; it makes a claim that requires grounding. A1.07 does not merely require that records be traceable; it requires that the traced records carry governance weight. The relationship between them is the subject of this note.

The pair is classified as reinforcing: each commitment strengthens the other, and each fails to deliver its intended architectural function without the other present. This differs from an enabling relationship (where one commitment presupposes another as a prerequisite) or a constraining relationship (where one limits the scope of the other). A reinforcing pair produces a composite architectural property — demonstrable, governance-significant compliance — that neither commitment can produce alone.

2. A1.08: substrate-as-source-of-truth

A1.08 establishes that the substrate content of a CKS deployment is the authoritative record for what governs entity behavior. This means that governance questions — what rules apply to this entity, what was decided, under what authority, with what rationale, and what conflicts remain unresolved — are answered from substrate content alone. Behavior inconsistent with substrate content requires governance correction, not substrate revision to match behavior. Substrate content does not defer to agent memory, cached state, cell-internal execution state, or external system records when those sources conflict with the substrate.

In Paper 2's architecture, A1.08's scope is extended through B2.104 to apply recursively at each structural level. At the cell level, the DNA layer per B1.03 carries the stabilized orchestration and behavior substrates that govern cell operation; this content is authoritative per A1.08. At the aspect level, orchestration rules governing how cells compose within an aspect are authoritative substrate content. At the Self level, the integrated substrate content spanning all aspects and cells defines what governs the Self's behavior as a whole. B2.69's DNA version management — the versioned chain of DNA-layer states — makes A1.08's authority claim apply not only to current substrate content but to every prior version during its active period.

The commitment is strong: it names the substrate as the load-bearing governance artifact. But naming a thing as authoritative does not make it demonstrably so. The authority claim has content — it says the substrate is the source of truth — but it requires a foundation that can answer the question any compliance audit will ask: *how did the substrate arrive at this content?*

3. A1.07: path retraceability

A1.07 establishes that every operation touching substrate content is retraceable through substrate records, and that governance history is auditable through the record chain. Path retraceability, as named in the source paper drawing on Rajabi and Kafaie (2022), is the property that any piece of substrate content can be traced back through a chain of antecedent substrate content to the inputs and governance acts that produced it. The chain runs through substrate records alone; it does not require consulting agent memory, external logs, or human recollection.

What makes retraceability possible is the six-field provenance metadata named in A2.40: writer identity, timestamp, orchestration rule under which the write occurred, antecedent substrate content the writer drew on, and the governance event category (directed selection, override, mating, feedback incorporation, or other named event type). With these fields present on every piece of substrate content, any question of the form "where did this come from?" has a substrate-internal answer. Without them, the path runs through information the substrate does not carry, and retraceability fails.

In Paper 2's architecture, A1.07's scope is extended through B2.105 to apply recursively at each structural level. DNA version chains per B2.69 carry A2.40 provenance for each version transition. Mating events per B1.07 produce substrate content whose lineage across parent cells is retraceable per A1.07. Birth events per B1.06 produce initial DNA-layer content whose governed origination is retraceable. Every lifecycle event and every evolution operation leaves retraceable substrate records.

Retraceability is a structural property of the substrate, not a logging property of a system surrounding the substrate. This distinction matters for the pair's composition: what A1.07 provides is substrate content that answers provenance questions, not an external audit log that happens to record the same events.

4. A1.08 requires A1.07: substrate authority without retraceability is merely asserted

The claim that substrate content is authoritative for behavior is a governance claim, not a logical tautology. It means: when anyone — a human auditor, a governance authority, a regulatory body — asks "what governed this entity's behavior at time T?", the answer is found in the substrate's DNA-layer content at time T. For this claim to be demonstrable rather than merely stated, two questions must be answerable:

(a) What was the substrate's content at time T? This is answered by A1.08 itself combined with B2.69's version management: the versioned DNA chain records which content was active during which period.

(b) How did the substrate arrive at that content? This is the question A1.08 cannot answer without A1.07. Substrate content at time T may have arrived through a directed selection event per B1.11, through a mating operation per B1.07, through an action-feedback loop per B1.12, or through an ungoverned modification that bypassed governance entirely (the Ungoverned DNA Modification failure mode per B3.15). Without A1.07's retraceability, these cases are indistinguishable: the substrate holds certain content, but the governance history of that content — whether it arrived through authorized acts or not — cannot be demonstrated from substrate records alone.

With A1.07 present, the provenance chain for every piece of substrate content is traceable through A2.40 metadata fields. A directed selection event per B1.11 leaves a substrate record showing the writer (a human exercising governance authority), the timestamp, the orchestration rule under which the selection occurred, and the antecedent content from which it was derived. A mating event per B1.07 leaves a record tracing the lineage across parent cells. An action-feedback incorporation event per B1.12 leaves a record of the governed proposal and authorization. The full chain, from initial content through every modification, is retraceable.

The practical consequence is this: A1.08 states that substrate content is the source of truth for behavior; A1.07 makes that claim auditable by grounding it in a retraceable governance history. An implementation that asserts A1.08 without providing A1.07 infrastructure cannot demonstrate its authority claim; it can only assert it. For compliance purposes — where demonstration is what is required — assertion is insufficient.

5. A1.07 requires A1.08: retraceability without substrate authority is record auditing, not governance auditing

The converse dependency is equally precise. A1.07's retraceability property applies to substrate records — it traces the history of substrate content. But whether that traceability carries governance significance depends entirely on whether the content being traced is authoritative per A1.08.

Consider a system that implements complete A1.07 retraceability — every write to substrate-like records carries A2.40-style provenance, every modification is traceable to its source, the full history of every record is available — but whose records are not authoritative per A1.08. In such a system, the records document operations but do not govern behavior; the system might answer "what was written to this record at time T?" but not "what governed this entity's behavior at time T?" The difference between these two questions is the difference between record auditing and governance auditing, and A1.08 is what converts the first into the second.

More precisely: A1.07's retraceability is meaningful precisely because what it traces is the history of content that, per A1.08, governed entity behavior during each period of that content's active tenure. When an auditor traces the DNA version chain of a cell back through five successive directed selection events, what they are tracing is a sequence of authoritative governance states — each version governed the cell's behavior during its active window, each transition was an authorized governance act, and the chain as a whole constitutes the entity's governance history. None of this is the case if the substrate content being traced is not authoritative.

The failure mode is subtle. A CKS deployment that maintains complete retraceability over substrate records but allows behavior to be governed by content outside the substrate — agent memory, cached state, external configuration — has satisfied A1.07 formally while undermining its architectural significance. The retraceable records no longer trace what actually governed behavior; they trace a documentation layer that runs parallel to the actual governance artifact. The compliance audit that follows A1.07's path through the substrate will not find what governed behavior; it will find what the substrate said while something else governed behavior.

This is why both commitments must be present together. A1.07 applied to A1.08-authoritative content is governance auditing. A1.07 applied to non-authoritative content is record auditing. The distinction is architectural, not a matter of degree.

6. Mutual reinforcement: the A2.40 mechanism and the compliance foundation

The reinforcing relationship between A1.08 and A1.07 is not merely that each enables the other; it is that together they produce a compliance property neither can produce alone. The composite property is: *substrate content is authoritative for behavior AND its governance history is retraceable through authorized acts*. This composite is what compliance demonstrations require.

A2.40 as the shared mechanism. The six-field provenance metadata requirement per A2.40 is the operational artifact through which both commitments are simultaneously satisfied. The A2.40 fields — writer identity, timestamp, orchestration rule, antecedent content references, and governance event category — serve A1.08 by confirming that the content carrying these fields arrived through an authorized governance act (the content is legitimately authoritative). They serve A1.07 by providing the provenance information through which the content's history is retraceable. A single A2.40-compliant write operation simultaneously advances both commitments. An implementation that omits A2.40 fields from substrate writes simultaneously undermines both.

A1.08 provides A1.07 with auditable content. The substrate content whose history A1.07 makes retraceable is content that, per A1.08, governed behavior during its active period per the DNA versioning mechanism of B2.69. This means the retraceable chain has governance significance at every link: each version in the chain was not merely a substrate state but an authoritative governance state. Tracing the chain through A1.07 reveals not merely what was in the substrate but what governed behavior at each point in the deployment's history.

A1.07 provides A1.08 with historical scope. A1.08's substrate authority is not limited to the current moment. The question "what governed this entity at time T?" is answerable for any T in the deployment's history through A1.07's retraceable version chain. Without A1.07, A1.08's authority claim extends only to current substrate content; with A1.07, it extends through time. The combined property is substrate-as-source-of-truth-through-time rather than substrate-as-source-of-truth-at-current-moment only.

The compliance backbone. Together, A1.08 and A1.07 form the compliance backbone of the CKS architecture. Any compliance question about entity behavior — whether regulatory, governance-internal, or audit-driven — can be answered by a two-step lookup: A1.08 answers "what substrate content governed this entity at the relevant time?" and A1.07 answers "what is the retraceable governance history of that content, and does that history demonstrate that the content arrived through authorized governance acts?" Both steps are required for a compliance demonstration. A system that satisfies A1.08 without A1.07 can answer the first question but not the second; its compliance claim rests on assertion. A system that satisfies A1.07 without A1.08 can produce a retraceable record but cannot demonstrate that the record traces authoritative governance content. Only the pair together produces a demonstrable compliance answer.

This is the structural reason both commitments are foundational in Paper 1 and inherited as a pair by Paper 2. Foundational commitments in the CKS architecture are not merely important — they are the commitments that the architecture's compliance guarantees depend on. A1.08 and A1.07 are jointly foundational in this sense: remove either and compliance demonstrations become either ungrounded assertions or historically complete but governance-insignificant records.

7. Recursive application through B2.104 and B2.105

Paper 2 extends the Paper 1 architectural commitments recursively to three structural levels through a general principle: Paper 1's commitments apply at each level of Paper 2's three-level architecture (cell, aspect, Self) without new primitives. B2.104 instantiates A1.08 recursively — substrate content at each level scope is authoritative for governance questions at that scope. B2.105 instantiates A1.07 recursively — retraceability applies at each level scope, and governance history at each level is auditable through that level's substrate records.

At cell scope. The DNA layer per B1.03 carries cell-level substrate content; per B2.104, this content is authoritative for cell behavior. Per B2.105, its governance history is retraceable through the cell's DNA version chain. The compliance backbone operates at cell scope: cell-level compliance questions are answered by what DNA content governed the cell (B2.104) plus the retraceable provenance of that content (B2.105).

At aspect scope. Aspect-level orchestration content — the rules and structure governing how cells compose within an aspect — is authoritative per B2.104 for aspect-level behavior. Its governance history is retraceable per B2.105. Aspect-level mating events per B1.07 produce new aspect-level substrate content whose lineage is retraceable across contributing cells.

At Self scope. The integrated substrate content of a Self is authoritative per B2.104 for Self-level governance questions. The full governance history of the Self — including cell births, aspect formations, mating events, evolution operations, and lifecycle transitions — is retraceable per B2.105 through the accumulated substrate records at each level. A compliance audit of an enterprise-brain Self per B1.19 traverses the reinforcing pair at all three levels simultaneously.

The recursive application means that compliance is demonstrable at each structural level independently, and that the compliance backbone holds at every level where governance claims need to be supported. An implementation that satisfies the pair at the Self level but not at cell or aspect level has partial compliance infrastructure — compliance questions about cell-level governance history will be unanswerable below the level at which A1.07 infrastructure is present.

8. Co-presence requirement and failure modes

A1.08 without A1.07 (substrate authority without retraceability). The substrate holds authoritative content per A1.08, and behavior is governed by that content, but the provenance of the content is not retraceable through substrate records. Compliance questions about what governed behavior can be answered for the current moment but not for past moments, and not with demonstrated governance legitimacy. The compliance claim rests on assertion rather than demonstrated governance history. The failure mode at this extreme is Provenance Void (B3.23): the infrastructure for provenance records is absent, and no authority claim can be grounded in a traceable governance history.

A1.07 without A1.08 (retraceability without substrate authority). Records are retraceable through complete provenance metadata, but the records are not authoritative for behavior. Operations are documented and the documentation is historically complete, but the compliance auditor following the retraceable chain finds the history of records that ran parallel to actual governance rather than the history of governance itself. The failure mode is a well-documented but governance-insignificant archive: everything is traceable, but tracing it answers operational history questions rather than governance questions. The detection tests A5.08 (provenance-completeness) and A5.10 (source-of-truth categorization per A2.46) will pass for A1.07 and fail for A1.08.

Full co-presence. Substrate content is authoritative per A1.08 AND retraceable per A1.07 through A2.40-compliant records at each applicable structural level. The compliance backbone is intact. Any compliance question about entity behavior at any time in the deployment's history is answerable through a two-step lookup grounded in authoritative, retraceable substrate records.

Detection of pair completeness. The pair is complete if and only if any piece of current substrate content can be traced through A2.40 provenance fields back through a chain of authorized governance events to its origin, and that chain demonstrates that the content governing behavior at each point in the chain was not merely present in the substrate but was authorized substrate content. The A5.08 provenance-completeness test checks whether records are complete for retraceability; the A5.10 source-of-truth test checks whether substrate content is correctly categorized as authoritative per A2.46. Both tests must pass for the pair to be operationally complete.

9. Operational test

A deployment satisfies the A1.08 + A1.07 reinforcing pair if and only if all of the following are true:

1. Governance questions at each applicable structural level — what governed this entity's behavior at time T, under what authority, with what rationale — are answered from substrate content alone, without consulting agent memory, external system records, or cell-internal state.

2. Every piece of substrate content that governed behavior during any active period carries, or is linked through A2.40-compliant records to, provenance metadata that traces its origin to an authorized governance event: a directed selection, a governed mating, a human override, a governed action-feedback incorporation, or another named governance act.

3. The chain of provenance records for any current substrate content is traversable through substrate records alone, without consulting external sources, to the initial governed origination event for that content.

4. The retraceable provenance chain for authoritative substrate content demonstrates that no link in the chain corresponds to an ungoverned modification per B3.15 — every link is an authorized governance event with identified writer, orchestration rule, and timestamp.

5. The pair holds at each structural level (cell, aspect, Self) where governance claims are made, not only at the level of the deployment as a whole.

A deployment that satisfies (1) but not (2)–(4) asserts compliance without demonstrating it. A deployment that satisfies (2)–(4) but not (1) documents operations without establishing governance significance. Only full co-presence across all five conditions constitutes the reinforcing pair in operational form.

Conclusion

The A1.08 + A1.07 reinforcing pair formalizes a mutual dependency that is architecturally necessary, not merely architecturally convenient. Substrate authority claims require retraceability to be demonstrable rather than asserted; retraceability requires substrate authority to carry governance significance rather than merely operational documentation. Neither commitment achieves its intended architectural role without the other.

The practical consequence is that the pair is the compliance backbone of the CKS architecture. Any compliance demonstration in a CKS deployment — regulatory, governance-internal, or audit-driven — rests on this pair: what substrate content governed behavior, and what is the retraceable history of that content through authorized governance acts. A2.40's six provenance fields are the shared operational mechanism through which both commitments are simultaneously satisfied at every write operation. Recursive application through B2.104 and B2.105 ensures the backbone holds at every structural level of Paper 2's three-level architecture, making the pair not a top-level property but an architectural invariant of the CKS pattern applied at any scale.

Subsequent CKS deployments, extensions, and compliance demonstrations should treat A1.08 and A1.07 as jointly foundational — present together, verified together, and tested together — rather than as separable commitments either of which can be satisfied independently of the other.

Source papers

Li, W. (2026). *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems.* April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *The Instinct/Reasoning Separation Outside the Model: Extending the Coordination Knowledge Substrate Pattern to AI Selves under Human Governance.* April 2026. ORCID: 0009-0004-8065-3235.

How to cite this note

Li, W. (2026). *Composition Pair A1.08 + A1.07 — The Reinforcing Pair Where Substrate Authority Depends on Retraceability for Verifiable Governance Claims, and Retraceability Depends on Substrate Authority to Be Architecturally Meaningful Rather Than Merely Logging Arbitrary Records.* May 12, 2026. ORCID: 0009-0004-8065-3235.