

FAI Governance for Competitive Intelligence

Series: CKS Derivation Notes — Phase D2, Note D2.65 (#560)

Derives from: D2.20 (cooperation/competition variant), D1.13 (preserve tier)

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 15, 2026

License: Creative Commons Attribution 4.0 International (CC BY 4.0)

Attribution

This work derives from and formalizes the inter-Self coordination architecture introduced in “Inter-Self Coordination via Shared Substrate / Full Aspect Integration” (Li, April 2026), the third paper in the CKS theory series following “Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems” (Li, April 2026) and “The Instinct/Reasoning Separation Outside the Model” (Li, April 2026).

Abstract

D2.65 operationalizes the competition variant (D2.20) for the specific case where the FAI event’s purpose is generating comparative governance intelligence rather than achieving a merged resolution. In the competitive intelligence configuration, the conflict registry is the primary output: a dense registry showing where two Selves’ governance approaches differ is the intended result, not a residue of incomplete resolution. This requires a preserve-tier-dominant conflict routing configuration, symmetric rich sharing scope, and four specific governance requirements: conflict classification configuration, full carry-through for all preserved conflicts, prompt post-mortem review, and explicit governance review of any subsequent DNA absorption. The note distinguishes competitive intelligence FAI from knowledge transfer FAI (D2.64), states the misrouting anti-pattern that defeats the event’s purpose, and closes with an operational test for verifying that a competitive intelligence event is correctly configured.

1. D2.65 as operational decomposition of D2.20

D2.20 established that the competition variant of Full Aspect Integration (FAI) uses the same shared-substrate architecture as the cooperation variant under different orchestration rules. Both variants are the same architectural object — a temporary shared substrate spanning two or more Selves’ governance perimeters — operating under governance-configured rule sets that shape what the event produces.

D2.20’s competition variant is an orchestration configuration, not a different mechanism. The shared substrate is constructed in the same way, aspects are contributed at governance-configured sharing

scope, conflicts surface as first-class substrate state during the merge, and the three-tier conflict-handling mechanism (D1.13–D1.16) applies throughout. What differs is the orchestration configuration that routes surfaced conflicts: cooperation-variant events optimize for mutual benefit through orchestration that typically routes conflicts toward resolution; competition-variant events route most conflicts toward preservation, making the conflict registry the governance output rather than a residue to be cleared.

D2.65 decomposes the competition variant into the specific orchestration configuration appropriate for competitive intelligence: the case where the event’s governing purpose is producing comparative intelligence about how two Selves’ governance approaches differ. Not every competition-variant event serves this purpose; a competition-variant event could also serve mutual benchmarking without the full intelligence-extraction orientation D2.65 specifies. The competitive intelligence configuration is the highest-fidelity expression of the competition variant’s preserve-tier routing logic applied to a deliberate comparative governance purpose.

2. The competitive intelligence configuration

2.1 Symmetric rich sharing scope

Competitive intelligence FAI requires symmetric rich sharing scope: both Selves contribute substantive aspects in the same domain. This distinguishes the competitive intelligence configuration from knowledge transfer events (D2.64), where one Self contributes and one receives, and from shallow sharing configurations where neither Self exposes enough content to enable meaningful comparison.

Symmetric rich sharing is not an incidental property of the configuration; it is a prerequisite for the conflict registry to carry governance intelligence. If only one Self contributes substantive content, there is nothing to compare. If both contribute but in different domains, the conflicts that surface reflect domain non-overlap rather than governance approach differences. The comparative intelligence value depends specifically on both Selves contributing substantively in the same domain, so that the conflicts surfaced during the merge reflect genuine differences in how each approaches the same governance problem.

The sharing scope configuration specifies which aspects each Self contributes to the shared substrate. For competitive intelligence events, governance should configure sharing scope so that the aspects contributed by both Selves cover the same domain with sufficient depth to generate approach-difference conflicts. Shallow scope will produce a sparse conflict registry even when governance approaches genuinely differ — not because the Selves have converged but because the event did not reveal the divergence. This is the contribution depth gap, and it is a governance configuration failure.

2.2 Preserve-tier-dominant conflict routing

The orchestration rules for a competitive intelligence FAI event are configured to route most conflicts to the preserve tier (D1.13) rather than the resolve tier. This is the defining feature of the competitive intelligence configuration.

The preserve tier inherits from Paper 1's substrate-level conflict preservation: conflicts routed to the preserve tier remain first-class addressable substrate state within the shared substrate and, upon dissolution, carry through to each participating Self's home substrate as annotated content. They are not cleared, merged, or marked as errors. They persist as the record of where two governance approaches differ — which is exactly the record the competitive intelligence event is designed to produce.

The conflict routing configuration is orchestration substrate content authored under joint authority of participating Selves' governance. Governance specifies, in the orchestration rules, which conflict types route where. A well-configured competitive intelligence event specifies: conflicts of type A (approach differences) → preserve; conflicts of type B (factual inconsistencies) → resolve; conflicts of type C (domain boundary conflicts) → escalate or preserve as configured. This per-type routing is what Section 3 develops in detail.

2.3 The conflict registry as primary output

In a standard FAI event, the conflict registry — the collection of preserved conflicts as substrate state — is a record of what was not resolved: aspects of the shared substrate that need governance attention. In a competitive intelligence FAI event, the conflict registry is the primary output. A dense conflict registry showing where the two Selves' governance approaches differ is the intended result of a successful event.

This reframing has direct implications for how governance evaluates the event. For a standard event, conflict density is something to manage: a sparse registry after dissolution signals successful conflict handling. For a competitive intelligence event, the evaluation is different: a dense registry signals that both Selves contributed substantively, that genuine approach differences exist, and that the event surfaced them successfully. A sparse registry signals either governance convergence (both Selves' approaches have aligned, which is valuable intelligence in itself) or insufficient contribution depth (the sharing scope was too shallow to reveal the divergence that likely exists, which is a configuration gap to address before the next event).

Neither outcome is automatically bad; both are informative. What a governance observer must diagnose is which condition produced the sparse registry — convergence or shallow scope. That diagnosis depends on comparing the actual sharing scope against what the domain warranted. If the sharing scope was configured richly and the registry is still sparse, convergence is the plausible reading. If the sharing scope was shallow, the sparse registry is an artifact of undercontribution and the intelligence value of the event is diminished.

3. Four governance requirements specific to competitive intelligence FAI

Requirement 1 — Conflict classification configuration

Preserve-tier-dominant routing without conflict classification produces a uniform registry that conflates approach differences with factual inconsistencies and domain boundary conflicts. This conflation reduces the governance intelligence value of the registry: an observer examining the registry cannot distinguish which preserved conflicts reflect genuine governance philosophy differences (high intelligence value) from which reflect one Self having incorrect data (low intelligence value, more productively resolved than preserved).

The orchestration rules for a competitive intelligence event must therefore be configured to classify conflicts by type before routing them. The classification governs the routing: approach-difference conflicts route to the preserve tier because they are the primary intelligence payload; factual inconsistency conflicts route to the resolve tier because preserving them adds registry noise without comparative intelligence; domain boundary conflicts route according to governance judgment about whether the boundary itself is a finding worth preserving.

Conflict classification is an orchestration rule configuration, and the classification definitions themselves are substrate content under joint authority of participating Selves' governance. Joint authority is necessary because classification requires agreement on what constitutes an approach difference versus a factual inconsistency — a determination that one Self's governance cannot make unilaterally without distorting the comparison. The classification configuration is authored before the event and applies during the merge. If the pre-event classification configuration is incomplete, conflicts may be misrouted during the event and the registry may not accurately reflect the approach differences that exist.

Requirement 2 — Full carry-through for all preserved conflicts

Preserved conflicts in the shared substrate must carry through (D1.16) to both Selves' home action layers upon dissolution. Full carry-through means that every preserved conflict — every comparative intelligence signal the event produced — is propagated as annotated content into each participating Self's home substrate, where it becomes available for home governance to act on.

Without carry-through commitment, preserved conflicts dissolve with the shared substrate and produce no downstream governance effect. The competitive intelligence event generates a dense conflict registry during the event, but if that registry does not travel home, the intelligence it represents is not available to either Self's governance for analysis, decision, or evolution. The event's value is contained within its own lifespan rather than feeding the governance processes that justify running it.

Full carry-through for competitive intelligence events means that the carry-through depth configuration is set to propagate all preserved conflicts — not just a subset. Selective carry-through is appropriate for some event types where only certain preserved conflicts are relevant to home governance. For competitive intelligence, where the registry's comparative intelligence is the point, selective carry-through would filter out exactly what the event was designed to produce. The configuration should specify complete registry carry-through as the governance commitment for competitive

intelligence events.

Requirement 3 — Prompt post-mortem review

Competitive intelligence events benefit specifically from post-mortem review (D2.39) conducted promptly after dissolution while the comparative analysis is fresh. The conflict registry produced by a competitive intelligence event requires governance interpretation: the registry records that conflicts exist and classifies them by type, but governance judgment is needed to interpret what the approach-difference conflicts reveal about each Self's governance philosophy and what — if anything — each Self should consider doing differently.

This interpretation is context-sensitive and time-sensitive. The governance participants who ran the event have direct knowledge of the contributing aspects and the organizational context that produced each Self's governance approach; that context is most accessible immediately after the event. A post-mortem review deferred weeks or months after dissolution is working with colder context and risks governance amnesia about the substantive content that made specific conflicts informative.

The post-mortem for a competitive intelligence event specifically analyzes the conflict registry patterns: which domain showed the most approach-difference conflicts; whether the distribution of conflict types matched the classification configuration's expectations; whether the contribution depth was sufficient; whether any approach differences revealed by the registry represent genuine improvements one Self should consider absorbing. This analysis is the governance step that converts a conflict registry into actionable governance intelligence.

Requirement 4 — Absorption governance

When a competitive intelligence FAI event reveals that another Self's governance approach represents a genuine improvement over one's own, the natural governance response is to consider absorbing elements of that approach through the DNA evolution mechanism (D2.11, Paper 3 §7). DNA absorption from competitive intelligence events must be explicitly governance-reviewed rather than treated as a routine ingestion operation.

The competitive intelligence event surfaces differences; it does not adjudicate which differences represent improvements. That adjudication is a governance judgment. Two Selves can have genuinely different governance approaches and both be correct within their respective contexts: a larger Self may use a more elaborate orchestration configuration precisely because its scale warrants it; a smaller Self may use a simpler configuration that is optimal for its current scope. Absorbing the larger Self's orchestration complexity would be a governance error for the smaller Self even though the difference showed up as a conflict in the registry.

Absorption governance for competitive intelligence events therefore involves an explicit review step: governance examines the preserved approach-difference conflicts from the registry, interprets what each represents in context, and makes a directed selection judgment about whether any element warrants absorption. This review does not occur automatically upon dissolution and carry-through; it

requires governance initiative. The carry-through annotates home substrate with the preserved conflicts; governance must act on those annotations. Configuring a competitive intelligence event without a downstream absorption review step leaves the event's most consequential governance decision — what, if anything, to learn from the comparison — ungoverned.

4. Competitive intelligence as governance benchmarking through governed exchange

Competitive intelligence FAI is governance benchmarking through governed substrate exchange. An organization running a competitive intelligence FAI event against an organizational peer is comparing its governance architecture against the peer's using the shared substrate as the comparison medium — not by adopting externally imposed standards or by hiring consultants to evaluate both organizations against a common rubric, but by running a governed event in which both Selves contribute aspects under joint authority, conflicts surface as first-class substrate state, and the resulting registry is a governance-internal record of where their approaches differ.

This distinguishes competitive intelligence FAI from external benchmarking in two ways. First, the comparison medium is the organizations' own substrate content rather than an external evaluator's framework. The conflicts that appear in the registry are conflicts between what each Self has actually built into its governance substrate — not responses to a survey or scores on a rubric designed by a third party. Second, the comparison is symmetric and mutual: both Selves contribute and both receive the registry's intelligence, under joint authority that both governance structures have authorized. Neither Self is being evaluated by the other; both are being compared against each other in a governed substrate space that both jointly administer.

The benchmarking value of the event is bounded by the quality of the conflict classification configuration and the interpretation depth of the post-mortem review. Classification determines whether the registry accurately captures approach differences rather than noise; post-mortem review determines whether governance draws useful conclusions from what the registry reveals. A competitive intelligence event with poor classification or no post-mortem is a governed event that produces an ungoverned output — the registry exists but produces no governance intelligence downstream.

The broader population-level implication (Paper 3 §9) is that repeated competitive intelligence FAI events across a network of Selves, each producing conflict registries and post-mortems, constitute a governance intelligence infrastructure at population scope. No individual event reveals population-level patterns; accumulated events, with their registries and post-mortems, constitute a distributed record of how governance approaches vary, converge, and evolve across the network. This is beyond the scope of D2.65, which addresses the single-event configuration, but it is the motivation for treating competitive intelligence FAI as a governed practice rather than an ad hoc activity.

5. Anti-pattern: competitive intelligence misrouting

Anti-pattern: Configuring a competitive intelligence event with resolve-tier-dominant conflict routing rather than preserve-tier-dominant routing.

Form: The orchestration rules are configured — either by default inheritance from a cooperation-variant template or by explicit misconfiguration — to route most conflicts to the resolve tier. The shared substrate executes the merge, conflicts surface, and the orchestration rules resolve them as they arise. The event concludes with a largely resolved shared substrate and a sparse conflict registry.

Why this defeats the event’s purpose: The competitive intelligence event was configured to produce a dense conflict registry as its primary output. Resolve-tier-dominant routing collapses the approach-difference conflicts into resolved outputs, eliminating the comparative intelligence those conflicts would have carried. The governance observation that was the event’s purpose — where do our governance approaches genuinely differ? — is answered with silence, not because the approaches are similar but because the configuration eliminated the evidence of their differences before it could be recorded.

Diagnostic: After a competitive intelligence event, check the conflict registry density. If the registry is sparse and the sharing scope was rich, check the orchestration rules for the conflict routing configuration. Resolve-tier-dominant routing in a competitive intelligence event is the most likely explanation for a sparse registry from rich contributions.

Correction: Reconfigure the orchestration rules before the next event to specify preserve-tier-dominant routing for approach-difference conflict classes. Do not inherit cooperation-variant orchestration templates for competition-variant events without reviewing the conflict routing configuration.

6. Operational test

For a competitive intelligence FAI event, an observer performing governance verification asks three questions:

Question 1: Can the observer find preserve-tier-dominant conflict routing in the orchestration rule configuration? The configuration should specify that approach-difference conflicts route to the preserve tier. The routing rules should be inspectable as substrate content under joint authority of both Selves’ governance.

Question 2: Can the observer find a dense conflict registry as the primary output of the event, with conflicts classified by type? The registry should contain approach-difference conflicts as a substantial share of its entries. The classification annotations should allow an observer to distinguish approach-difference conflicts from factual inconsistency conflicts and domain boundary conflicts.

Question 3: Can the observer find carry-through annotations in both Selves’ home action layers? The preserved conflicts should have propagated to each participating Self’s home substrate upon

dissolution, where they appear as annotated content available for home governance to act on. If the registry is dense but the carry-through annotations are absent from one or both home substrates, the event's intelligence value did not reach the governance processes it was designed to serve.

An event that passes all three checks is correctly configured as a competitive intelligence event. An event that fails any check has a configuration gap that should be identified and corrected before the next event of the same type.

Summary

D2.65 operationalizes the competition variant (D2.20) for the competitive intelligence use case. The competitive intelligence configuration is defined by three elements: symmetric rich sharing scope from both participating Selves, preserve-tier-dominant conflict routing, and the conflict registry as primary output rather than residue. Four governance requirements make this configuration work: conflict classification configuration enables intelligent routing that distinguishes approach-difference conflicts from less informative types; full carry-through commitment ensures preserved conflicts reach both Selves' home governance; prompt post-mortem review converts the registry into actionable intelligence while context is fresh; and explicit absorption governance ensures that the event's findings about approach differences are addressed through directed selection rather than treated as routine ingestion. The competitive intelligence anti-pattern — resolve-tier-dominant routing in an event configured for comparison — collapses the intelligence payload the event was designed to produce. The operational test checks for preserve-tier routing in configuration, conflict density and classification in the registry, and carry-through annotations in both home substrates.