

Composition Pair 11: Competitive Intelligence and Post-Mortem Review

Series D — Phase D4, Note D4.12 | Defensive Publication #617

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 15, 2026

License: Creative Commons Attribution 4.0 International (CC BY 4.0)

Attribution

This work derives from and formalizes the inter-Self coordination architecture introduced in Inter-Self Coordination via Shared Substrate / Full Aspect Integration (Li, April 2026), the third paper in the CKS theory series following Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems (Li, April 2026) and The Instinct/Reasoning Separation Outside the Model (Li, April 2026).

Abstract

This note formalizes the governance requirements that arise when two Paper 3 operational commitments are applied simultaneously: Competitive Intelligence Full Aspect Integration (D2.65), which configures an FAI event to produce a conflict registry as its primary governance output, and Post-Mortem Governance (D2.39), which requires structured review after any FAI event to produce governance improvement records. The combination produces three non-obvious requirements not present in either commitment alone: (1) the post-mortem is the mandatory intelligence extraction mechanism for competitive events — without it, the conflict registry remains raw data and the event's primary purpose is not achieved; (2) the post-mortem requires a three-way conflict pattern classification (approach difference, approach quality, domain incompatibility) that cannot be automated and constitutes the human governance judgment component the composition specifically demands; and (3) the competitive intelligence report is a distinct, persistent governance artifact — separate from standard post-mortem improvement records — that should be stored in the substrate as a referenceable record of how the Self's governance compares to a given partner's. Any system claiming to implement governed AI governance benchmarking must address all three requirements.

1. Pair Identification

Commitment A — Competitive Intelligence FAI (D2.65): A preserve-tier-dominant FAI event configuration whose primary governance purpose is comparative analysis of two participating Selves' governance approaches. The event runs to completion with conflicts preserved as first-class state in the shared substrate under Paper 3's three-tier conflict-handling mechanism (Claim 3). The resulting conflict registry — a structured record of where the two Selves' governance configurations differ — is the primary output. The event is not designed to resolve those conflicts; it is designed to surface and record them as comparative data.

Commitment B — Post-Mortem Governance (D2.39): A structured review commitment applied after any FAI event completes. The post-mortem produces governance improvement records: structured findings

about what the event revealed regarding the Self's own governance practices, and what changes, if any, the review recommends. The commitment applies to all FAI event types; D2.39 establishes the review obligation and its recommended completion window.

The governance scenario activating both simultaneously: a competitive intelligence FAI event completes, leaving a dense conflict registry documenting governance approach differences between two Selves. Both commitments now apply — the event was specifically designed for comparative intelligence (Commitment A), and the post-mortem review obligation (Commitment B) is active. The question this composition raises is: what does post-mortem review mean, and what must it produce, when the event it is reviewing was a competitive intelligence event?

2. The Governance Scenario

Self A and Self B conduct a competitive intelligence FAI event. Their governance configurations differ in several respects: conflict-handling tier-selection rules, escalation-path configurations, action-layer expression thresholds, and DNA-layer review cadences. The event's preserve-tier dominance means that rather than resolving these differences through orchestration or escalating them to joint human authority, the shared substrate retains both Selves' configurations side-by-side, with each divergence point identified and annotated with provenance.

The event dissolves. Per Paper 3's four-locus evolution feed (Claim 4), content propagates to each Self's home substrate under home governance authority. The shared substrate exits operation, leaving behind a conflict registry of substantial density — not a list of errors or failures, but a structured record of how two governed AI systems approach the same governance problems differently.

Post-Mortem Governance (D2.39) is now active. Standard post-mortem review asks: given what this event produced, what governance improvements should result for this Self's own practices? That question applies here. But the competitive intelligence purpose of the event creates additional review obligations that the standard post-mortem structure does not address. This is the governance scenario the composition requires.

3. Non-Obvious Governance Requirements

Requirement 1 — Post-Mortem as the Intelligence Extraction Mechanism

In a cooperation FAI event, the post-mortem review is governance best practice: it converts the event record into structured improvement findings. The event itself has delivered its primary output — the merged governance content that each Self absorbs into its home substrate. The post-mortem adds analytical value on top of a completed primary purpose.

In a competitive intelligence FAI event, the relationship is different in kind, not just in degree. The conflict registry produced by the event is raw data. It records *that* the two Selves' governance configurations differ, and *where* they differ, with full provenance. It does not record *what those differences mean* for either Self's governance — whether a given conflict pattern indicates something worth learning from, something to deliberately not adopt, or a domain-specific divergence that carries no transferable lesson.

The post-mortem is the mechanism that performs this conversion. Without the post-mortem analysis, the preserved conflicts remain a record but not a governance learning. The event's primary purpose —

generating comparative intelligence — is not achieved by the conflict registry alone; it is achieved only when the post-mortem converts the conflict registry into interpretable findings.

This makes the post-mortem mandatory for competitive intelligence events in a stronger sense than for cooperation events. In cooperation events, skipping the post-mortem is a governance quality failure — the improvement opportunity is missed. In competitive intelligence events, skipping the post-mortem is a purpose failure — the event's designed output is not produced. Governance systems implementing competitive intelligence FAI must enforce the post-mortem as a required phase, not a recommended one.

Requirement 2 — Three-Way Conflict Pattern Classification

Standard post-mortem review asks: what governance improvements should result from this event? This question has a single analytical dimension: what does this event's record suggest this Self should do differently?

The competitive intelligence post-mortem requires an additional analytical layer with a specific three-way structure.

Class 1 — Approach difference. Some conflict patterns indicate that Self A and Self B have arrived at genuinely different but comparably defensible governance approaches to the same problem. Neither is governance-better than the other; they represent different solved forms of a shared design challenge. These patterns are worth learning from in the sense that they expand the reviewing Self's awareness of the governance design space.

Class 2 — Approach quality difference. Some conflict patterns indicate that one approach is governance-better than the other — more robust, more auditable, more consistent with the relevant operational context. These patterns carry a different learning signal: the reviewing Self should consider whether to adopt or adapt the superior approach.

Class 3 — Domain incompatibility. Some conflict patterns indicate that both approaches are valid but were designed for different operational contexts. The divergence does not reflect a design difference or a quality difference; it reflects that the two Selves operate in contexts where different governance configurations are appropriate. These patterns carry no transferable lesson; they mark the boundary of what is comparable between the two governance approaches.

This three-way classification cannot be automated. It requires governance practitioners who understand both Selves' operational contexts — the constraints that shaped each configuration, the problems each was designed to solve, the environments in which each operates. Automating the classification without that contextual knowledge would produce misreadings in all three directions: treating domain incompatibility as quality difference, treating quality difference as mere approach difference, or treating genuine approach differences as incompatibility and dismissing them.

The three-way classification is the human governance judgment component that the composition specifically requires. Neither Commitment A alone (which produces the conflict registry but imposes no review structure) nor Commitment B alone (which requires review but applies standard improvement-record framing) demands this classification. The composition creates the demand.

Requirement 3 — Competitive Intelligence Report as a Distinct Governance Artifact

Standard post-mortem review produces improvement records: structured findings about what this Self's governance should change, in what direction, on what timeline. These records are home-substrate artifacts — they inform the Self's own evolving practices and are governed within the Self's home authority structure.

The competitive intelligence post-mortem produces a second artifact that is distinct in kind: the competitive intelligence report. This report is the comparative analysis of the two Selves' governance approaches, organized around the three-way classification from Requirement 2, and documenting what the reviewing Self learned about how its governance compares to its partner's.

The competitive intelligence report is not a variant of the improvement record. Its purpose and reference horizon differ along two dimensions. First, the improvement record looks inward — it documents changes this Self should make. The competitive intelligence report looks outward — it documents how this Self's governance compares to another's, which is useful regardless of whether any internal changes result. A competitive intelligence event could reveal that the two Selves' governance approaches are comparably strong in different ways (predominantly Class 1 patterns), producing a rich competitive intelligence report and few or no improvement record entries.

Second, the improvement record is a point-in-time finding. The competitive intelligence report is a referenceable comparative analysis that gains value over time. When the same governance partner is encountered again in future FAI events — whether competitive intelligence or cooperation — the prior competitive intelligence report provides context: what was already known about this partner's governance approach, what classification was assigned to previous conflict patterns, and what has changed since the prior event. The report belongs in the substrate as a persistent, referenceable record, not archived as a one-time review output.

The requirement to produce and separately store the competitive intelligence report is specific to this composition. No single-commitment implementation of either D2.65 or D2.39 generates this artifact or its storage requirement.

4. Prior-Art Significance

Competitive intelligence is a well-developed practice domain. Post-mortem review is a well-developed governance practice. What is not established in prior art is the specific architectural pattern created by their composition in a governed AI coordination context: the post-mortem as the intelligence extraction mechanism, the three-way conflict pattern classification as a required analytical structure, and the competitive intelligence report as a persistent substrate artifact distinct from improvement records.

The prior-art significance is in the conjunction. A system that implements competitive intelligence benchmarking between AI systems without the three-way classification is not performing the analysis the composition requires — it is performing a simpler, less informative comparison that conflates approach differences, quality differences, and domain incompatibilities. A system that produces post-mortem improvement records from a competitive intelligence event without separately producing a competitive intelligence report conflates two governance artifacts with different purposes and different reference horizons. A system that treats the conflict registry as the output of a competitive intelligence event, rather than the input to its post-mortem analysis, has not completed the architecture the composition specifies.

Any system claiming to implement governed AI governance benchmarking — any architecture in which AI systems with distinct governance configurations are compared, and that comparison is governed — must address all three requirements: the mandatory post-mortem, the three-way classification, and the competitive intelligence report as a distinct persistent artifact. These requirements are not implementation details; they are the architecture.

5. Operational Test

For a competitive intelligence FAI event with post-mortem review, a governance observer can verify the composition's requirements by asking three questions.

Test (a) — Three-way conflict pattern analysis present: Does the post-mortem record show that each conflict pattern in the conflict registry was classified as approach difference, approach quality difference, or domain incompatibility? An absence of this classification indicates that the analysis collapsed the three-way structure into a simpler comparison, or that the classification was not performed and the post-mortem produced only standard improvement record content. The classification need not be recorded at the level of individual conflict entries — it may be reported in aggregate — but the analytical structure must be present and traceable in the post-mortem record.

Test (b) — Competitive intelligence report produced as a distinct artifact: Does the post-mortem produce two separately recorded outputs — standard improvement records (home-substrate artifacts documenting recommended changes to this Self's practices) and a competitive intelligence report (a comparative analysis documenting how this Self's governance compares to its partner's, organized around the three-way classification)? Conflation of the two into a single document, or production of only one, indicates the composition's output requirement was not met.

Test (c) — Competitive intelligence report completed within the review window: Is the competitive intelligence report produced within the post-mortem review window established by D2.39? Because the report is the primary intelligence output of the event, delay beyond the review window degrades the report's value — the context of the event is freshest immediately after dissolution — and constitutes a governance quality failure specific to competitive intelligence events. The report's completion timestamp should appear in the substrate record alongside the standard post-mortem completion record.

All three tests are verifiable from substrate records. No test requires access to the two Selves' internal governance configurations or to the content of the competitive intelligence report itself — the tests examine the structure and completeness of the governance record, not its substance.

End of Note D4.12.