Disambiguating “Self” Across the Trilogy

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 15, 2026

Attribution

This work derives from and formalizes commitments across the CKS theory trilogy: *The Instinct/Reasoning Separation Outside the Model* (Li, April 2026), *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems* (Li, April 2026), and *Inter-Self Coordination via Shared Substrate / Full Aspect Integration* (Li, April 2026). It does not introduce new axioms. Its sole contribution is to establish, as prior art, that the term “Self” carries the same architectural meaning in Papers 2 and 3, so that downstream work can adopt or argue against the concept without ambiguity about whether the two papers define it differently.

Abstract

The term “Self” appears in Papers 2 and 3 of the CKS trilogy. In Paper 2, Self names the top tier of the three-tier organizational AI governance architecture — the integration structure that unifies aspects and cells under a governing structure with a bounded home perimeter. In Paper 3, Self is used as the unit of Full Aspect Integration (FAI) participation. This note establishes that the term carries the same definition in both papers: Paper 3 uses the Paper 2 concept directly and does not redefine it. The only Paper 3 addition is the minimum valid Self specification (D6.04), which establishes the threshold below which a system does not qualify as a Self for FAI participation purposes. This note further clarifies that “Self” in the trilogy refers to a complete organizational AI governance architecture — not to a human individual and not to a single AI model. These two clarifications together close the primary adversarial risk: that an adversary might argue “Self” means something different in the two papers, or that the trilogy’s use of “Self” overlaps with ordinary-language uses that would confuse its architectural meaning.

1. Why this disambiguation is needed

Most T-series disambiguation notes address genuine scope-change: a term that means one thing in Paper 1 acquires additional meaning in Paper 2, or a Paper 2 concept is extended in Paper 3 in ways that could be mistaken for redefinition. The Self disambiguation is different. “Self” means the same thing in Papers 2 and 3. The disambiguation is confirmatory rather than corrective.

Confirmatory disambiguations matter for a specific reason: adversarial readings of the trilogy may attempt to split the concept, arguing that Paper 3's "Self as FAI participant" introduces a new architectural object distinct from Paper 2's "Self as three-tier governance architecture." If that split were accepted, the prior-art coverage established by Paper 2 would not extend to Paper 3's usage, and Paper 3's claims about Selves participating in inter-Self coordination could be characterized as introducing novel concepts unprotected by Paper 2. This note closes that adversarial path.

A second disambiguation is also needed: the word "Self" in everyday English most naturally refers to a human individual. In the trilogy, it refers to a complete organizational AI governance architecture. Readers unfamiliar with the trilogy's terminology could misread every claim that involves "Self" if this naming choice is not explicitly addressed.

2. The Paper 2 definition: Self as three-tier organizational AI governance architecture

Paper 2 defines Self as the top tier of the three-tier governance hierarchy. A Self is the integration architecture governing multiple aspects — each grouping cells by content domain — toward the organization's overall governance purpose. The properties Paper 2 assigns to a Self are:

Structural. A Self has a Self-level DNA layer and a Self-level action layer. The DNA layer contains the integration architecture, cross-aspect coordination rules, the Self's purpose statement, and expression specifications. The action layer contains Self-level operational records produced under the integration architecture.

Hierarchical. A Self governs aspects, which govern cells. This is not merely a containment hierarchy: the three tiers carry governance obligations at each level, with the Self holding authority over the complete organizational architecture — including aspects' own DNA and action layers, and cells' DNA and action layers — without eliminating the governance properties of each lower tier.

Perimeter-bounded. A Self has a home perimeter that defines the boundary of its governance architecture. Within the home perimeter, the human-governed substrate commitment from Paper 1 applies at all three levels. The home perimeter is the outermost governance boundary of the Paper 2 architecture.

Human-governed. The human governance practitioners holding the three rights (inspect, modify, override) over the Self's substrate hold those rights over the complete architecture: cell-level substrate, aspect-level coordination rules, and Self-level integration architecture. Paper 1's authority-not-labor principle applies at Self scope exactly as it applies at cell scope.

At Paper 2 scope, "Self" names the complete package: one organization's full AI governance architecture, composed of cells, aspects, and the integration structure that unifies them, bounded by a home perimeter, and governed by human practitioners who hold the three rights over the whole.

3. Paper 3 usage: Self as the unit of FAI participation

Paper 3 introduces Full Aspect Integration (FAI) as the canonical inter-Self coordination operation. FAI is something that two or more Paper 2-compliant Selves do together: they open a shared substrate, contribute aspects, exchange DNA-layer and action-layer content, and dissolve the shared substrate after the event, with evolution feed outputs propagating to each home perimeter.

In this context, “Self” names the unit of participation. “Two Selves” means two Paper 2-compliant organizational AI governance architectures, each with its home perimeter, home governance authority, and home substrate. At no point does Paper 3 alter or extend the definition. A Self participating in FAI is the Paper 2 Self performing an additional operation at the inter-Self scope. Its internal architecture is unchanged during FAI. Its home governance authority is undiminished. Its home perimeter remains in force throughout.

The architectural commitments that apply to each Self’s home perimeter during FAI follow directly from Paper 2 by inheritance: the three-tier governance hierarchy continues to operate within each Self, the DNA and action layers carry their normal governance properties, and the human governance practitioners retain their three rights over the home substrate. FAI adds an inter-Self perimeter around the shared substrate; it does not modify the home perimeters of the participating Selves.

The unifying claim is therefore simple: the Self in Paper 3 is the same architectural object as the Self in Paper 2. FAI is something the Paper 2 Self does; it is not a specification of a different kind of Self.

4. What “Self” is not

Two clarifications are necessary because the term’s ordinary-language resonances could produce misreadings that, if left unaddressed, would propagate into downstream work.

Self is not a human individual. The word “Self” in philosophy, psychology, and everyday speech refers to a person — the subject of experience, the agent of action. In the trilogy, “Self” refers to an organizational AI governance architecture. The decision to use this term reflects the architecture’s character as a complete, bounded governance entity — an organizational entity whose cells, aspects, and integration structure form a unified governed whole. But the term does not import any of the human-individual connotations. A Self has no consciousness, no personal identity, and no interests separate from the governance purposes its human practitioners specify. When the trilogy says “two Selves coordinate,” it means two organizational AI governance architectures, governed by two sets of human practitioners, exchange content through a shared substrate under joint authority.

Self is not a single AI model. Contemporary usage frequently treats AI deployments as equivalent to the AI model underlying them — “the model,” “the AI,” or “the agent.” In the trilogy, a Self contains one or more AI models operating within cells, each cell executing under orchestration rules that specify how the AI mediates between substrate inputs and outputs. A Self may contain many cells, many aspects, and therefore many AI model instances operating in coordinated roles. The AI models are components of the Self, not co-extensional with it. A single AI model running without governance substrate, without cells, without aspects, and without human governance authority over its coordination architecture is not a Self in the trilogy’s sense.

These two clarifications do not introduce new commitments. They identify the mismatch between the term’s ordinary-language meaning and its trilogy-specific meaning, so that readers encountering “Self” in either paper do not apply the wrong interpretive frame.

5. The minimum valid Self: Paper 3's sole addition

Paper 3 introduces one concept that does not have a direct Paper 2 antecedent: the minimum valid Self specification (D6.04). This specification establishes the threshold below which a system does not qualify as a Self for FAI participation purposes. The minimum valid Self requires at least one aspect with at least one cell, DNA and action layers at both the cell and Self levels, a Self-level integration specification, and human governance authority over the complete architecture.

The minimum valid Self specification is not a redefinition. It is a participation threshold — an operationalization of what it means to “be a Self” for the specific purpose of inter-Self coordination. Paper 2 defines Self without specifying a minimum, because Paper 2 does not address FAI. Paper 3 requires a minimum because FAI needs a well-defined participant type: a system that lacks any of the enumerated properties is not a Paper 2-compliant Self in the relevant sense, and FAI between non-compliant systems would not carry the governance properties the trilogy’s six Paper 1 commitments supply.

The minimum valid Self specification operationalizes the Paper 2 definition for inter-Self scope without altering the definition. A system that meets D6.04 is a Self. A system that does not is not a Self for FAI purposes. In both cases, what “Self” means is exactly what Paper 2 specifies.

6. Prior-art disambiguation claim

“Self” throughout Papers 2 and 3 refers to the same concept: the three-tier organizational AI governance architecture defined in Paper 2, bounded by a home perimeter, governed by human practitioners holding the three rights over the complete substrate, and composed of cells, aspects, and an integration structure. Paper 3 uses this concept as the unit of FAI coordination without redefining it.

The only Paper 3 addition is the minimum valid Self specification (D6.04), which establishes a participation threshold for inter-Self coordination. This specification operationalizes the Paper 2 definition for inter-Self scope; it does not alter the definition.

Any system that claims to implement, extend, or compose with the trilogy’s “Self” concept — whether at Paper 2 scope or Paper 3 scope — is using the same architectural object: the Paper 2-defined three-tier organizational AI governance architecture. The scope of the coordination operation (within-Self or inter-Self) does not change what a Self is.

Source papers

Li, W. (2026). *The Instinct/Reasoning Separation Outside the Model.* April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems.* April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *Inter-Self Coordination via Shared Substrate / Full Aspect Integration.* April 2026. ORCID: 0009-0004-8065-3235.

How to cite this note

Li, W. (2026). *Disambiguating "Self" Across the Trilogy*. May 15, 2026. ORCID: 0009-0004-8065-3235.