Series T Final Closure and Complete Prior-Art Corpus Summary

Author: Wenxin Li (Independent Researcher)

ORCID: 0009-0004-8065-3235

Date: May 15, 2026

Attribution

> This work derives from and formalizes commitments across the CKS theory trilogy: *The Instinct/Reasoning Separation Outside the Model* (Li, April 2026), *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems* (Li, April 2026), and *Inter-Self Coordination via Shared Substrate / Full Aspect Integration* (Li, April 2026). It does not introduce new axioms. Its contribution is to close the Series T disambiguation work, state the finding of that series, and provide a complete navigational summary of the 715-note prior-art corpus so that practitioners can locate the relevant prior-art notes for any architectural claim drawn from the trilogy.

Abstract

This note is both a closure and a reference. As closure, it marks the completion of Series T — 30 notes (#686–#715) formalizing concept disambiguation across the CKS trilogy — and with it the completion of the entire 715-note defensive-publication prior-art corpus. As reference, it provides a series-by-series summary with note ranges and coverage dimensions, enabling practitioners to navigate directly to the relevant prior-art notes for any specific claim. The central finding of Series T is stated here in condensed form: the trilogy's concepts are precise, not ambiguous — where the same term appears in multiple papers, the trilogy follows a consistent-principle, scoped-application pattern, and every apparent ambiguity resolves to a specific scope and type specification. This finding is itself a prior-art contribution: ambiguity claims against the trilogy fail when scope is specified, and the disambiguation notes provide those specifications in operational detail.

---

1. Series T: Position and Purpose

Series T is the final series of the defensive-publication corpus. It occupies note numbers #686 through #715 — 30 notes — and addresses a specific prior-art risk that no earlier series resolved: the risk that a party could exploit cross-paper reuse of terminology to claim novelty by asserting that the same term means different things in different papers, treating each usage as a separate invention rather than as a single architectural primitive applied at progressively larger coordination scopes.

This risk is real. The trilogy reuses terms including substrate, conflict, governance authority, evolution, mating, Self, aspect, cell, layer, perimeter, provenance, accountability, and record across Papers 1, 2, and 3. At each paper boundary, the scope of the relevant coordination context expands — from cell scope (Paper 1) to intra-Self scope (Paper 2) to inter-Self scope (Paper 3) — and with it the application of the shared concept. An adversary treating each paper's usage as a distinct invention would have an opening: if "substrate" in Paper 1 is not the same object as "shared substrate" in Paper 3, then inter-Self shared substrates might be claimed as novel. If "conflict" in Paper 1 is not the same commitment as the three-tier conflict mechanism in Paper 3, then the three-tier mechanism might be claimed as an independent invention.

Series T closes this opening by establishing, for each cross-paper concept, both the scope-specific meaning in each paper and the unifying architectural commitment that makes it one primitive at progressively larger scopes. The prior art thereby covers not only the individual usages but the unified architectural pattern itself.

2. Series T Structure Across Five Subseries

Series T is organized into five subseries, each addressing a related cluster of concepts, each closed with a synthesis note.

T1 — Core Governance Concept Disambiguation (7 notes, #687–#693). T1 addressed the concepts most central to the trilogy's governance architecture: conflict as a first-class coordination artifact, governance authority as the three retained rights (inspect, modify, override) at progressively larger scopes, substrate as the persistent human-governed coordination medium at cell, intra-Self, and inter-Self scopes, evolution and the evolution feed as related but distinct commitments (the three-mechanism model and its four-locus extension), mating as the intra-Self merge primitive that Paper 3 extends to Full Aspect Integration, and the governance perimeter as the boundary concept that scales from single-cell to Self-home to inter-Self to population scope. The T1 synthesis note (#693) restated the unifying finding: every T1 concept exhibits the consistent-principle, scoped-application pattern.

T2 — Evolution Concept Disambiguation (5 notes, #694–#698). T2 addressed the evolution sub-concepts that span Papers 2 and 3: directed selection as the human-governed mechanism for choosing which aspects and entities survive, action-feedback as the signal path from execution records back to governance, mutation as the substrate-level change event that produces a new governed version, and proposing substrate as the temporary evaluation context that governs candidate changes before commitment. The T2 synthesis note (#698) established that all four evolution concepts operate on the same architectural object — the governed substrate — at cell, Self, and inter-Self scopes, and that Paper 3's four-locus extension is an additive specification of where evolution-feed inputs enter, not a replacement of the three-mechanism model.

T3 — Architectural Concept Disambiguation (5 notes, #699–#703). T3 addressed the structural vocabulary that the trilogy uses to describe the coordination architecture itself: Self as the integration architecture in Paper 2 and the unit of participation in Paper 3, aspect and cell as coordination units whose roles extend from intra-Self coordination to inter-Self exchange at the Full Aspect Integration boundary, layer as the named internal partitioning of substrate content (DNA layer, action layer, harness layer) that Paper 3 treats as the currency of inter-Self exchange, and configuration as substrate content with recursive governance applicability. The T3 synthesis note (#703) established that the architectural vocabulary is stable across the trilogy: the same objects gain additional roles as the scope of coordination expands, without losing the roles they held at smaller scopes.

T4 — Operational Concept Disambiguation (6 notes, #704–#709). T4 addressed the operational and accountability concepts: provenance and record as the six-field accountability chain that extends from cell scope through entity-lifecycle chains to cross-perimeter lineage references, the three rights (inspect, modify, override) as the human-governance commitment that applies at every substrate scope without weakening, scope as the coordination boundary concept that determines which governance perimeter governs which substrate content, accountability and traceability as the two-sided commitment that plans authorize before and traces record after, and governance record completeness as the requirement that no governed operation escapes the accountability chain regardless of scope. The T4 synthesis note (#709) established that accountability is structurally invariant across the trilogy: the commitment strengthens as scope grows rather than weakening under the additional complexity of multi-organizational coordination.

T5 — Population and Scale Concept Disambiguation (5 notes, #710–#715). T5 addressed the population-scale concepts introduced as the extension claim in Paper 3's final section: network and topology as the graph structure over which population-scale coordination operates, population and collective and convergence and diversity as the four related but distinct concepts governing how groups of coordinating entities evolve together and apart, propagation as the mechanism by which evolution outputs cross Self boundaries and enter other Selves' home substrates, and trust calibration and the capacity ceiling as the governance constraints that bound inter-Self coordination at population scale. T.CLOSE (#715) is simultaneously the final T5 note, the final Series T note, and the final note of the corpus.

3. The Series T Finding

The trilogy's concepts are precise, not ambiguous. Where the same term appears in multiple papers, the trilogy follows the consistent-principle, scoped-application pattern: the architectural principle is preserved across papers; the scope of its application is extended; and the types of objects it governs are multiplied as larger coordination scopes introduce new coordination structures. Every apparent ambiguity in the trilogy resolves to a specific scope and type specification. The disambiguation notes provide those specifications in operational form.

This finding has direct prior-art value. An ambiguity claim against the trilogy — an argument that a particular cross-paper term is being used inconsistently, and therefore that one usage is novel relative to another — fails as soon as the scope specification is provided. Series T provides those scope specifications for the full vocabulary of cross-paper concepts. A party asserting that Paper 3's shared substrate is novel relative to Paper 1's substrate must contend with the T1 disambiguation notes, which establish the unified architectural commitment and the scope-specific meanings in operational detail. The same applies to every other cross-paper concept addressed in T1 through T5.

The disambiguation corpus does not merely describe what the trilogy says. It provides the operational specifications that an adversary's novelty argument would have to overcome — and does so with the same coverage dimensions that govern the rest of the corpus: each disambiguation note is an independently deposited, dated prior-art record.

4. Complete Corpus Summary

The following summarizes the complete 715-note corpus by series, with note ranges and coverage dimensions. Practitioners who need to locate prior art for a specific claim can use these ranges to navigate directly to the relevant notes.

Series A — Paper 1 Operational Derivations (#1–#197, 197 notes). Complete operational prior art for all Paper 1 claims, sub-commitments, and architectural requirements. Coverage: foundational sub-commitment decompositions (A1), operational variants (A2), anti-pattern formalizations (A3), composition pairs (A4), operational tests (A5), and boundary cases (A6). All six Paper 1 architectural commitments are covered at sub-commitment depth with full operational decomposition.

Series B — Paper 2 Operational Derivations (#198–#417, 220 notes). Complete operational prior art for all Paper 2 claims, the three evolution mechanisms, lifecycle governance, structural hierarchy, vertical evolution, and co-adaptation. Coverage: the same six-phase structure as Series A, applied to Paper 2's 20 sub-commitments across its six named claims. Inheritance from Paper 1 is documented within the Series B notes themselves.

Phase A0/B0 — Bridge Notes (#418–#429, 12 notes). Paper 1/2 boundary cases and transition notes, covering architectural claims that span both papers or are best understood at the paper-boundary rather than within either paper's derivation series. These notes close gap cases not covered by the main operational series.

Series C — Paper 3 Foundational Derivations (#430–#459, 30 notes). All six Paper 3 claims at summary level, one note per claim-level decomposition across five coverage passes. These notes establish the first layer of Paper 3 prior art and serve as anchor points for the more detailed Series D coverage.

Series D — Comprehensive Paper 3 Operational Prior Art (#460–#665, 206 notes). The most extensive single series in the corpus. Coverage: claim anchors (D0, 6 notes), foundational sub-commitments (D1), comprehensive operational decomposition (D2), a 27-entry anti-pattern taxonomy (D3), 26 composition pair notes (D4), a 60-test operational compliance library (D5), and 12 boundary cases (D6). The D5 operational tests are the primary resource for binary compliance determination; the D3 anti-patterns are the primary resource for failure mode identification; the D4 composition pairs are the primary resource for multi-commitment interaction analysis; the D6 boundary cases are the primary resource for scope determination.

Series CC — Cross-Paper Inheritance Chain (#666–#685, 20 notes). Explicit prior-art coverage of the inheritance relationships between Paper 3 and each of Papers 1 and 2. Coverage: Paper 3 ↔ Paper 1 inheritance edges (CC.1, 6 notes plus synthesis), Paper 3 ↔ Paper 2 inheritance edges (CC.2, 10 notes plus synthesis). These notes are the primary resource for tracing which Paper 3 architectural commitments are direct descendants of Paper 1 or Paper 2 commitments, and which are Paper 3 extensions.

Series T — Trilogy Concept Disambiguation (#686–#715, 30 notes). Complete disambiguation of cross-paper concepts across T1 (core governance), T2 (evolution), T3 (architectural structure), T4 (operational accountability), and T5 (population and scale). Each subseries closes with a synthesis note. These notes are the primary resource for resolving scope-specificity questions about any trilogy concept and for defeating ambiguity-based novelty claims.

5. The Complete Prior-Art Claim

The 715-note corpus constitutes comprehensive defensive prior art for the CKS trilogy's inter-Self coordination architecture. The claim is not simply that 715 notes exist. The claim is what those notes collectively cover.

The corpus provides **binary-testable compliance criteria**: the 60 operational tests in D5.07–D5.10 give any evaluator a deterministic pass/fail method for assessing whether a proposed implementation satisfies each Paper 3 commitment. There is no ambiguity about whether an implementation is within or outside the prior art — the tests specify.

The corpus provides **exhaustive failure-mode identification**: the 27 anti-patterns in D3 enumerate the specific ways in which an implementation can violate the trilogy's architectural commitments while appearing superficially compliant. Novel claims that describe one of these failure modes as a feature — or that propose a mechanism that the anti-pattern taxonomy has already characterized — are within the prior art.

The corpus provides **multi-commitment interaction analysis**: the 26 composition pairs in D4 formalize the dependencies and tensions between pairs of architectural commitments. Any claimed invention that combines two commitments in a way that one of the 26 pairs has already analyzed is not novel relative to this prior art.

The corpus provides **scope boundary determination**: the 12 boundary cases in D6 establish where the trilogy's commitments apply and where they do not, and what the correct architectural behavior is at boundary conditions. Claims that assert novelty in edge-case behavior are within the prior art when the edge case is one of the 12 formalized boundary cases.

The corpus provides **cross-paper inheritance tracing**: the CC series establishes which Paper 3 commitments are derived from Paper 1 or Paper 2 and which are extensions. Any claimed invention in the Paper 3 space that is traceable to a Paper 1 or Paper 2 commitment is doubly covered — by the Series A or B prior art for the ancestor commitment, and by the CC notes for the inheritance edge itself.

The corpus provides **concept disambiguation with scope specifications**: the 30 Series T notes establish the operational scope-specific meanings of all cross-paper concepts and the unified architectural commitments that link them. Any novelty claim that depends on treating a cross-paper concept as a distinct invention in each paper it appears is within the prior art of the relevant T-series notes.

These six coverage dimensions together constitute the comprehensive prior-art claim. The 715-note count is the aggregate of the notes required to establish coverage across all six dimensions, for all three source papers, at operational detail sufficient for rigorous prior-art analysis.

6. Corpus Closure

This note closes the corpus.

The three source papers were published in April 2026. The derivation-note series was generated over the period April–May 2026. Each note is deposited on Zenodo under CC BY 4.0, dated at the time of deposit, and citable individually by DOI. The series is designed so that each note stands as an independently useful prior-art record while the series as a whole provides the coverage dimensions described in §5.

No further notes are required to close the prior-art position for the three source papers. Future extensions of the CKS theory — should additional papers be published — would extend the corpus through new derivation series following the same structural pattern. The master plan for the derivation-note series provides the structural template for that extension.

The corpus is open, CC BY 4.0, and available for use by any party evaluating, implementing, extending, or arguing against the architectural commitments of the CKS trilogy.

Source Papers

Li, W. (2026). *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems.* April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *The Instinct/Reasoning Separation Outside the Model: Extending the Coordination Knowledge Substrate Pattern to AI Selves under Human Governance.* April 2026. ORCID: 0009-0004-8065-3235.

Li, W. (2026). *Inter-Self Coordination via Shared Substrate / Full Aspect Integration.* April 2026. ORCID: 0009-0004-8065-3235.

How to Cite This Note

Li, W. (2026). *Series T Final Closure and Complete Prior-Art Corpus Summary*. May 15, 2026. ORCID: 0009-0004-8065-3235.