

The Instinct/Reasoning Separation Outside the Model: Extending the Coordination Knowledge Substrate Pattern to AI Selves under Human Governance

Wenxin Li *Independent Researcher* ORCID: 0009-0004-8065-3235

April 2026

Abstract

Paper 1 (Li, 2026) defended the Coordination Knowledge Substrate (CKS) as a design pattern at the coordination and decision layer, with the substrate handling coordination and governance, the LLM handling high-dimensional reasoning, and the division drawn at the governance boundary at cell scope. This paper extends the CKS pattern from cell scope to AI Self scope and toward the enterprise brain Self as an architecturally coherent design pattern extending Paper 1's cell concept, defending six architectural commitments.

The first separates fast-pattern instinct (the LLM, functioning as a System-1 analogue) from deliberate reasoning (the CKS substrate, functioning as a System-2 analogue) into independently-evolving layers composed into one Self under unified human governance. The second commits structural modularity at every level: cells composing into aspects, aspects into Selves, with relational rather than intrinsic role membership; DNA and action as distinguishable substrate layers within every cell with distinct governance semantics; expression as governed selection over which substrate content activates per goal. The third commits lifecycle as governed primitives — birth, mating with three pattern variants, and death with two distinct types — operating at every level of composition. The fourth commits evolution as three mechanisms in productive tension across multiple levels and on horizontal and vertical axes: instinct evolution operating on the LLM and infrastructure layer, DNA evolution operating on the orchestration substrate, and action-feedback evolution closing the loop from recorded experience back into governed substrate refinement. The fifth commits human governance as multi-shaped — distinct governance shapes co-determined with distinct mechanisms, the instinct/reasoning boundary as governed substrate content, and distinct death-type governance processes. The sixth offers the vision claim that the enterprise brain Self is architecturally coherent as a design pattern extending Paper 1's cell concept — a functioning unit composed of cells through aspects under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability where integration-glued enterprise architectures take a different structural approach. The vision is offered pending debate and observation in real-world enterprise deployment.

The contribution is a design pattern, not an empirical evaluation. The unifying architectural move is the separation of instinct and reasoning into independently-evolving layers under unified human governance — Paper 1's *outside-the-model vs. inside-the-model* spine extended through Self lifecycle, evolution, and into the enterprise brain Self. Paper 2's claims at cell scope inherit from Paper 1's proof-of-concept at cell scope; the architectural commitments at aspect, Self, and enterprise brain scope are defended architecturally and rendered concretely through a thought experiment in healthcare continuing the POC's lineage, with empirical validation in real-world enterprise deployment as downstream work. Paper 2 sits at the intersection of dual-process AI, structural-modularity-as-architectural-commitment, biology-of-evolvability, governance-as-architecture, and enterprise AI deployment lineages, articulating a design pattern available to be adopted, varied, composed with adjacent patterns, or argued against.

Table of Contents

1. Introduction

- 1.1 The architectural problem at Self scope
- 1.2 Contribution statement — extending CKS from cell to Self under human governance
- 1.3 Roadmap

2. Framing and Scope

- 2.1 The cross-claim spine
- 2.2 Biology as bounded load-bearing terminology
- 2.3 What Paper 2 is not

3. Related Work

- 3.1 Dual-process AI
- 3.2 Modular AI architectures and governance-as-architecture
- 3.3 Biology of evolvability and AI evolution
- 3.4 Enterprise AI deployment
- 3.5 Where Paper 2 sits

4. Claim 1 — Separation of instinct and reasoning into independently-evolving layers under unified human governance

- 4.1 Claim statement and the extension from cell to Self scope
- 4.2 The named foil — conventional LLM-only architecture and its structural limits
- 4.3 Closest-neighbor engagement — SOFAI and the dual-process AI lineage
- 4.4 Three differentiation axes
- 4.5 Reviewer-objection responses
- 4.6 Paper 1 inheritance and downstream placement

5. Claim 2 — CKS Selves as architecturally modular by commitment rather than by evolution

- 5.1 Claim statement and the foundational-modularity framing
- 5.2 Multi-level composition with relational roles
- 5.3 DNA/action layer distinction within every cell
- 5.4 Expression as governed selection
- 5.5 Reviewer-objection responses
- 5.6 Paper 1 inheritance and downstream placement

6. Claim 3 — Lifecycle operations as governed primitives at every level of CKS Selves

- 6.1 Claim statement and the lifecycle-as-governed-primitive framing
- 6.2 Birth as governed origination
- 6.3 Mating as one governable primitive over content combination
- 6.4 Death as governed retirement with two distinct architectural results
- 6.5 Reviewer-objection responses
- 6.6 Paper 1 inheritance and downstream placement

7. Claim 4 — Evolution as three mechanisms in productive tension under unified human governance

- 7.1 Claim statement and the productive-tension framing
- 7.2 The three mechanisms
- 7.3 Multi-level simultaneous evolution
- 7.4 Dual-axis evolution — horizontal and vertical
- 7.5 Reviewer-objection responses
- 7.6 Paper 1 inheritance and downstream placement

8. Claim 5 — Human governance across CKS evolution as multi-shaped commitment

- 8.1 Claim statement and the authority-vs-labor extension to evolutionary dynamics
- 8.2 Per-mechanism governance shape
- 8.3 The instinct/reasoning boundary as governed substrate content
- 8.4 Distinct death-type governance processes
- 8.5 Reviewer-objection responses
- 8.6 Paper 1 inheritance and downstream placement

9. Claim 6 — The enterprise brain Self as architecturally coherent design pattern extending Paper 1's cell concept

- 9.1 Claim statement and the vision-claim framing
- 9.2 The wholeness move
- 9.3 The topology consequence move
- 9.4 Three supporting commitments
- 9.5 Reviewer-objection responses
- 9.6 Paper 1 inheritance, Claims 1–5 inheritance, and closing

10. The Enterprise Brain: A Thought Experiment in Healthcare

- 10.1 Setting and the cell-aspect-Self structure rendered
- 10.2 Architectural commitments threaded through the rendering
- 10.3 Honest limits and positioning

11. Discussion and Implications

- 11.1 What the composed architecture commits to
- 11.2 Scope and limits
- 11.3 Open questions for future work

12. Conclusion

References

1. Introduction

1.1 The architectural problem at Self scope

Paper 1 defended a design pattern at the scope of a single coordination cell: a unit of work in which substrate handles coordination and governance, the LLM handles high-dimensional reasoning, and the division between them is drawn at the governance boundary rather than the capability boundary. The pattern stands on its own at single-cell scope. It was demonstrated through one proof-of-concept in regulated-industry reporting, instantiated in commodity spreadsheet infrastructure with a commercially available LLM integration — a deliberately modest setting chosen so the architectural commitments could be defended without depending on bespoke tooling or unusual deployment conditions. At that scope, the cell holds; the design pattern is available to be adopted, varied, composed with adjacent patterns, or argued against on its own terms.

What Paper 1 did not address is what happens when the cell stops being the largest architectural object the design pattern reasons about. Three scope-shift directions sit outside Paper 1 and motivate Paper 2. First, cells compose into larger structures — groupings of cells that operate as coordinated wholes around shared purposes, and integrated arrangements that hold multiple such groupings together as one functioning intelligence. Second, those structures persist through life and change over time under continued human governance, raising architectural questions about what lifecycle primitives the design pattern must commit to and what those primitives must support. Third, AI deployment moves from a single coordination cell at one workflow to an enterprise scope where multiple coordinated wholes operate alongside human organizational structures, where coherence across the wholes becomes its own architectural concern and the cost of integrating cell-level work into a recognizable functioning unit is no longer negligible.

The architectural question Paper 2 takes up is what an LLM-era AI system that is governed across its lifecycle and evolution looks like at this larger scope. The question is real, and it is sharpening rather than fading as LLM capability advances. Contemporary LLM-only architectures collapse coordination, governance, conflict-handling, and reasoning into one model in which errors can only be addressed by modifying the model itself. As LLM capability sharpens, the temptation to absorb more of the architecture into the model intensifies — every workflow component the model can plausibly subsume becomes a candidate for absorption, and every absorption further centralizes the locus of error correction inside weights humans cannot inspect, modify, or version-control with the granularity that governance requires. The commitments Paper 1 named at cell scope — coordination, governance, conflict-handling,

the reasoning that humans should govern regardless of capability — face exactly this absorption pressure as the system grows.

Paper 2’s response is not that the absorption pressure should be resisted as such; the LLM’s capability is essential to every scope at which CKS architectures operate, and a sharper LLM makes the Self more capable rather than less. The response is that some commitments — coordination, governance, conflict-handling, the reasoning humans should govern regardless of capability — must remain architecturally separable from the model layer for the same reasons Paper 1 identified at cell scope, but now extended to a complete AI Self with a lifecycle and evolutionary dynamics. The unifying architectural move is the separation of fast-pattern instinct from deliberate reasoning into independently-evolving layers under unified human governance: the LLM carries the fast pattern-matched responses, the CKS substrate inherited from Paper 1 carries the explicit human-governed reasoning, and the two layers compose into one Self with their independent evolutionary dynamics held safe by governance commitments specified at every level of composition. Paper 1’s *outside-the-model vs. inside-the-model* spine extends through Paper 2 unchanged: coordination, governance, and reasoning live outside the LLM in the persistent, human-governed substrate; fast-pattern instinct lives inside the LLM; and the architecture’s value at Self scope comes from the same principled separation Paper 1 defends at cell scope.

1.2 Contribution statement — extending CKS from cell to Self under human governance

Paper 1 (Li, 2026) introduced the Coordination Knowledge Substrate as a design pattern at the coordination and decision layer for AI systems, with two levels of structure: the substrate as the atomic structured representation of one informational task, and the cell as the functional whole that aggregates substrates with human-authored orchestration rules into a coordinated unit of work. Paper 1 defended six architectural commitments — substrate-based hybrid composition with a governance boundary, conflict preservation at the substrate level, human governance over content and rules, AI as substrate mediator, tool-agnosticism at the substrate layer, and linear-cost storage and composition. The cross-claim spine Paper 1 deploys throughout, *outside-the-model vs. inside-the-model*, names where coordination, governance, and decision state live in CKS architectures and where they do not. Three things from Paper 1 are what readers need to read Paper 2 productively: the substrate/cell two-level structure as foundational architectural commitment; the six commitments as the design pattern’s content held throughout Paper 2’s architecture; and the outside-the-model spine as the architectural distinction Paper 2 builds on at higher scopes.

Paper 2’s contribution is six architectural commitments defended through six claims. The first separates fast-pattern instinct (the LLM, functioning as a System-1 analogue) and deliberate reasoning (the CKS substrate, functioning as a System-2 analogue) into independently-evolving layers composed into one Self under unified human governance. The second commits structural modularity at every level — cells composing into aspects, aspects into Selves, with relational rather than intrinsic role membership; DNA and action as distinguishable substrate layers within every cell; expression as governed selection over which substrate content activates per goal. The third commits lifecycle as governed primitives — birth, mating with three pattern variants, and death with two distinct types — operating uniformly at every level. The fourth commits evolution as three mechanisms in productive tension across multiple levels and on horizontal and vertical axes: instinct evolution operating on the LLM and infrastructure layer, DNA evolution operating on the orchestration substrate, and action-feedback evolution closing the loop from recorded experience back into governed substrate refinement. The fifth commits human governance as multi-shaped — distinct governance shapes co-determined with distinct mechanisms, the instinct/reasoning boundary as governed substrate content, and distinct death-type governance processes. The sixth offers the vision claim that the enterprise brain Self is architecturally coherent as a design pattern extending Paper 1’s cell concept — a functioning unit composed of cells through aspects under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability where integration-glued enterprise architectures take a different structural approach — pending debate and observation in real-world enterprise deployment.

The contribution genre is design-pattern writing in the same sense Paper 1 inhabits: an articulation of a recurring structure of design problem and design response, named so that subsequent work can adopt it, vary it, compose it with adjacent patterns, or argue against it. Paper 2 is not a new system, not a vendor product, not a novel algorithm, and not an empirical evaluation at scale.

1.3 Roadmap

The paper is organized as a single integrated theory paper rather than as Paper 1's three-part core-and-extension structure. Paper 2's six architectural commitments compose into one design pattern at Self scope, and the cross-claim work each commitment does in supporting the others reads more naturally as one arc than as separated parts.

Section 2 frames the cross-claim spine, the role of biology vocabulary as bounded load-bearing terminology, and what Paper 2 is and is not. Section 3 positions Paper 2 within five recognized lineages — dual-process AI, modular AI architectures and governance-as-architecture, biology of evolvability and AI evolution, enterprise AI deployment, and the configuration these lineages do not jointly occupy — without re-engaging the per-claim closest-neighbor work the claim sections develop in their own territories.

Sections 4 through 9 defend the six architectural commitments in turn. Sections 4 through 8 develop the five core commitments — separation of instinct and reasoning into independently-evolving layers (§4); structural modularity at every level with relational role membership and the DNA/action layer distinction within every cell (§5); lifecycle as governed primitives operating uniformly at every level (§6); evolution as three mechanisms in productive tension on horizontal and vertical axes (§7); and human governance as multi-shaped with distinct governance shapes co-determined with distinct mechanisms (§8). Section 9 offers the vision claim that the enterprise brain Self is architecturally coherent as a design pattern extending Paper 1's cell concept — a functioning unit composed of cells through aspects under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability where integration-glued enterprise architectures take a different structural approach — pending validation through real-world enterprise deployment.

Section 10 renders the §9 vision claim concretely through a thought experiment in healthcare, continuing Paper 1's POC lineage at enterprise brain scope. The thought experiment is an architectural rendering rather than empirical demonstration; its role is to show what the composed architecture looks like when realized at enterprise scope, not to settle whether realization at scope is achievable in any particular deployment. Section 11 reflects on what the composed architecture commits to, what it does not commit to, scope limits, and open questions for future work. Section 12 closes.

2. Framing and Scope

2.1 The cross-claim spine

The cross-claim spine running through Paper 2's six claims is the separation of fast-pattern instinct from deliberate reasoning into independently-evolving layers composed into one Self under unified human governance. Each claim either establishes structures the separation requires, develops dynamics the separation enables, or specifies governance properties that keep the separation safe. The instinct layer is the LLM operating as a System-1 analogue — pattern-matched, fast-path, no-reasoning-needed responses. The reasoning layer is the CKS substrate operating as a System-2 analogue — explicit, deliberate, human-governed processing of what humans choose to encode. The two layers compose into one Self under unified human governance, with the LLM's capability essential at every scope and the substrate providing the architecture that makes the LLM's capability work in coordinated, governed, evolving systems. The hybrid commitment Paper 1 §2 establishes at cell scope — substrate handles coordination and governance, LLM handles high-dimensional reasoning, division at the governance boundary — extends throughout Paper 2 as the foundation for the instinct/reasoning separation at Self scope.

2.2 Biology as bounded load-bearing terminology

Paper 2 deploys biology vocabulary — DNA, action, expression, birth, mating, death, instinct, evolution, lineage — where it does specific architectural work that the surrounding architectural development carries forward. The biology vocabulary functions as bounded load-bearing terminology rather than as decorative analogy: each term carries architectural content the claim sections specify, with the biology parallel earning reader comprehension because the structural shape it names — cells composing into larger structures, life cycles, evolutionary mechanisms operating

on heritable content — is intuitive to absorb. The vocabulary’s interpretation is bounded by the architectural commitments Claims 2 through 5 supply, with terms like *DNA* (orchestration substrate carrying Paper 1’s specific commitments), *mating* (governed substrate-content combination across three pattern variants), and *death* (governed retirement with two distinct architectural results) carrying CKS’s specifications rather than biology’s referents. CKS differs from biology at multiple specific points where design choices diverge from biology’s evolved constraints: directedness alongside undirectedness (biology has only mutation; CKS has both directed selection and undirected mutation in productive tension); modularity by architectural commitment rather than evolved (biology had to evolve modular architectures over deep time; CKS is modular by commitment from instantiation); reversibility (biology’s extinct species cannot return; CKS’s archived cells remain addressable and reactivatable); cross-lineage compatibility (biology has species barriers; CKS has orchestration-governed cross-lineage combinations); structural reorganization on operational timescales (biological body plans take millions of years to evolve; CKS Selves restructure through governed reorganization). The differences reflect different conditions — biology’s constraints emerged from undirected evolution under fitness pressure; CKS is committed by design under human governance — and the architectural commitments themselves are CKS’s own, defended in their specific configurations across Claims 1 through 6.

2.3 What Paper 2 is not

Six disambiguations bound Paper 2’s scope. Paper 2 is not a replacement for LLM-based AI: the LLM instinct layer is essential at every scope, and Paper 2 defends the hybrid composition pattern at Self scope continuing Paper 1’s hybrid commitment rather than proposing an alternative. Paper 2 is not a rigid three-level hierarchy: structural roles are relational and purpose-defined, with the same underlying CKS artifact able to participate as cell, as part of an aspect, or in multiple aspects simultaneously depending on purpose, and cross-level access patterns enabled where purpose requires. Paper 2 is not automatic evolution: all three evolution mechanisms operate under human governance in different shapes, with mutation integrated through verification, directed selection human-governed, and action-feedback evolution human-mediated. Paper 2 is not a metaphor dressed up as architecture: the biological progression from cell through aspect to Self is architectural commitment, not decorative analogy, with the biology vocabulary functioning as bounded load-bearing terminology while the architectural work is CKS’s own. Paper 2 is not about inter-Self autonomous dynamics: scope ends at the enterprise brain Self under unified human governance, with how multiple enterprise brain Selves interact across organizational boundaries left as a class of extension theories that would each stand on its own claim set rather than on arguments Paper 2 develops. Paper 2 is not empirical demonstration at scale: it is a design-pattern paper in the genre of Paper 1, defending architectural commitments and rendering the enterprise brain vision concretely through a thought experiment continuing Paper 1’s POC lineage, with empirical validation in real-world enterprise deployment as downstream work.

3. Related Work

Each of Paper 2’s six claims engages closest-neighbor work in its specific territory, and the per-claim engagement is where the detailed differentiation work lives. This section integrates those engagements into paper-level positioning that shows where Paper 2 sits within recognized lineages, without re-engaging the per-claim closest-neighbor work the claim sections develop in their own territories. Five lineages converge on the architectural territory Paper 2 occupies; the subsections that follow name where Paper 2 sits within each and, in §3.5, what configuration the lineages do not jointly occupy.

3.1 Dual-process AI

The closest research lineage for Paper 2’s foundational instinct/reasoning separation is the dual-process AI tradition. SOFAI (Booch et al. 2021), PlanSOFAI (Fabiano et al. 2024), and the broader S1/S2 architectural family commit to separating fast pattern-matching from deliberate reasoning at the component level. Predating this contemporary lineage, the cognitive-architectures tradition has decades of peer-reviewed work on dual-process structures: SOAR (Newell, Laird) with its working-memory architecture, ACT-R (Anderson) with its declarative/procedural distinction, CLARION (Sun) with its explicit/implicit knowledge separation, and ICARUS (Langley) with its knowledge-and-skill memory split each commit to architectural separations of fast and deliberate cognition that anticipate the dual-process AI

tradition’s component-level commitments. The 2024–2026 expansion of the lineage is substantial and active: Bengio’s System-2 Deep Learning research program keeps S2 inside neural architectures through attention, working-memory analogues, and world models; Reasoning on a Spectrum (Ziabari et al. 2025) produces S1-aligned and S2-aligned LLM variants through alignment; D-Mem places fast and slow memory paths in dual-process framing; and the System Implications of Neuro-Symbolic-Probabilistic Architectures (ACM 2025) research direction recommends independent versioning, monitoring, and deployment of layered components. Adjacent to dual-process AI on the conceptual side, the extended-cognition tradition (Clark and Chalmers 1998) and distributed cognition (Hutchins 1995) frame cognitive processes as extending beyond the individual into external structures — a framing CKS’s outside-the-model commitment inherits at AI substrate scope without depending on the lineage’s philosophical foundations. The lineage supplies established warrant for the architectural shape Paper 2’s Claim 1 commits to.

What Paper 2 adds within the lineage — developed in §4 — is the configuration in which the System-2-analogue layer lives in a persistent human-governed coordination substrate in Paper 1’s specific sense, motivated by governance commitments rather than by capability gaps, with independent evolvability of the two layers as architectural requirement rather than incidental modularity.

3.2 Modular AI architectures and governance-as-architecture

Modular composition has decades of precedent in software architecture (microservices, hierarchical state machines, feature-flag governance, configuration management). Within AI systems specifically, the recent emphasis on agent frameworks and multi-agent composition (AutoGen, LangGraph, the Microsoft Agent Framework, the Claude Agent SDK, A2A “agent cards”) foregrounds modular construction. Foundational multi-agent systems literature on agent organization architectures — MOISE+ (Hannoun and Boissier) on organizational structures for MAS, AGRE/AGREEN (Ferber and Gutknecht) on agents-groups-roles formalisms, Demazeau’s Vowels framework — supplies organizational primitives for multi-agent composition that contemporary agent frameworks rebuild in domain-specific form. Adjacent to this is the governance-as-architecture direction in which governance content is treated as architectural object rather than overlay — the agent-memory cluster around Epsilla AgentStudio, Graph Digital, Atlan’s metadata control plane, and Unstructured.io’s separation of governance and observability articulates this in industry voice; Cognitive Core (Seck 2026) approaches it from architecture-time configuration; object-capability formalisms (Miller’s *Robust Composition*) and qua-entity formalisms (Steimann 2000; OntoUML’s role stereotype) supply foundational architectural primitives for relational structural roles.

Paper 2’s Claims 2 and 5 sit within these lineages. What Paper 2 adds — developed in §5 and §8 — is the configuration in which modularity at every level is architectural commitment rather than evolved property, the DNA/action distinction within every cell as governable layers within Paper 1’s substrate framework, and governance-shape co-determined with mechanism shape rather than overlaid as mechanism-agnostic policy.

3.3 Biology of evolvability and AI evolution

The Major Transitions in Evolution framework (Maynard Smith and Szathmáry 1995), the Extended Evolutionary Synthesis position on inclusive inheritance (Jablonka and Lamb; Laland and Uller), evo-devo work on body-plan evolution (Davidson and Erwin 2006), and evolvability research on modularity-as-evolvability-enabler (Hartwell et al. 1999; Kirschner and Gerhart 2005; Wagner) supply the conceptual scaffold Paper 2’s Claims 3 and 4 deploy for lifecycle and evolution. On the AI side, evolutionary computation (Holland; Goldberg; Rechenberg and Schwefel’s evolution strategies; Hansen’s CMA-ES; Koza’s genetic programming), open-ended evolution and quality-diversity (Lehman and Stanley’s novelty search; MAP-Elites; Soros and Stanley’s open-endedness conditions), neural architecture search (NEAT; DARTS; AutoML-Zero), hierarchical evolutionary algorithms (Potter and De Jong’s CCEA; Vanchurin et al.’s multi-level learning frame), and the recursive-self-improvement direction (Schmidhuber’s Gödel Machine; the broader RSI discourse) supply lineages of architectural alternatives for AI evolution.

What Paper 2 adds — developed in §6 and §7 — is the configuration of three evolution mechanisms in productive tension under unified human governance, with directedness alongside undirectedness, structural reorganization on operational timescales, lineage primitives operating at every level of composition concurrently, and the action-feedback loop closing through governed substrate-edit machinery rather than autonomous dynamics.

3.4 Enterprise AI deployment

The contemporary literature on enterprise AI deployment frames the architectural problem at organizational scope. Major-vendor enterprise AI deployment patterns — Azure AI Foundry, Vertex AI, AWS Bedrock with SageMaker, Salesforce Agentforce, Anthropic Enterprise, OpenAI Enterprise — commit to multi-tenant model-as-a-service or control-plane-with-multi-model-orchestration with governance overlays. Multi-LLM gateway architectures (LiteLLM, OpenRouter, Portkey, Agentgateway) treat multi-provider deployment as resilience commitment rather than as substrate commitment.

A second cluster of lineages reaches enterprise AI architecture from broader software-architectural and organizational sources. Team Topologies (Skelton and Pais 2019) commits team boundaries and interaction modes as first-class architectural design variables; the microservices-vs-modular-monolith literature (2018–2026) treats coordination cost as architectural concern; the distributed systems and SRE lineage (Lamport; Brewer’s CAP; Vogels; Google SRE) treats failure-domain isolation as architectural commitment with engineering patterns enforcing intent. The broader organizational design lineage — Galbraith’s star model and lateral capability framing, Mintzberg’s organizational structures and configurations, Argyris and Schön’s organizational learning, Senge’s learning organization — frames how organizations are architected as systems with explicit structural and adaptive design variables, supplying the architectural vocabulary of organizations into which AI deployment compositions are placed. The MIT Sloan Intelligent Choice Architectures direction positions decision-rights and meta-decision-rights as evolving alongside AI choice-environments. Constitutional architectures (Bai et al. 2022; contemporary corporate-constitution writing) frame deployment-scope governance as substrate that composes with provider-side training-time infrastructure.

A third cluster characterizes the topological shape traditional enterprise architectures take, and the substrate-sharing alternatives that have emerged alongside it. Enterprise integration patterns (Hohpe and Woolf’s *Enterprise Integration Patterns* canon) characterize the structural shape — point-to-point mappings, Canonical Data Model translation layers, asynchronous messaging — that traditional enterprise architectures share. Adjacent substrate-sharing lineages — Inmon and Kimball’s data warehousing, the Databricks/Iceberg/Delta Lake lakehouse cluster, the post-microservices retrospective with the Amazon Prime Video monolith-rebuild case as its popular crystallization — name substrate-sharing as architectural axis with explicit trade-offs. What Paper 2 adds — developed in §9 — is the configuration in which the per-Self architectural commitments §4–§8 establish carry into the enterprise brain Self as a functioning unit, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability where integration-glued enterprise architectures take a different structural approach.

3.5 Where Paper 2 sits

Paper 2 sits at the intersection of these five lineages, not as a single contribution to one lineage. The dual-process AI tradition supplies the architectural shape of the foundational separation; structural-modularity-as-architectural-commitment supplies the structural vocabulary; governance-as-architecture supplies the governance vocabulary; biology-of-evolvability and AI-evolution lineages supply lifecycle and evolution scaffolding; enterprise AI deployment supplies the application context. Paper 2’s contribution is the configuration in which all of this is articulated as one design pattern at AI Self scope, building on Paper 1’s substrate-cell foundation through the cross-claim spine §2.1 names. Each claim’s per-claim related-work engagement establishes the contribution within its specific territory; this section’s contribution is to name where Paper 2 sits at the paper level, in the empty configuration the lineages do not jointly occupy.

4. Claim 1 — Separation of instinct and reasoning into independently-evolving layers under unified human governance

4.1 Claim statement and the extension from cell to Self scope

Claim 1 states Paper 2’s foundational architectural commitment: a CKS-governed AI Self separates fast-pattern instinct from deliberate reasoning into two independently-evolving layers composed into one Self under unified human governance. The instinct layer is the LLM, operating as a System-1 analogue — pattern-matched, fast-path,

no-reasoning-needed responses. The reasoning layer is the CKS substrate inherited from Paper 1 (Li, 2026), operating as a System-2 analogue — explicit, deliberate, human-governed processing of what humans choose to encode. Conventional LLM-only AI systems collapse these into one layer in which the model is the entire system, which means errors can only be addressed by modifying the model itself. CKS keeps them separate, and the separation does specific architectural work: the reasoning layer can route around bad instinct, Paper 1’s conflict-preservation commitment catches what instinct would otherwise silently merge, and the human-governed substrate provides corrective signal without requiring weight modification.

The architectural move is an extension rather than a replacement. Paper 1 draws the substrate/LLM division of labor at the governance boundary rather than the capability boundary: the substrate carries what humans govern — coordination, rationale, authority policies, conflict state — while the LLM carries high-dimensional reasoning humans cannot or should not encode directly. Paper 1 defends this commitment at cell scope and does not re-open it here. What Paper 2 Claim 1 adds is the scope extension. The governance boundary that cleanly splits substrate from LLM at the level of a single coordination cell generalizes to the level of an entire AI Self, where the full population of the Self’s instinct behavior and the full population of its deliberate reasoning are held as distinct architectural layers rather than blended into one model-as-system. Paper 1’s cross-claim spine — *outside-the-model* vs. *inside-the-model* — applies throughout Paper 2: coordination, governance, and reasoning live in the substrate outside the LLM, fast-pattern instinct lives inside the LLM, and the architecture’s value comes from the principled separation between them.

The System-1/System-2 framing this section uses is pedagogical scaffold, not theoretical grounding. Claim 1 does not depend on Kahneman’s dual-process theory being correct about human cognition, on Booch et al.’s SOFAI architecture being correct about machine cognition, or on any particular account of fast-vs-slow thinking being the right decomposition of reasoning in humans or machines. The framing imports warrant from an established way of thinking readers absorb quickly; the architectural work — what the two layers are, how they compose, how they evolve, how governance holds across them — is CKS’s own. This is the same move Paper 1 Claim 3 makes when it positions paraconsistent logic and Dung-style argumentation as lineage warrant without letting the defense of conflict preservation depend on either being architecturally load-bearing. Lineage supplies warrant; the architectural contribution is specific to CKS.

4.2 The named foil — conventional LLM-only architecture and its structural limits

The foil Claim 1 positions against is the monolithic “LLM-as-the-whole-system” pattern in which the model is the entire architecture and everything the system does — coordination, reasoning, governance, conflict handling, memory across sessions — lives inside the LLM’s inference. The 2024–2026 literature establishes a strong and growing consensus that this pattern has structural limits rather than merely capability limits, and the foil is well-supported.

The theoretical backbone is Mohsin et al. (2025), which argues that hallucination, long-context compression failure, reasoning degradation, retrieval fragility, and multimodal misalignment are intrinsic to the probabilistic-sequence-model paradigm trained with likelihood objectives over open-ended domains. The paper grounds these limits in diagonalization arguments, uncomputability of self-reference and program analysis, finite Kolmogorov complexity bounds, and PAC-style sample complexity, and — critically for Claim 1 — names them as architecture-class properties of probabilistic sequence models rather than engineering debt training can close. The recommended response is to treat LLMs as components in tool-augmented, retrieval-bounded, and solver-integrated systems. Quirke et al. (2025), a position paper, reaches a structurally similar conclusion: the single-frontier-LLM-as-entire-system interface is monolithic and structurally limited in capability, transparency, and safety. A cluster of reasoning-failure syntheses reinforces the picture. Song et al. (2026) surveys hundreds of studies and argues reasoning failures derive from the sequence-model-plus-likelihood objective, with chain-of-thought often improving final-answer accuracy while leaving step-wise reasoning unreliable. Apple’s 2025 “Illusion of Thinking” work shows reasoning models collapse at high complexity even when given explicit algorithms, and contemporaneous work on architectural limits in symbolic computation and reasoning documents correct rule statement coexisting with incorrect rule application inside the same model.

What matters for Claim 1 is that these are architecture-class findings rather than training-debt findings. If the foil’s limits were capability gaps closable by better training, the architectural response would be premature — a better LLM would absorb the reasoning layer Claim 1 commits to keeping separate. The foil literature denies that reading.

The limits are properties of the architecture class. The appropriate response is therefore architectural, and Claim 1 names one: separate the LLM from the reasoning layer as independently-evolving components under unified human governance, with the hybrid composition the response rather than a temporary scaffold awaiting an LLM capable of carrying both.

4.3 Closest-neighbor engagement — SOFAI and the dual-process AI lineage

The closest architectural neighbor is SOFAI (Booch et al. 2021), which foregrounds exactly the split Claim 1 commits to — an S1 solver aligned with fast pattern-based machine learning, an S2 solver for deliberate reasoning, and a metacognitive module arbitrating between them — with the specificity of a five-component design, navigation and epistemic-planning experiments, and a research program continuing through CACM 2025 and CP 2024 follow-ups. The shared commitments are substantial: the S1/S2 architectural split, the alignment of S1 with current ML/LLM-class systems in the paper’s own framing, the treatment of S2 as a deliberate-reasoning component distinct from S1, modular composition at the component level. SOFAI is the single most important architectural neighbor for Claim 1, and the honest framing is that Claim 1 belongs to the same broad design family.

Where the two works separate is on three axes the next subsection develops in full. The compact version here: SOFAI’s motivation is capability, not governance — the paper frames the split as a response to generalizability, adaptability, explainability, causal analysis, and common-sense reasoning gaps rather than a response to any desire to keep deliberate reasoning under human oversight. The metacognitive module is machine-driven: resource budgets, confidence scores, and expected-reward comparisons are computed at runtime from models rather than arbitrated by human-governed policy. The paper is explicit that “S2 never self-activates”, a phrase that locates decision logic in the metacognitive module rather than in S2 itself and carries architectural weight paraphrase cannot fully transmit — but that decision logic sits inside the machine, not in a human-governed substrate humans can inspect, modify, or override without retraining.

PlanSOFAI (Fabiano et al. 2024) instantiates the SOFAI architecture for classical planning, with neural components (a case-based planner and a transformer fine-tuned for planning) as S1 solvers and classical planners (Fast Downward, LPG) as S2 solvers, and the NuCLeaR workshop summary describes the metacognitive module as “a metacognitive process for governance” over S1/S2 selection. The lexical overlap with CKS’s use of *governance* is incidental: PlanSOFAI’s referent is machine-metacognitive arbitration of solver invocation at runtime, while CKS’s referent is human authority over substrate content, orchestration rules, and modification rights. Different referents, lexical collision only. Claim 1 does not treat PlanSOFAI as a separate architectural neighbor beyond this note — its commitments are SOFAI’s commitments with a neuro-symbolic instantiation, and the differentiation axes that hold for SOFAI hold for PlanSOFAI.

The broader 2024–2026 dual-process AI literature crowds the general vicinity without converging on Claim 1’s configuration. Ziabari et al.’s “Reasoning on a Spectrum” (2025) produces S1-aligned and S2-aligned variants of the same base LLM through DPO/SIMPO alignment; both reasoning modes live inside the model, with no external component. D-Mem (2026) places a fast and slow memory path in a dual-process framing, but both paths are machine-governed rather than human-governed. Bengio’s System-2 Deep Learning research program keeps S2 inside neural architectures through attention, working-memory analogues, and world models. A 2026 practitioner essay on “Products of System 2” positions Prolog-style logic programs as the stable residue of deliberate reasoning, which comes closest to CKS’s external-substrate move but stops short of formal governance commitment. “System Implications of Neuro-Symbolic-Probabilistic Architectures” (ACM 2025) is the closest peer-reviewed approach to independent-evolvability framing — it recommends independent versioning, monitoring, and deployment of a neural layer, a symbolic reasoning core, and a probabilistic layer — but treats this as systems-engineering practice rather than architectural requirement. The cognitive-architectures tradition (SOAR’s working memory; ACT-R’s declarative/procedural distinction; CLARION’s explicit/implicit knowledge separation; ICARUS’s knowledge-and-skill memory split) supplies decades of peer-reviewed dual-process structure that anticipates SOFAI’s component-level commitments, with declarative memory in ACT-R and explicit knowledge in CLARION as the closest cognitive-architectures analogues to a deliberate-reasoning substrate. The differentiation axes that hold for SOFAI hold for cognitive architectures with the same shape: the substrate is engineered configuration held during a cognitive run rather than a persistent human-governed substrate humans modify during operation, and independent-evolvability of the deliberate-reasoning component from the pattern-matching component is not formalized as architectural requirement. None of the works in this broader

literature puts reasoning capacity in a human-governed coordination substrate in Paper 1’s sense, and none formally commits to independent evolvability as a design axiom.

4.4 Three differentiation axes

Three architectural axes differentiate Claim 1 from the SOFAI-lineage neighbors and the broader dual-process AI literature. The axes are load-bearing individually and carry the differentiation jointly.

The first axis is the motivation for separating the LLM from an external reasoning component. Capability-motivated separations frame the split as a response to what the LLM cannot yet do well; the external component fills a capability gap and would in principle be absorbed if training could close the gap. Governance-motivated separations frame the split as a response to what humans should retain authority over regardless of LLM capability; the external component preserves authority architecturally rather than capability practically. The dominant architectural tradition is capability-motivated. SOFAI, PlanSOFAI, Logic-LM, IBM’s Logical Neural Networks toolkit, most neurosymbolic work, Bengio’s research program, Reasoning on a Spectrum, D-Mem, Dual-Process Scaffold Reasoning, Toolformer, and Re-Act all motivate the separation this way. Claim 1’s commitment is different in kind. The separation is motivated by the question of what humans should govern rather than by the question of what the LLM cannot do. Capability-motivated separations remain compatible with — indeed suggest — eventually training the LLM to absorb the external component; skill migration from S2 to S1 is an emergent behavior SOFAI explicitly names. Governance-motivated separations architecturally preserve human authority over substrate content regardless of LLM capability, which is the property that lets the two layers coexist safely as the LLM improves. Paper 1’s Claim 2 is the direct antecedent: the substrate/LLM division drawn at the governance boundary rather than the capability boundary. Claim 1 of Paper 2 extends that commitment from cell-level hybrid to Self-level instinct/reasoning separation.

Within the peer-reviewed literature, Constitutional AI (Bai et al. 2022) is the most prominent governance-motivated separation work and warrants direct engagement here rather than only at §8.5’s Claim 5 territory: it commits to safety-oriented separation through a training-time constitution as antecedent infrastructure for harmless behavior, with self-improvement supervised by the constitution rather than by per-decision human feedback. The architectural axis differentiating Claim 1 from Constitutional AI is the temporal locus of the substrate governance lives in. Constitutional AI bounds the substrate to training-time artifact — the constitution shapes weights through RLAI and is fixed at deployment. CKS commits the substrate to runtime governable artifact, with humans retaining authority to inspect, modify, and override substrate content during operation. Both are governance-motivated; they differ on whether governance lives at training time or at runtime. A small but load-bearing strand of 2024–2026 industry work moves in CKS’s runtime-governable direction — Epsilla/AgentStudio’s framing of memory as active governance system rather than storage, with policy that must not be dynamically modified by the agent; Graph Digital’s rule that “workers read doctrine, they do not change it”; Atlan’s treatment of governed metadata catalogs as accurate and auditable where unstructured vector memory is not; runtime AI governance platforms enforcing prompt firewalls and egress controls as explicit non-trust of LLM decisions. These sources articulate as industry practice what Claim 1 commits to as formal architectural commitment, with Constitutional AI standing as the peer-reviewed governance-motivated neighbor on the training-time-substrate side of the temporal axis.

The second axis is where reasoning capacity lives architecturally. If it lives in the LLM’s weights, it cannot be inspected, modified, or overridden without retraining. If it lives in an external symbolic system held fixed as engineered configuration, it cannot be modified by the humans who use the system during operation. Only if it lives in a human-governed coordination substrate — persistent, inspectable, modifiable, overridable, cross-session-addressable — can humans maintain architectural authority over the reasoning layer during operation. The 2024–2026 dual-process AI literature does not converge on this configuration. Bengio keeps S2 inside neural architectures. Reasoning on a Spectrum puts S2 inside an aligned LLM variant. D-Mem puts S2 in a second machine-governed memory path. SOFAI’s S2 solvers are external to the LLM but constructed as engineered configuration held fixed during operation rather than as live substrate. CKS’s commitment — S2 capacity living in a persistent, human-governed coordination substrate in Paper 1’s sense — is unoccupied in the surveyed literature.

The third axis is whether independent evolvability of the neural and non-neural layers is an incidental property of modular design or an architectural requirement. The two layers need to evolve on different timescales by different actors: the LLM through commercial or upstream upgrades chosen by vendors, the substrate through human-governed modification under organizational authority. Independent evolvability is the property that makes this divergence

safe. Upgrading the LLM must not require revalidating the substrate, and modifying the substrate must not require retraining the LLM. Only when the two layers are architecturally decoupled in this way can Claim 1’s separation hold as the LLM improves. The surveyed 2024–2026 literature treats this property as systems-engineering practice or as flexibility, not as architectural requirement. The ACM 2025 Neuro-Symbolic-Probabilistic Architectures paper is the closest approach — it explicitly recommends independent versioning, monitoring, and deployment, and recommends that ontologies and symbolic components be governed like software artifacts — but stops short of treating independent evolvability as a formal design axiom. SOFAI’s and PlanSOFAI’s modularity is framed as plug-and-play flexibility. Modular AI platforms (Unframe AI, Modular’s unified compute layer) describe modularity in componentwise terms rather than neural/non-neural specifically. No 2024–2026 source surveyed here states independent evolvability of neural and non-neural components as a hard architectural requirement. Claim 1 crystallizes as formal commitment what the recent literature recommends as practice — and the operational consequence of the difference is that revalidation discipline is treated as architectural property rather than as deployment hygiene: an LLM upgrade is non-conformant with Claim 1 if it requires revalidating the substrate, and a substrate modification is non-conformant if it requires retraining the LLM. Treating independent evolvability as practice leaves these revalidation paths as engineering choices the deploying organization can take or leave; treating it as architectural requirement makes their absence non-conformance with the design pattern. This is Paper 1 Claim 2’s governance boundary read forward into evolution time: if the substrate and the LLM are governed by different authorities on different timescales, they must evolve independently, and the architecture must formalize this or the governance boundary collapses.

4.5 Reviewer-objection responses

Two objections anticipate themselves, and the responses take distinct shapes.

The first objection — *this is just hybrid AI with extra vocabulary* — has surface plausibility. The capability-motivated hybrid tradition is extensive, well-published, and genuinely structurally similar to CKS at the neural/non-neural-component level. SOFAI’s lineage and the neurosymbolic literature supply the reviewer concrete precedents. The response has three independent components, each sufficient on its own, with the conjunction carrying the differentiation. The hybrid-AI tradition is almost entirely capability-motivated, and CKS’s commitment is governance-motivated; the small governance-motivated strand Claim 1 aligns with is recent, industry-voice, and distinct in motivation from the dominant capability-motivated tradition. Even within the capability-motivated tradition, the reasoning capacity does not live in a human-governed coordination substrate — SOFAI’s symbolic planners are engineered configuration, Bengio’s S2 sits inside neural networks, Reasoning on a Spectrum’s S2 lives inside an aligned LLM variant. The hybrid-AI tradition does not formally commit to independent evolvability as design requirement; CKS does. None of the three axes alone carries the differentiation, but the conjunction is specific and checkable. Paper 1 already defended the governance-boundary commitment against this exact objection at cell scope, distinguishing CKS from Cognitive Core (Seck 2026), MACI (Chang 2026), and control-plane work (Stevens, Sakura Sky 2025) at §4.4. A reviewer attacking Paper 2 Claim 1 at the “extra vocabulary” level is attacking Paper 1 Claim 2 and should be directed to Paper 1’s treatment; Paper 2 Claim 1 extends that already-defended commitment from cell to Self scope rather than reasserting it. The honest framing is that CKS joins a specific governance-motivated strand within a broader hybrid-AI landscape rather than that CKS is entirely novel.

The second objection — *System-1/System-2 analogies in AI are decades old* — is genuine. The dual-process framing has a long history in AI discourse, SOFAI instantiates the split architecturally from 2021 forward, Bengio’s System-2 Deep Learning research program predates SOFAI, and Kahneman’s own comments on deep learning as System-1-like are widely cited. The response does not contest the lineage. It makes the same move Paper 1 Claim 3 makes against paraconsistent logic at §5.1: the lineage supplies warrant from an established way of thinking, and the architectural contribution is the specific CKS configuration. The configuration has three conditions, none of which the lineage satisfies jointly. The LLM is specifically positioned as the System-1-analogue fast-pattern layer, which many dual-process AI works satisfy. System-2 is specifically a human-governed coordination substrate in Paper 1’s sense, which SOFAI, Bengio, Reasoning on a Spectrum, D-Mem, and the broader neurosymbolic lineage all fail. Independent evolvability of the two layers is an architectural requirement rather than incidental modularity, which the surveyed literature does not formally commit to anywhere. The conjunction is where the architectural contribution lives, not any single condition. The S1/S2 framing earns its place on warrant, not on load-bearing theoretical grounding; the architectural argument would hold with the scaffold replaced.

Three related objections surface on specific terms and warrant briefer treatment. “Your *governance* is what governed-RAG already does” collapses Claim 1’s commitment to independent evolvability of neural and non-neural layers with the broader category of policy-aware retrieval and citable outputs; the governance territory is shared, but Claim 1’s specific architectural move within that territory is distinct on the first and third axes above. “Why not just train a better LLM” collapses the reasoning-layer’s governance function with its capability content; scaling does not close the function, because the function exists to provide corrective signal without requiring weight modification, which is by construction what scaling cannot deliver; and Paper 1 §4 supports this argument by drawing the substrate/LLM division at the governance boundary rather than the capability boundary. “PlanSOFAI’s metacognitive module already does governance” is the lexical-overlap misreading addressed in §4.3.

4.6 Paper 1 inheritance and downstream placement

Claim 1 of Paper 2 inherits from Paper 1 (Li, 2026) without redefinition: the six architectural commitments (substrate as human-governed coordination artifact, schema-level decomposition at the governance boundary, conflict preservation as first-class substrate state, linear-cost storage and composition, tool-agnostic instantiation, AI-as-substrate-mediator), the authority-vs-labor distinction that governance is an authority architecture rather than a review workflow, and the outside-the-model/inside-the-model cross-claim spine that applies throughout Paper 2. The specific inheritance load-bearing for Claim 1 is Paper 1’s Claim 2 commitment: the substrate/LLM division drawn at the governance boundary rather than the capability boundary. Claim 1 extends that commitment from the cell-level hybrid Paper 1 defends to the Self-level instinct/reasoning separation Paper 2 requires. No new theoretical commitment is added at this step; the extension is a scope move, not a foundation move. Paper 1’s engagements with OIDA (Bottino et al. 2026) at §5.2 and Knowledge Objects (Zahn & Chana 2026) at §6.2 — adjacent substrate-based architectures that do not make the System-1/System-2 framing or the instinct/reasoning separation Claim 1 defends — carry forward without modification, and Claim 1 references those treatments rather than redoing them.

The honest framing of Claim 1 as scope move warrants a brief contribution-defense note. Claim 1 is Paper 2’s first claim despite extending an already-defended Paper 1 commitment because every subsequent claim in Paper 2 depends on the Self-scope framing Claim 1 establishes. Claim 2’s structural machinery (cells composing into aspects into Selves; DNA/action layer distinction; expression as governed selection) operates on the Self the instinct/reasoning separation makes architecturally available; Claim 4’s three evolution mechanisms (instinct evolution, DNA evolution, action-feedback evolution) operate on the two separated layers; Claim 5’s multi-shaped governance operates per-mechanism over the separated layers; Claim 6’s enterprise brain Self extends the per-Self separation to organizational scope. Claim 1 is foundational not because it adds new theoretical content but because the Self-scope framing is what subsequent claims operate over. A reviewer treating Claim 1 as scaffolding rather than as load-bearing first claim is treating the scope-move nature of the extension as if it disqualifies the claim from foundational status; the extension is what makes the rest of Paper 2 well-formed.

The actor-neutrality that holds throughout Paper 1 holds here as well. Human-governed does not mean human-authored. LLMs can draft, populate, and update substrate content under human governance; what humans own is authority over substrate structure, orchestration rules, and modification rights. Claim 1’s commitment to the reasoning layer as human-governed substrate is a commitment about where authority sits, not about who produces every substrate element.

Claim 1 establishes the foundation. Downstream work in Paper 2 depends on the separation holding. As §5 develops, the instinct/reasoning separation is supported by structural levels that compose cells into larger coordinated wholes, with distinct layers inside each cell that the separation requires. As §7 develops, the two layers evolve on different timescales through distinct mechanisms, and that independence of evolution is what the third differentiation axis above secures as architectural commitment. As §8 develops, human governance takes different shapes across evolution mechanisms, operating over what Claim 1 commits to keeping separate. As §9 develops, the separation supports an enterprise deployment pattern in which powerful LLM capability and human-governed coordination coexist at organizational scope. Claim 1 makes the separation available; the rest of Paper 2 uses it.

5. Claim 2 — CKS Selves as architecturally modular by commitment rather than by evolution

5.1 Claim statement and the foundational-modularity framing

Claim 2 states Paper 2’s structural commitment: a CKS-governed AI Self is architecturally modular by commitment rather than by evolution, with modularity required at every level of composition rather than treated as an emergent property the architecture has to wait for. Three components carry the commitment. Cells compose into aspects, aspects compose into Selves, with structural roles relational and purpose-defined rather than intrinsic — the same underlying CKS artifact can participate as a cell in one arrangement, as part of an aspect in another, or as a component of different aspects simultaneously, depending on what the structure is for. Within every cell, two distinguishable substrate layers operate: DNA as the stabilized orchestration content that defines what the cell is governed to do, and action as the recorded experience of what the cell has actually done, with distinct governance semantics holding between them. Expression is the per-deployment design choice over which DNA-layer substrates activate for a given cell goal, governed by orchestration substrate that is itself substrate content under Paper 1 (Li, 2026) commitments. Across all three components, modularity is architecturally required at instantiation, not measured after the fact.

Claim 2 does not re-argue Claim 1’s instinct/reasoning separation; it establishes the structural machinery that separation depends on. Without multi-level composition with relational roles, the Self scope at which Claim 1 operates cannot hold together as architectural composition under unified governance. Without the DNA/action layer distinction, the substrate cannot carry what Claim 1 commits to it carrying — governed reasoning content distinct from recorded execution state. Without expression, the substrate cannot be selectively activated per cell goal. Each component is a structure Claim 1’s separation requires; together they make the separation architecturally available.

One disclaimer about the term *Self* belongs in the opening rather than at first architectural development. CKS commits *Self* to architectural composition under unified human governance — the integrated whole that holds multiple aspects as facets of one CKS-governed intelligence — not to consciousness, subjectivity, psychological selfhood, or any extension of autopoiesis to minimal selfhood or sensorimotor agency. The biological *Self* in autopoiesis and biological-individuality discussions is organizationally closed without being conscious; microbes and plants are autopoietic individuals without being conscious in the ordinary sense. Reader connotations around *Self* drift toward consciousness easily; the bound is placed preventively here so the architectural development that follows can use *Self* freely as architectural commitment.

5.2 Multi-level composition with relational roles

Three architectural levels structure the Self: cell, aspect, Self. The cell is the atomic unit Paper 1 defends — a CKS artifact with substrates, orchestration rules, and the six architectural commitments Paper 1 establishes. All Paper 1 commitments hold at the cell level without modification. An aspect is a coordination arrangement of cells serving a particular purpose, operating over its constituent cells as content domain. Multiple aspects can coexist within one Self, potentially sharing cells across arrangements; a person’s competitive sports mode and calm study mode are distinct aspects drawing on the same underlying capacities, coexisting as facets of one person. The Self is the integrated whole — not a bigger aspect, not a rigid top-of-hierarchy, and not one end of a strict access hierarchy, since higher levels operate over lower levels as content domain but access patterns are not strictly hierarchical and the Self can access cells directly when purpose requires. The Self contains the multiple coexisting structural perspectives and their collective intelligence as one unified whole.

Two architectural commitments shape how these three levels relate to one another. The first is governance-boundary inheritance at every level. Paper 1 §4 draws the substrate/LLM division at the governance boundary at the cell level: substrate handles coordination and governance, LLM handles high-dimensional reasoning humans cannot or should not encode, with the division at what humans should govern rather than at what the LLM cannot do. Claim 2 commits that this same division holds at every composition level. Cells have substrate/LLM at the governance boundary as Paper 1 specifies; aspects coordinate cells through substrate-mediated composition under human governance, with the orchestration of cells into a purpose-defined arrangement itself substrate content; Selves hold multiple aspects under unified human governance with substrate carrying the composition relationships across levels. Microservices+DDD, hierarchical state machines, and contemporary multi-level architecture commit to layered governance, but the spe-

cific content of governance at each layer is left to the implementation; what they do not commit to is recursive inheritance of the same substrate/LLM division at every composition level. CKS's commitment is that the same architectural move — Paper 1 §4 generalized — holds at every level.

The second commitment is relational structural roles. Existing modular-software literature commits to intrinsic structural membership: a service belongs to one bounded context, an aggregate to one service, a component to one composite. CKS's commitment is the inverse — the same underlying CKS artifact participates as a cell in one arrangement, as part of an aspect in another, or as a component of different aspects simultaneously, depending on what the structure is for. The artifact's identity is the bearer; the role-instances inhere in the bearer per the structural arrangement that calls on them; governance properties propagate per arrangement rather than per unit.

This commitment is honestly framed as a synthesis of three closest-adjacent strands rather than as territory in a vacuum. Object-capability systems (Miller's *Robust Composition*; the 2017 ECOOP capability-based module system paper) carry one underlying object referenced through multiple capabilities, each conferring distinct rights; the capability-facet pattern — wrap an underlying object with a narrower interface — is the architectural mechanism for "same underlying state, different authority per arrangement." Formal-ontology work on roles and qua-entities (Steimann 2000; BFO; OntoUML's role stereotype) carries roles as anti-rigid externally-dependent entities, with multiple roles inhering in the same bearer simultaneously and each tied to its own external context. The 2026 *Architecting Agentive Communities* paper carries community specifications of roles plus deontic governance tokens (burden, permit, embargo) attached to roles within community patterns rather than to agents globally — the same agent fills different roles in different communities, each role carrying its own obligations and permissions per arrangement. CKS's relational-roles commitment is the synthesis of these three: same underlying unit of identity, multiple structural role-instances per arrangement, governance attached per arrangement, in an executable substrate at multi-level composition scope.

A vocabulary collision deserves a brief note here. The term *aspect* in CKS is not aspect in the aspect-oriented programming sense (Kiczales et al. AspectJ), where aspects are cross-cutting concerns weaving into base code through pointcut-advice composition. CKS aspects are coordination arrangements of cells serving a particular purpose, drawing on the same underlying capacities, coexisting as facets of one Self — not cross-cutting concerns over base cell code.

One contrast belongs here briefly as well. Biology evolved modular cell architectures over deep time because modularity supports evolvability (Hartwell et al. 1999; Kirschner and Gerhart 2005); designed systems can adopt modularity from the start because modular decomposition is part of the engineering design repertoire (Vincenti 1990). CKS is modular by architectural commitment from the start at every level — not a property the architecture hopes for, evolves toward, or measures after the fact, but a design requirement at instantiation time. The architectural novelty is not "modular by commitment" considered alone — designed systems have always been modular by commitment from the start; that is what software engineering is — but the specific configuration in which multi-level composition with relational role membership, the DNA/action layer distinction within every cell, and expression as governed selection over which DNA-layer substrates activate per goal compose into one architectural commitment. The biology contrast frames the conditions under which CKS can commit to the configuration upfront (engineered design under human governance) rather than the configuration itself, which is the novelty Claim 2 defends.

5.3 DNA/action layer distinction within every cell

Within every cell, two distinguishable substrate layers operate. The DNA layer carries stabilized orchestration content — harness logic, conflict-handling rules, lifecycle policies, schemas — that defines what the cell is governed to do. The action layer carries recorded task instances and their outputs — what actually happened when the DNA met an actual task — that accumulates as the cell's lived experience. Both layers are substrate content under Paper 1's commitments. What distinguishes them is governance semantics: DNA is human-governed at origin and modification, while action is produced by the cell's operation, with one-way authority — DNA constrains and shapes what action can be, and action does not autonomously rewrite DNA. The loop closing action evidence back into DNA refinement is governed and human-mediated; as §7 develops, that loop is one of three evolution mechanisms operating on the Self's internal layers.

The architectural shape — governed policy on one side, operational state on the other, one-way authority from policy

to state — is recognized in two contemporary literatures Claim 2’s commitment engages directly. The first is a 2024–2026 industry/practitioner cluster of agent-memory architectures that has converged on this distinction explicitly. Epsilla AgentStudio’s framing crystallizes the pattern: “Memory is Governance, not Storage.” The architectural rule is a three-way separation of state, memory, and policy in which policy is the system’s external specification and is not modified by the agent’s memory operations — if memory could rewrite the agent’s permission rules, that would not be memory evolving but privilege escalation. Graph Digital’s four-typed-layer model carries the same shape, with worker agents reading doctrine but unable to write to it; doctrine modification flows through governed update processes rather than worker writes. Atlan frames the distinction as a metadata control plane (governed catalog as substrate of truth, vector storage as derivative); Unstructured.io defines governance and observability as a cross-cutting layer separate from cognition, control flow, memory, and tool execution. These sources are industry/practitioner voice rather than peer-reviewed architectural foundations, but together they constitute the strongest emerging-pattern literature for the DNA/action distinction.

The second literature is configuration management (Burgess on theory; Morris on Infrastructure as Code; Puppet, Ansible, Chef, Kubernetes in practice), which commits across the cluster to declarative policy defining desired state with runtime state reconciled toward the policy through controllers, the boundary enforced one-way and drift treated as a bug rather than a new configuration to adopt. Modern policy-as-code (Open Policy Agent, Conftest, Checkov) adds constraints over configuration as gates on changes to desired state.

The differentiation territory is on what kind of policy and what kind of state, within Paper 1’s specific substrate framework. The agent-memory cluster’s “doctrine” or “policy” varies across sources — skills and schemas at Graph Digital, permission boundaries at Epsilla, metadata and lineage at Atlan; configuration management’s policy is infrastructure desired-state. CKS’s DNA is orchestration substrate carrying Paper 1’s specific commitments — substrate-based, human-governed, conflict-preserving, AI-mediated, tool-agnostic, linear-cost — not policy or doctrine in some general sense but substrate with Paper 1’s properties. CKS’s action is recorded experience with its own substrate semantics — not workflow state to be reconciled but substrate content with the Paper 1 commitments holding at the layer. The architectural shape (policy and state with one-way authority) is shared with the neighbors; what CKS articulates is the specific commitment to two layers within a single cell, both substrate content under Paper 1’s framework, with distinct governance semantics. The differentiation depends on Paper 1’s substrate framework being accepted; Paper 2 does not redefine the framework here but cross-references Paper 1’s defense of the six commitments and treats DNA and action as the two layers within Paper 1’s substrate at cell scope. Claim 2 articulates as formal architectural commitment within Paper 1’s substrate framework what the agent-memory cluster treats as emerging best practice.

Two further bounds belong with the term *DNA*. The fully-determining-genome popular framing — DNA as master program that fully specifies cell behavior — is not what CKS commits to. Modern genomics has substantially revised that picture (alternative splicing; regulatory and catalytic RNA roles; epigenetic marks regulating expression without altering DNA sequence); DNA in contemporary biology is one substrate among others. CKS’s DNA is similarly *one stabilized substrate layer* alongside action, not a complete program that determines cell behavior. The actor question takes a parallel bound: DNA is human-governed, not human-authored — humans hold authority over DNA structure, orchestration rules, and modification rights, while LLMs may draft, populate, and update DNA content under that authority, with Paper 1 §3.3’s authority-vs-labor distinction holding throughout.

5.4 Expression as governed selection

Expression is the mechanism that determines which DNA-layer substrates activate for a given cell goal. Each cell carries a harness substrate — itself human-governed, fully inspectable, modifiable, and overridable — that selects which sub-substrates are active for current activity. Cells can carry the full Self’s DNA with selective expression (when lineage and reconstitution matter, as in regulated work) or partial slices (when storage and cognitive load matter more). This is a per-deployment design choice governed by orchestration substrate that is itself substrate content under Paper 1 commitments. Cells with shared underlying DNA differ in active behavior because what is expressed differs by deployment, not because the DNA differs.

The closest software-architecture neighbor is feature flags (Hodgson on Martin Fowler’s site). Feature flags commit to deployment-time activation as a deliberate architectural pattern — code is deployed with new features hidden behind toggles; the toggle is flipped to expose the feature to some or all users. Hodgson’s four-category taxonomy maps to different governance treatments (release, experiment, ops, and permissioning toggles). The architectural

shape — code plus flag wiring define potential behavior space; governed flag configuration selects which behavior is active per deployment; runtime state is processed by whatever subset is currently active — is shared with CKS expression. Three differences hold. Feature flags are typically capability-motivated (deployment safety, gradual roll-out, A/B testing) rather than governance-motivated; the closest governance treatment in the feature-flag literature is deployment-safety governance, not the substrate-content governance Claim 2 commits to. Feature flags activate code paths; CKS expression activates governed orchestration substrate content within Paper 1’s framework. Feature-flag configurations are typically per-binary; CKS expression is per-deployment design choice over what subset of governed substrate content the cell carries and which subset of carried content is active. Configuration management (IaC, Puppet, Kubernetes) shares the policy-versus-state shape from the previous section’s engagement; what differs at expression is selecting active substrates per deployment rather than reconciling runtime state toward policy. CKS expression is governed selection under orchestration substrate authority; orchestration substrate is itself substrate content under Paper 1 commitments, so what selects what to express is human-governed in the same way DNA is.

Biology vivifies the architectural shape. Davidson and Levine’s gene regulatory network work treats the genome as a static catalog of genes and cis-regulatory modules, with transcription factors as context signals: “Gene regulatory networks explicitly represent the causality of developmental processes.” Different cell types differ in regulatory state, not genomic content. The architectural picture — selective activation of available content per context — transfers cleanly to what CKS expression commits to. One bound is needed here: biological gene expression is mechanistically determined by molecular dynamics that arose through evolution; CKS’s expression is per-deployment design choice governed by orchestration substrate. The selective-activation shape is shared; the governance commitment is CKS’s. Expression governance, like DNA governance, is human-governed without being human-authored — orchestration rules can be LLM-drafted under human governance, and what humans own is authority over what expression mechanisms exist and how activation decisions are made.

5.5 Reviewer-objection responses

Three objections warrant full responses; three more take briefer forms.

Software has been modular for decades; what is new. Two architectural moves §5.2 developed. Governance-boundary inheritance at every level generalizes Paper 1 §4’s substrate/LLM division across composition levels; existing literature commits to layered governance but leaves the content of governance at each layer to the implementation. Relational structural roles is the synthesis of three closest-adjacent strands (object-capability facets, qua-entity formalism, deontic-token agent communities) into an executable substrate at multi-level composition scope; existing modular-software literature commits to intrinsic structural membership and CKS does not. The synthesis-as-novelty framing is honest about composition but understates the architectural work: each of the three synthesized strands operates in a different territory (object-capability for memory-safety and security; qua-entity formalism for ontological role-modeling; deontic-token agent communities for distributed agent obligations), and the architectural contribution is neither the strands individually nor their concatenation but the configuration in which all three operate together over a single human-governed substrate at multi-level composition scope under Paper 1’s commitments — a configuration none of the three strands occupies alone and none of the standard combinations (object-capability with role-modeling; deontic tokens with object-capability) pursues at the level of executable substrate. Neither move alone carries the differentiation; the conjunction is specific and checkable.

DNA/action is just RAG with a taxonomy, or policy-versus-state distinctions are not new. Concede the shared territory and locate the contribution within it. §5.3 developed the differentiation: CKS’s DNA is orchestration substrate carrying Paper 1’s specific commitments — not policy in some general sense but substrate with Paper 1’s properties — and CKS’s action is recorded experience with its own substrate semantics, not workflow state to be reconciled. CKS articulates as formal architectural commitment within Paper 1’s framework what the agent-memory cluster treats as emerging best practice.

Paper 2 is just scaling Paper 1 modularity up; what is new. Three architectural moves Paper 1 does not make. Multi-level composition with relational roles is between-unit modularity Paper 1 does not address; the DNA/action distinction is within-unit modularity Paper 1 does not commit to; expression as per-deployment design choice over which DNA-layer substrates activate is a mechanism Paper 1 does not introduce. Paper 2 defines additional modularity — between-unit, within-unit, and per-deployment — that Paper 1 leaves undefined.

Three briefer responses. *Three levels is arbitrary* misses that cell, aspect, and Self are architecturally driven scopes — cell is what Paper 1 defends, Self is where Claim 1’s separation operates, aspect is the intermediate scope at which coordination arrangements of cells serve particular purposes — rather than positions chosen for hierarchical-depth aesthetics. *Biology is decorative metaphor* misses the per-term work woven through the development above bounding what each term carries architecturally and what connotations it does not adopt. *This is just object-capability with a biology coat of paint* misses the honest synthesis-of-three-strands framing for relational roles and the independent Claim 2 commitments — governance-motivation, multi-level structural composition with named levels, DNA/action, expression — that compose with the relational-roles commitment.

5.6 Paper 1 inheritance and downstream placement

Claim 2 inherits all six Paper 1 architectural commitments at every level of the three-level structure: substrate as human-governed coordination artifact (Paper 1 §3); governance boundary and schema-level decomposition (Paper 1 §4) — the most load-bearing inheritance, with Paper 2 Claim 2 generalizing the substrate/LLM division across composition levels; conflict preservation as first-class substrate state (Paper 1 §5) preserved at every level; linear-cost storage and composition (Paper 1 §6) underwriting multi-level composition tractability; tool-agnostic pattern (Paper 1 §7); AI-as-substrate-mediator (Paper 1 §4.2). Conflict preservation at every level is architecturally substantive — aspects carry their own coordination conflicts, Selves carry conflicts across aspects, and Paper 1’s two-level conflict-handling structure (substrate-level preservation, cell-level resolution by orchestration rules) extends naturally to multi-level scope. The actor-neutrality from Paper 1 §3.3 holds throughout: human-governed does not mean human-authored; what humans own is authority over substrate structure, orchestration rules, and modification rights, not authorship of every element.

Claim 2 establishes the structural machinery downstream claims operate on. As §6 develops, lifecycle operations (birth, mating, death) apply at every level of this structural machinery. As §7 develops, evolution operates on the modularity Claim 2 commits to through three mechanisms in productive tension, with bidirectional axes (horizontal and vertical) holding within the structure this claim defends. As §8 develops, human governance takes different shapes across evolution mechanisms operating over the substrate content this claim commits to. As §9 develops, the structural machinery extends into the enterprise brain Self under the wholeness and topology-consequence commitments — cells composing through aspects into one functioning unit under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability. Claim 2 makes the modularity available; the rest of Paper 2 uses it.

6. Claim 3 — Lifecycle operations as governed primitives at every level of CKS Selves

6.1 Claim statement and the lifecycle-as-governed-primitive framing

Claim 3 states Paper 2’s lifecycle commitment: a CKS-governed AI Self treats birth, mating, and death as governed primitives that apply uniformly at every level of the structural machinery Claim 2 establishes. Birth is the creation of new cells, aspects, or Selves under human governance, with origination performed by humans directly or by LLMs operating under human direction. Mating combines parental content across both the DNA and action layers under orchestration substrate governance, with three pattern variants — union, selective merge, and lineage-preserved union — configuring how that combination is performed. Death has two categories with distinct architectural results: functional obsolescence is genuine deletion of a cell whose function is no longer needed, releasing its substrate resources; lineage supersession is retirement-with-archival of parent cells after a superior offspring emerges, with archived substrate content remaining substrate-addressable.

Across all three operations, the architecture commits to governance over the lifecycle decision and the type-appropriate architectural result, not the labor of executing the decision and not autonomous lifecycle dynamics. This positioning matters: lifecycle vocabulary in biology connotes processes that organisms undergo, and the easy reading of *birth*, *mating*, *death* in an AI architecture would be that the architecture itself is what makes these things

happen. Claim 3 commits something different. The lifecycle primitives are first-class as architectural property, and humans hold authority over each lifecycle decision; the labor of carrying out the decision is available to humans directly or to LLMs operating under human direction, and the substrate makes the type-appropriate result legible at every level of composition. Claim 3 does not re-argue Claim 1’s instinct/reasoning separation or Claim 2’s three-level structure; it operates lifecycle primitives at every level and across both layers of the architecture Claims 1 and 2 establish, and positions the primitives as the substrate downstream evolution and governance claims operate over.

6.2 Birth as governed origination

Birth’s architectural commitment is governance over creation, not labor of creation. A new cell, aspect, or Self comes into existence when humans authorize its creation under their governance authority; origination of the cell’s substrate content may be performed by humans directly or by LLMs operating under human direction, with humans holding authority over the orchestration rules under which LLMs operate. This commitment inherits directly from Paper 1 §3.3’s authority-vs-labor distinction: governance is an authority architecture, not a review workflow. The cell’s existence is governed at the moment of creation; its substrate content may be LLM-drafted, human-authored, or composed by both under whatever orchestration rules humans put in place.

The closest-neighbor literature in AI agent frameworks and software constructors treats creation as the simplest lifecycle primitive but leaves the actor question largely vague. Configuration-driven agent instantiation (AutoGen, LangGraph, the Microsoft Agent Framework, the Claude Agent SDK), UI-driven creation (OpenAI Agent Builder, StackAI, CrewAI AMP), A2A “agent cards” describing capability via skills — across this cluster, creation is configuration-plus-instantiation, with framings like “developer creates,” “operator instantiates,” or “factory produces” doing the actor work without distinguishing authority over creation from labor of creation. Software constructors (OOP allocate-initialize-register, Kubernetes pod create-and-schedule, React mount, Spring bean creation) follow the same pattern at the engineering-primitive level. None of these works formalizes the governance-vs-labor distinction Paper 1 §3.3 articulates, and none treats creation as a recursive primitive uniform across multiple structural levels under unified governance authority. Birth in CKS imports the engineering-primitive understanding and adds Paper 1’s authority commitment: creation is governed at every level, whether the substrate content is LLM-drafted or human-authored.

6.3 Mating as one governable primitive over content combination

Biology is the closest domain that singles out two-parent combination as a fundamental lifecycle primitive distinct from creation. Syngamy — gamete fusion into a zygote — is named-primitive in cell and organism lifecycle vocabulary; cell fusion in myoblasts and osteoclasts and mitochondrial fusion are named primitives at developmental and subcellular levels. Within AI specifically, evolutionary computation has treated crossover as fundamental primitive distinct from mutation since Holland 1975 (genetic algorithms), Goldberg 1989, and Koza’s genetic programming, and the broader EC tradition operates crossover as a load-bearing operator across CMA-ES, evolution strategies, and quality-diversity algorithms. CKS mating is closer to GP crossover than to either software-lifecycle merge logic or biological syngamy, in that all three (GP crossover, biological mating, CKS mating) operate as architectural primitive over heritable content; §7.5 develops the differentiation against EC at the mechanism-count and architectural-separation level. Software lifecycles and AI agent frameworks at the deployment-architecture level do not have a combination primitive at all — combination of parental content appears as application-level merge logic (Git two-parent merge, software product line feature derivation, schema migrations) bolted onto version-control infrastructure. CKS imports the combination primitive from the lineage where it has architectural status (biology and EC) into AI Self architecture under governance.

Mating in CKS combines parental content across both the DNA and action layers Claim 2 establishes within every cell, with the specific combination configured by orchestration substrate. Three pattern variants are available, each with distinct architectural commitments. The union pattern keeps everything from both parents and preserves all conflicts as first-class substrate state in the offspring; conflicts at merge time become substrate content the offspring carries forward, with cell-level orchestration rules in the offspring governing how cells handle them at execution. This inherits Paper 1 §5 conflict preservation directly: where Paper 1 establishes that contradictions encountered during AI-mediated coordination are preserved in the substrate as first-class addressable objects with their own identity and provenance, the union pattern extends that commitment to merge-time conflicts arising when two parents are united. Git two-parent merge with conflict markers is the closest computational neighbor, but Git markers are transient —

required to be resolved for the merge to complete, after which they disappear. The CKS union-pattern commitment is the inverse: merge conflicts become persistent first-class substrate state per Paper 1 §5, addressable across sessions in the offspring’s substrate, with provenance attached to the conflict and deferral as a terminal recorded outcome.

The selective-merge pattern minimizes the offspring’s content footprint with maximum curation: humans or LLMs operating under human direction pre-curate what crosses the boundary from each parent into the offspring, with humans holding authority over the curation rules themselves. The closest neighbor is software product-line engineering’s selective composition (Pohl, Böckle, and van der Linden; Czarnecki and Eisenecker), where feature models declare valid product configurations and product derivation selects subsets subject to compatibility constraints. The differentiation is on the actor question: SPL selective composition is constraint-driven, with curation framed as engineering work; CKS selective merge is governance-driven, with the curation rules themselves substrate content under human authority. The lineage-preserved-union pattern maximizes provenance: the offspring carries pointers to parent cells so every element of its substrate has traceable derivation back to its source. The closest neighbors are W3C PROV and Aethon’s reference-based replication, where structures descend through specialization and branching with the reference graph encoding lineage. PROV implements lineage as annotation over data; CKS lineage-preserved union and Aethon both commit to lineage as substrate content — the lineage pointers are themselves in the offspring’s substrate, not in a separate provenance graph layered above it.

The mating mechanism’s distinctive architectural work, beyond the three pattern variants, is collapsing what biology distributes across multiple analytically-distinguishable transmission mechanisms — genetic, epigenetic, behavioral, and cultural inheritance, treated by the Extended Evolutionary Synthesis position (Jablonka and Lamb; Laland and Uller) as four channels of one multi-layer inheritance architecture, each with its own mechanisms and fidelity profile but causally entangled rather than fully independent — into one governable primitive. CKS does not commit to genetic-style high-fidelity copying or cultural-style low-fidelity imitation as separate mechanisms but to one configurable combination primitive that biology’s distinctions become special cases of under orchestration substrate governance.

Cross-lineage mating compatibility is where CKS differs from biology on the mating axis. Biology’s general case in canonical multicellular eukaryotes, codified in Mayr’s Biological Species Concept and developed in Coyne and Orr’s *Speciation*, is reproductive isolation through layered prezygotic and postzygotic mechanisms making cross-species recombination rare and often maladaptive. The narrowing-contrast cases — horizontal gene transfer pervasive in prokaryotes; symbiogenesis as rare, transformative once-or-a-few-times events; hybridization as widely available where reproductive barriers are incomplete but rate-limited and lineage-biased toward plants — narrow the contrast but do not invalidate it. The exceptions require specific mechanisms, are mostly lineage-specific, and are rate-limited per generation. CKS commits to governance-configured cross-lineage combinations available across all cell types as architectural commitment, not as exception. One bound on the term: CKS mating is the broader combination primitive of which sexual-reproduction-style two-parent inheritance is one governance-configured pattern; the architecture does not commit to gametes, fertilization, meiosis, or two-parent-only configurations, and whether a mating instance is two-parent biological-style, multi-parent, or self-derived from one parent’s variants is configured by orchestration substrate.

6.4 Death as governed retirement with two distinct architectural results

Biology distinguishes two cell-level retirement strategies whose architectural shapes are clean. Apoptosis (Kerr, Wyllie, and Currie 1972) is genetically-controlled programmed cell death — caspase signaling triggers physical dismantling, with cell and internal state genuinely destroyed. Dormancy and seed-bank preservation is a reversible inactive state where metabolic activity is reduced but the cell remains viable, addressable, and reactivatable. The two are distinct life-history strategies, not slow-vs-fast variants of one primitive — apoptosis couples unaddressability with destruction, dormancy decouples participation from addressability.

Death in CKS is the lifecycle primitive that contains both shapes as architectural commitments to two distinct types under one governed retirement decision. The general statement that death in CKS is operational retirement rather than deletion of state applies cleanly to lineage supersession; functional obsolescence is genuine deletion. Conflating the types produces confused understanding of what the architecture commits to, and the two type-appropriate architectural results are the load-bearing distinction the rest of this subsection develops.

Functional obsolescence is the first death type. A cell whose function is no longer needed is deleted under a governed retirement decision; its substrate content is removed and the resources its substrate occupied are released. The closest computational neighbor is garbage collection (mark-and-sweep, generational GC, reference counting): unreachable objects are reclaimed, memory is reused, the system loses object state unless separately persisted. The closest biological neighbor is apoptosis. The deletion-to-free-resources shape is consistent across both: the architectural commitment is to deletion. CKS's commitment for functional obsolescence is the same; what CKS adds is the governance commitment over the deletion decision, with the per-type governance machinery developed by §8.

Lineage supersession is the second death type. After a superior offspring emerges from mating or evolution, the parent cells are retired but archived for audit and historical reference; archived substrate content remains substrate-addressable at the same identifier, with the system designed to resolve the identifier into the prior state. The closest computational neighbors share the architectural commitment to decoupling addressability from active participation: object-capability revocation through membrane and caretaker patterns (Mark Miller's *Robust Composition*; Van Cutsem and Miller on trustworthy proxies), where revocation removes or attenuates authority without requiring destruction of the target object; software deprecation and archival policies (PyPI project archival; GitHub repository archival as read-only with full history accessible); and database snapshot semantics with point-in-time recovery. The closest biological neighbor is dormancy and seed-bank preservation. The strongest biological-architectural framing is the Major Transitions in Evolution framework (Maynard Smith and Szathmáry 1995): "entities can replicate only as part of a larger unit after it". Ancestral entities — mitochondria and plastids in eukaryotic cells, cells in multicellular organisms, castes in insect societies — persist as components of new individuals while losing autonomous evolutionary individuality. The picture is substrate incorporation plus suppression of autonomous replication, not deletion.

What both death types share is the governance commitment over the retirement decision: death in CKS is a governed retirement decision with a type-appropriate architectural result, not an autonomous outcome of cell metrics or system dynamics. What they differ on is the result: deletion in functional obsolescence; archival with continued substrate-addressability in lineage supersession. The systems literature treats deletion and archival as distinct architectural commitments rather than parameterizations of one primitive — software engineering distinguishes EOL of support from EOL of addressability; capability literature treats revocation as separate from destruction; biology treats apoptosis and dormancy as distinct life-history strategies. CKS commits to both patterns simultaneously under one lifecycle primitive with two type-specific architectural results.

Reversibility for lineage supersession is where CKS differs from biology on the death axis. Biology's general case is irreversibility — extinction is forever in any meaningful evolutionary timeframe — and the moral and legal urgency of conservation rests partly on that irreversibility. The narrowing-contrast cases narrow but do not invalidate the contrast. Resurrection ecology (Daphnia hatched from sediment egg banks; microbes revived from permafrost) re-activates dormant cohorts of still-extant species using built-in biological snapshot mechanisms, not extinct species; Lazarus taxa is epistemic reversibility; de-extinction efforts produce functional analogs through CRISPR editing of living relatives' genomes, with the literature converging on the framing that "de-extinction may be more about creating valuable approximations than creating Jurassic Parks" and IUCN guidance preferring "ecological replacements" or "proxies." CKS's commitment for lineage supersession is reversibility-by-architectural-commitment: archived parent cells remain substrate-addressable at the same identifier, with the system designed to resolve the identifier into the prior state through PITR-style restoration, snapshot mounts, or capability reactivation as primary operations. The architectural-property contrast: biological "addressable" means physically present and locatable but does not guarantee viability, completeness, or identity preservation; computational "addressable" means logically referenceable via identifiers the system is designed to resolve into exact prior states given retained substrate. The reversibility framing applies to lineage supersession only; it does not extend to functional obsolescence.

6.5 Reviewer-objection responses

Three objections warrant full responses; five more take briefer forms.

Why mating at all? Why not just versioning? CKS mating is not committed to biology's two-parent genetic recombination but is a governance primitive over content combination across the DNA and action layers, with three pattern variants — union, selective merge, lineage-preserved union — each configurable for two-parent, multi-parent, or self-derived combination. Versioning gives one form of lineage tracking but does not give a combination primitive at all. Software lifecycle vocabulary and AI agent frameworks treat two-parent combination only as application-level

merge logic; biology is the only domain that commits to combination-of-two-parents as fundamental primitive, and CKS imports it under governance configuration.

Reversibility-as-where-CKS-differs-from-biology overclaims. The reversibility framing applies to lineage supersession, where archived parent cells remain substrate-addressable and reactivatable as architectural property; it does not extend to functional obsolescence. For lineage supersession, biology’s general case is irreversibility, and the narrowing-contrast cases (resurrection ecology reactivating dormant cohorts of still-extant species; Lazarus taxa as epistemic reversibility; de-extinction producing functional analogs the literature itself frames as not faithful replicas) narrow but do not invalidate it. The architectural-property contrast holds — biological “addressable” is physical persistence without semantic guarantees; CKS “addressable” is identifier-level reactivation contract.

Cross-lineage mating compatibility overclaims — biology has horizontal gene transfer and symbiogenesis. Biology’s general case in canonical multicellular eukaryotes is species barriers via reproductive isolation. The exceptions — HGT pervasive in prokaryotes, symbiogenesis as rare and transformative once-or-a-few-times events, hybridization rate-limited and lineage-biased toward plants — narrow the contrast but do not invalidate it. They require specific mechanisms, are mostly lineage-specific, and are rate-limited per generation. CKS commits to governance-configured cross-lineage combinations available across all cell types as architectural commitment, not as exception. The narrowing-contrast cases help frame the contrast: biology can support cross-lineage combination architecturally — just not as the general case in canonical eukaryotic multicellular organisms.

Four briefer responses round out the section, plus one engaging an adjacent process-architecture lineage. *Why these three operations and not others such as replication, hibernation, fission, or fusion, and isn’t lifecycle-at-every-level just recursion with biological vocabulary?* Birth, mating, and death compose lifecycle as governance primitive: creation, combination, retirement. Other operations either reduce to these (fission as birth-by-division-of-one-parent; fusion as mating; replication as birth-with-content-derived-from-one-parent) or are state transitions within continued existence (hibernation, sleep, dormancy do not retire a cell from the architecture). What recurses at every level is not biological vocabulary but the commitment — birth, mating, and death as governed primitives with Paper 1’s six commitments inheriting at every composition level. *Mating-as-one-primitive collapses distinctions that matter.* The collapse is from biology-as-substrate-for-inheritance to CKS-as-substrate-with-governable-inheritance, not from biology’s analytically-distinguishable transmission mechanisms to one biological mechanism; §6.3’s pattern-variants development carries the differentiation. *Death-as-operational-retirement is just deprecation with extra steps.* Deprecation and archival in software engineering are themselves treated as architecturally distinct from deletion; §6.4 commits to both patterns simultaneously under one lifecycle primitive with type-specific results, with deprecation-style archival as one of two CKS death types rather than a softer version of one undifferentiated retirement primitive. *Lifecycle primitives at every level is process architecture; BPM has been doing this for decades.* Business Process Management (van der Aalst on process mining and Petri nets; Workflow Management Coalition standards; BPMN as standard notation) treats process composition, lifecycle states, and process-level governance as architectural primitives. Claim 3 differs in scope (lifecycle primitives at every composition level uniformly, not at process scope only), in primitive content (mating with three pattern variants and death with two distinct types as CKS-specific governance primitives over substrate content), and in living as governed substrate content rather than as workflow definitions in a BPM engine.

6.6 Paper 1 inheritance and downstream placement

Claim 3 inherits from Paper 1 (Li, 2026) without redefense. The most load-bearing inheritance is Paper 1 §5’s commitment to conflict preservation as first-class addressable substrate state with cell-level resolution by orchestration rules; this carries directly into the union mating pattern’s commitment to merge-time conflicts as persistent first-class substrate state. Paper 1 §3.3’s authority-vs-labor distinction is the load-bearing inheritance for actor-neutrality across all three lifecycle operations: birth governance, mating governance (with selective merge specifically), and both death types’ decision processes are governed by humans, with origination, curation, and retirement available to humans directly or to LLMs operating under human direction. Paper 1 §4’s substrate/LLM division at the governance boundary holds at every level of the lifecycle primitives, as Claim 2 generalized. Paper 1 §6’s linear-cost commitment underwrites both the lineage-preserved-union pattern and lineage-supersession archival without superlinear scaling cost. All six Paper 1 architectural commitments hold for cells subject to lifecycle operations at every composition level.

Claim 3 establishes the lifecycle primitives downstream claims operate over. As §7 develops, three evolution mechanisms operate on the modularity Claim 2 commits to through the lifecycle primitives Claim 3 establishes — mutation, directed selection, and the productive tension between them all presuppose primitives that create, combine, and retire the cells, aspects, and Selves on which evolution operates. As §8 develops, human governance takes different shapes across the three evolution mechanisms and across the two death types’ distinct decision processes; what Claim 3 names as architectural commitment, §8 specifies the governance machinery for. As §9 develops, the lifecycle primitives extend through every level of the enterprise brain Self, with hierarchical reorganization operating as architectural primitive under the same machinery — the brain’s lifecycle as one unit composed of cells through aspects under unified human governance.

7. Claim 4 — Evolution as three mechanisms in productive tension under unified human governance

7.1 Claim statement and the productive-tension framing

Claim 4 states Paper 2’s evolutionary commitment: a CKS-governed AI Self evolves through three mechanisms operating in productive tension across multiple levels of architectural composition simultaneously, with horizontal and vertical evolution axes, all under unified human governance. Two terms in the claim statement carry biological connotations that need bounding before the architectural development proceeds. *Evolution* in the CKS sense includes both undirected and directed mechanisms operating in tension — biology’s evolution-vocabulary connotes random-mutation-plus-natural-selection across deep time without goals, but two of CKS’s three mechanisms commit to direction under explicit governance-defined goals. *Tension* here names compositional integration of architecturally distinct mechanisms operating on architecturally separated layers, not adversarial equilibrium between competing forces; the mechanisms do not compete for resolution but compose under governance integration.

The three mechanisms are *instinct evolution* operating on the LLM and substrate-platform infrastructure layer through upstream upgrades, *DNA evolution* operating on the CKS substrate’s stabilized orchestration content under governance-defined goals, and *action-feedback evolution* closing the loop from recorded action experience back into governed DNA refinement. They operate on architecturally separated layers — Claim 1’s instinct/reasoning separation, Claim 2’s DNA/action distinction within every cell — under unified governance, and they apply at every level of Claim 2’s three-level structure: cells, aspects, and Selves all evolve through the same three mechanisms at appropriate scope. Evolution in CKS Selves is the governed dynamics by which the lifecycle primitives Claim 3 establishes — birth, mating, the governed retirement primitives — generate succession over time.

Claim 4 does not re-argue the foundational commitments the prior claims defend. Claim 1’s instinct/reasoning separation, Claim 2’s three-level structure and DNA/action distinction, and Claim 3’s lifecycle primitives are taken as given; what Claim 4 adds is the evolutionary dynamics that operate over them. The per-mechanism governance machinery — verification substrate patterns for instinct integration, authority architectures for directed selection on DNA, the proposal-and-acceptance machinery for action-feedback’s substrate-change cycles, the per-type decision processes for the death types — is §8’s territory rather than this section’s. Claim 4’s commitments are that the three mechanisms exist, that they operate under governance, that productive tension among them is architectural commitment rather than incidental property, and that they apply concurrently at every level of composition.

7.2 The three mechanisms

Instinct evolution

Instinct evolution sharpens what the cell can do at the instinct/infrastructure layer through upstream upgrades the Self does not control. Two kinds of upgrade carry the mechanism. Commercial or upstream LLM upgrades arrive when the LLM provider releases a more capable model, and the Self consumes the new model under verification rather than authoring it. Substrate-platform infrastructure upgrades — for example, migrating from a spreadsheet substrate to a dedicated platform — sharpen what the cell can do at the infrastructure layer without changing DNA

content. Both kinds arrive from upstream and are mutation-like in being undirected from the Self’s perspective: capability jumps that may be good or bad in any specific case, integrated through governed responses.

The localization is precise — only this one mechanism is mutation-like; the other two are directed. The actor question takes the same shape Paper 1 §3.3 holds throughout: humans authorize the upgrade integration through verification substrates, retirement of superseded reasoning where appropriate, and retention of reasoning as fallback or as active verification where appropriate; humans do not implement the LLM or the substrate-platform itself. Consumption is governed; what §8 specifies is the per-pattern verification machinery.

DNA evolution

DNA evolution operates on the CKS substrate’s stabilized orchestration content under explicit goals defined by human governance. Selection criteria for what DNA changes are accepted, retained, or reverted are themselves substrate content — the criteria are governable and modifiable under the same authority architecture that governs the content they evaluate. Direction here means goal-defined-by-human-governance: not goal-defined-by-internal-system-state, not goal-defined-by-fitness-function, not goal-defined-by-emergent-bias. Direction is external (humans set the goals) and persistent (governance authority is ongoing rather than front-loaded into a configuration that runs autonomously).

DNA evolution is the canonical directed-selection form CKS commits to. Biology has nothing directed in this sense; what biology calls fitness, CKS replaces with governance-defined success criteria authored and maintained by human governance and themselves substrate content. DNA changes can be drafted by LLMs operating under human direction, with humans holding authority over the orchestration rules under which LLMs propose changes and over the acceptance decisions; what §8 specifies is the per-step authority architecture.

Action-feedback evolution

Action-feedback evolution closes the loop from recorded action experience back into governed DNA refinement. Action-layer accumulation produces evidence — what worked, what failed, what surprised — and the loop from action evidence to DNA refinement is governed and human-mediated rather than autonomous. The architectural commitment is that the loop exists, that it is governed, and that humans or LLMs operating under human direction propose DNA changes from action patterns, with humans holding authority over the proposal-and-acceptance machinery and over the orchestration rules under which LLMs propose. What §8 specifies is who proposes, what authorization process, what verification, what reversion paths.

The closest biological analog is West-Eberhard’s (2005) plasticity-first framing: phenotypes leading genotypes via environmental induction with subsequent selection and genetic accommodation. The shape — operational state shaping what is later stabilized — is structurally close. The architectural delta is the directionality and the governance: in biology, environment-induced plasticity feeds into selection that then operates undirected on heritable variation; in CKS, recorded action evidence feeds into governed proposals humans authorize as substrate edits, with the loop discrete and auditable rather than emerging through differential reproduction. The closest software-engineering neighbor is configuration-management drift-detection — Kubernetes controllers, GitOps/ArgoCD — which commits to declarative state with runtime reconciliation toward the policy under automation. CKS’s action-feedback loop differs from both neighbors on two axes: the actor question (governed proposals authorized under multi-party authority, not automated reconciliation) and the target of refinement (governed substrate content, not declarative configuration).

Productive tension as architectural result

The three mechanisms produce something none alone can deliver. Instinct evolution provides exploratory capacity — undirected change discovers possibilities directed selection would not search; DNA evolution provides directional capacity — purposeful change toward governance-defined goals; action-feedback evolution provides empirical feedback — operational experience grounding both the directional change and the consumption of new instinct layers. The mechanisms operate on architecturally separated layers, so they do not interfere; they compose because governance integrates them.

Biology's evolution operates only through one mechanism — natural selection acting on undirected mutation, with EES extensions adding niche construction, developmental plasticity, and inclusive inheritance as biasing and bootstrapping processes operating on natural selection's foundation rather than as parallel directed-selection mechanisms. CKS commits to undirected and directed selection mechanisms operating in parallel under unified governance. Biology cannot have directed selection because biology has no goals; CKS has directed selection because CKS has governance-defined goals, with selection criteria themselves substrate content. Productive tension is the first architectural property where CKS differs from biology's evolved constraints.

7.3 Multi-level simultaneous evolution

The three mechanisms apply at every level of Claim 2's three-level structure concurrently. Cells evolve through their internal DNA/action loops; aspects evolve through their constituent cells' evolution plus aspect-level orchestration changes; Selves evolve through their constituent aspects' evolution plus Self-level structural reorganization. The levels' evolution is loosely coupled — cells can evolve at one rate while aspects evolve at another and Selves at a third — with governance integration across levels rather than selection-pressure conflict. Different humans with different authority scopes govern different levels, and cross-level coordination flows through humans authorizing substrate changes affecting multiple levels.

The architectural lineages this commitment sits adjacent to engage at the level itself. The Major Transitions framework (Maynard Smith and Szathmáry 1995) names new levels of selection emerging from ensembles of pre-existing entities, structurally analogous to Claim 2's stacked cell/aspect/Self composition; what biology lacks and CKS adds is governance integration across the levels. Hierarchical evolutionary algorithms — Potter and De Jong's (2000) “an architecture for evolving such subcomponents as a collection of cooperating species” — couple levels through shared or nested fitness functions, with subcomponent fitness defined in terms of system-level performance. Vanchurin et al.'s (2022) framing of evolution as multilevel learning preserves a single-loss-function frame across separated scales. CKS gives different levels distinct, governed roles not collapsed into one objective, with Paper 1 §6's linear-cost composition underwriting the tractability of concurrent evolution at scope.

7.4 Dual-axis evolution — horizontal and vertical

Two axes characterize how evolution can change the Self. Horizontal evolution refines content within existing structure: existing cells refine their DNA; existing aspects refine their cell composition; existing Selves refine their aspect arrangement. Vertical evolution reorganizes structure itself: cells split or merge; aspects gain or lose constituent cells; Selves gain or lose aspects; relational roles reconfigure. Both axes operate at every level. Both axes operate through the three mechanisms.

The substrate-content commitment Claim 2 establishes is what makes vertical evolution architecturally available. The orchestration substrates that define composition are themselves substrate content humans govern in the same way DNA content is governed. Structural reorganization is therefore a substrate-edit operation, not a developmental-genetic-program rewrite or a graph-genome representation-rewrite. NEAT incrementally elaborates structure on a graph-genome representation; DARTS uses continuous relaxation of architecture representation with gradient descent over the search; AutoML-Zero operates on program-code representations. In each case the structural-evolution representation is categorically distinct from the runtime content. CKS's vertical evolution operates on substrate content of the same kind as the content evolved within structure — one architectural commitment carries both axes.

Structural evolvability on operational timescales is the second architectural property where CKS differs from biology's evolved constraints. Biological body plans take millions of years to evolve through evo-devo: “evolution of body plans must depend upon change in the architecture of developmental GRNs” (Davidson and Erwin 2006), and within a species the body plan is essentially fixed. CKS Selves restructure through governed reorganization on operational timescales as a routine architectural property — substrate-edit operations rather than the extensive often-lethal genetic rewiring biology requires for body-plan-level change.

7.5 Reviewer-objection responses

Five objection clusters consolidate eight anticipated objections, with two paired-response opportunities executed.

Three mechanisms as decomposition; productive tension as GA-with-extra-steps. The two related EC objections collapse Claim 4’s three mechanisms into operator-level decomposition of one mechanism (mutation+selection over a population), or treat productive tension as GA’s exploration-exploitation tension already exhibited at operator level. The differentiation holds on three architectural axes. Mechanism count: GA’s mutation and selection are operators within one search algorithm; CKS’s three mechanisms operate on architecturally separated layers. Goal source: GA’s fitness function is human-authored at design time and fixed during the run; CKS’s governance-defined goals are themselves substrate content under continuous human authority. Substrate: GA evolves parameters within fixed encoding; CKS evolves substrate content on dual horizontal+vertical axes. The architectural-separation between mechanisms — not the operator-level decomposition — is what carries the differentiation.

Action-feedback as recursive self-improvement or RLHF. Two related objections target the action-feedback mechanism specifically. Recursive self-improvement frameworks — Schmidhuber’s Gödel Machine rewriting any part of its own code once it has proof the rewrite is useful, and the broader RSI discourse where loss-of-control is the named risk — commit to autonomous improvement loops where the system improves itself within fixed scaffolding once configured. RLHF and online learning operate on model weights via gradient updates from feedback signals: feedback produces numeric rewards or preference losses, and gradient descent updates weights as automation. Constitutional AI and RLAIF (Bai et al. 2022) are different in their commitments along this same architectural axis, “training a harmless AI assistant through self-improvement, without any human labels identifying harmful outputs” — an architectural choice that achieves what it sets out to achieve while taking a different position than Claim 4 takes on the actor question.

The hostile reviewer’s specific challenge — *if Constitutional AI achieves the safety property without humans in each loop, why is human authorization at every step architecturally necessary?* — folds against the antecedent-infrastructure-vs-runtime-substrate distinction §4.4 introduces and §8.5 develops compactly. Constitutional AI commits the human-authored substrate (the constitution) at training time as antecedent infrastructure that shapes weights through RLAIF; once weights are fixed at deployment, the substrate’s authority lives inside the model rather than in a runtime artifact humans modify during operation. CKS commits the substrate to runtime-governable artifact, with humans retaining authority during operation. Both achieve what they set out to achieve at the architectural level the question is posed at; the differentiation is the temporal locus of substrate governance, not which approach is correct on the actor question. The two architectures answer different questions. Action-feedback evolution operates on substrate content (the governed knowledge and orchestration objects the LLM consults, the rules under which it operates, the criteria by which its outputs are evaluated) through governance-mediated proposals and explicit authorization. Once configured, RSI systems improve themselves and RLHF-style loops update weights from feedback signals; CKS systems require ongoing human authorization for each substrate change, with humans or LLMs operating under human direction proposing the changes and humans holding authority over the proposal-and-acceptance machinery.

Multi-level simultaneous evolution as hierarchical genetic algorithms. Hierarchical EAs (CCEA) and multi-level selection theory commit to selection-pressure flow across levels with potential conflict between levels: subcomponent fitness defined in terms of system-level performance, with cross-level coordination resolved (if at all) via further selection or evolved suppression mechanisms. CKS commits to loosely-coupled rates with governance integration across levels rather than selection-pressure conflict. Levels can evolve at different rates governed by different humans with different authority scopes; cross-level coordination flows through humans authorizing substrate changes affecting multiple levels rather than through shared evaluation functions or evolved policing.

Quality-diversity already has both directed and undirected mechanisms in tension. The most sophisticated objection comes from the quality-diversity literature, which explicitly combines undirected exploration with directed selection. Lehman and Stanley’s novelty search rejects explicit objectives in favor of “search with no objective other than continually finding novel behaviors”; MAP-Elites builds an archive of high-performing solutions across designer-chosen dimensions of variation; Soros and Stanley specify four conditions under which standard evolutionary search can remain open-ended. The QD-with-archive line — MAP-Elites with structured archives, hierarchical archives, and behavior-space partitioning — produces population structure across operators that approaches architectural separation more closely than the operators-within-one-search framing alone admits. The differentiation holds on three axes the QD literature does not occupy. No governance: QD assumes human-authored descriptors, criteria, minimal conditions, and environment rules, with no mechanism for evolving or governing those criteria themselves. No multi-level architecture: QD operates on one population at one level of representation, with multi-level structure appearing

as side effect (Picbreeder’s user community vs. branches; Cully et al.’s offline repertoire vs. online adaptation) rather than as evolutionary process over evaluators or institutions. No architectural separation between mechanisms: QD combines directed and undirected pressures as operators within one search algorithm — same population, same variation operators, same selection loop, with the criterion swapped or the archive structured. The MAP-Elites-with-structured-archive case approaches architectural separation but the archive remains internal to the search algorithm rather than a separately-governed substrate layer. CKS’s three mechanisms operate on architecturally separated layers under governance integration. The architectural-separation move is what carries the differentiation, not any single one of the three axes.

Where-CKS-differs-from-biology overclaims. Two related objections charge that EES extensions instantiate directed components and that biology’s rapid-speciation cases narrow the structural-evolvability timescale contrast. The biology-overclaims discipline serves both. EES extensions — niche construction, developmental plasticity, inclusive inheritance — extend the Modern Synthesis but operate on natural selection’s foundation rather than as parallel directed-selection mechanisms. EES “direction” means biases in variation and modified selective environments organisms construct, not goal-defined selection by an internal teleological agent or external authority — selection itself remains undirected even when developmental plasticity or niche construction shapes the variation it operates on. Rapid speciation cases — cichlid radiations over thousands of generations within a teleost body plan, Galápagos finch beak shifts within a finch body plan, polyploid plants emerging in one generation with shifted traits within an existing plant body plan — narrow the timescale contrast slightly without invalidating the structural-evolvability point. Even biology’s textbook rapid-radiation cases produce morphological diversification within constrained morphospaces over thousands to millions of years, and body-plan-level reorganization requires the GRN-architecture changes Davidson and Erwin specify. CKS’s vertical-axis evolution operates on substrate content humans govern, not on developmental-genetic programs requiring extensive often-lethal genetic rewiring.

7.6 Paper 1 inheritance and downstream placement

Claim 4 inherits Paper 1’s six architectural commitments at every level of the evolution mechanisms it establishes. Paper 1 §3.3’s authority-vs-labor distinction holds across all three mechanisms — humans hold authority over substrate structure, orchestration rules, and modification rights, while LLMs operating under human direction may draft, populate, and update substrate content. Paper 1 §6’s linear-cost composition underwrites multi-level simultaneous evolution being tractable. Paper 1 §4’s substrate/LLM division at the governance boundary holds throughout — instinct evolution operates on the inside-the-model layer through upstream upgrades; DNA evolution and action-feedback evolution operate on the outside-the-model layer through governed substrate-content modification.

Inheritance from prior claims is direct. Claim 1’s instinct/reasoning separation is what the three mechanisms operate within. Claim 2’s three-level structure is what multi-level simultaneous evolution operates across; Claim 2’s DNA/action layer distinction is what the action-feedback loop operates between; Claim 2’s commitment that the orchestration substrates defining composition are themselves substrate content is what makes vertical-axis evolution architecturally available. Claim 3’s lifecycle primitives are what evolution operates over: birth as the creation primitive for offspring instantiation, mating as the inheritance primitive whose three pattern variants (union, selective merge, lineage-preserved union) configure variation generation, and the death types Claim 3 establishes as governed retirement primitives. The causal arrows from evolution mechanisms to the death types are Claim 4’s specific work: lineage supersession follows from DNA evolution producing better orchestration, and capability supersession — the evolution-driven succession-variant within Claim 3’s lineage-supersession territory where a substrate’s function is now handled elsewhere — arises through instinct evolution sharpening what previously required reasoning or through DNA evolution consolidating substrates. SOFAI’s S2-to-S1 skill migration is the closest neighbor for the instinct-evolution — capability-supersession arrow, with SOFAI’s automated-internal migration via metacognitive controller contrasting with CKS’s governed-and-substrate-mediated capability-supersession decision.

As §8 develops, human governance takes different shapes across evolution mechanisms — verification substrates governing instinct integration, the standard authority architecture governing directed selection on DNA, human-mediated approval governing action-feedback’s substrate-change cycles, and the two death types’ distinct decision processes specified per type, with the instinct/reasoning boundary itself substrate content humans tune. As §9 develops, the three mechanisms extend through every level of the enterprise brain Self, with the same evolutionary dynamics underwriting hierarchical reorganization as architectural primitive at enterprise scope and the brain’s

evolution as one unit under unified human governance. Claim 4 makes evolution available; §8 specifies how it is governed.

8. Claim 5 — Human governance across CKS evolution as multi-shaped commitment

8.1 Claim statement and the authority-vs-labor extension to evolutionary dynamics

Claim 5 states Paper 2’s governance commitment: a CKS-governed AI Self holds human governance across its evolution as multi-shaped — distinct governance shapes operating on the distinct evolution mechanisms Claim 4 establishes. Three components carry the commitment, each addressing a different facet of how human authority operates over the evolving Self. The first specifies governance shape per evolution mechanism: verification-substrate machinery for instinct integration, authority architecture for DNA evolution, proposal-and-acceptance machinery for action-feedback’s substrate-change cycles. The second specifies that the boundary between instinct and reasoning Claim 1 establishes is itself governed substrate content rather than fixed architectural property. The third specifies distinct death-type governance processes — Claim 3’s death types operate under decision processes appropriate to their distinct underlying causes, with capability supersession (Claim 4) the succession-variant of lineage supersession.

The unifying architectural move is Paper 1 §3.3’s authority-vs-labor distinction extended from substrate-content-and-orchestration-rules to evolutionary dynamics. Paper 1 (Li, 2026) §3.3 renders the commitment as portable phrase: governance is an authority architecture, not a review workflow. Humans hold authority over substrate structure, orchestration rules, and modification rights; LLMs operating under human direction may draft, populate, and update substrate content. Claim 5 extends this across evolutionary dynamics: humans hold authority over what evolves, in what direction, through what processes, with the labor available to humans directly or to LLMs operating under human direction. This is where Paper 2’s commitment to governance-as-architecture rather than governance-as-policy-overlay is most concretely defended — the governance shapes are not policy documents written about an evolving system; they are architectural machinery the system instantiates as substrate content under human authority.

Before developing the three components, one bound on the term *governance* belongs in the opening. The term carries connotations from three adjacent domains, none of which CKS adopts. From biological regulation and homeostasis, governance imports feedback-driven stability through automatic correction — distributed and mechanism-driven without authority decomposition. From political theory, governance imports democratic legitimacy through reasoned consensus and accountability-as-constraint. From corporate compliance, governance imports retrospective audit, fiduciary duty, and risk-and-control framing oriented to fraud prevention. The architectural shape Claim 5 commits to is authority-driven rather than feedback-driven, expertise-grounded rather than consensus-grounded, prospective approval rather than retrospective audit, and operational-technical rather than democratic-legitimate or fiduciary-legal. The bound is placed preventively here so the three-component development that follows can use *governance* freely as architectural commitment.

8.2 Per-mechanism governance shape

The first component states what the multi-shaped commitment specifies architecturally: each of Claim 4’s three evolution mechanisms operates under a distinct governance shape, with shape co-determined with mechanism shape rather than overlaid on it. The unifying authority-vs-labor commitment holds across all three; what differs is the governance machinery the commitment instantiates. The three sub-passages that follow develop each per-mechanism shape in turn — verification-substrate machinery for instinct evolution, authority architecture for DNA evolution, proposal-and-acceptance machinery for action-feedback evolution — and the section closes with the central differentiation against the mechanism-agnostic AI-governance frameworks the published landscape currently dominates.

Instinct evolution is governed through verification-substrate machinery. The mechanism’s distinguishing feature is that capability changes arrive from upstream — LLM provider releases, infrastructure platform upgrades — rather than from goals the Self pursues. The governance shape responds to that asymmetry: verification substrates determine integration patterns when upgrades arrive, with parallel-run patterns evaluating whether prior CKS reasoning

can be retired now that new instinct handles what reasoning previously handled, retained as fallback for cases where instinct fails, or retained as active verification on safety-critical paths. The verification-substrate term bears bounding: CKS verification provides operational assurance through authority checkpoints, not epistemic certainty through formal proof, statistical guarantees through probabilistic bounds, or regulatory compliance through specification conformance. The closest engineering-pattern neighbor is ML deployment verification — canary, blue-green, shadow launches, A/B testing — which commits to verification-as-engineering-pattern for risk-mitigating changes the deploying engineer controls. CKS verification-substrate machinery commits to verification-as-governance-machinery for authority-exercising review of capability changes others control; the engineering-pattern literature has no equivalent for the upstream-controlled-substrate scenario.

DNA evolution is governed through an authority architecture extending Paper 1 §3.3 directly. The architecture decomposes who can propose substrate changes, who can authorize, what verification regime applies, what reversion paths exist. The authority-architecture term bears bounding: CKS authority architecture combines Weber’s legal-rational legitimacy (authority from positions and rules), organizational decision-rights clarity (explicit role specification with audit trails), and access-control enforcement (technical gates plus audit logs) without importing political-legitimacy framing, empowerment optimization, or pure-prevention access control. The closest software-engineering neighbor exhibiting the recursive-governability pattern Claim 5 commits to is Apache PMC governance, where “PMCs vote on new committers and PMC members for their project” — the rules defining who can authorize project changes are themselves under governed change control; GitHub CODEOWNERS exhibits the same pattern at smaller scope. CKS extends this from software-development to AI-Self evolutionary authority architecture: the orchestration rules under which substrate changes get proposed and authorized are themselves substrate content under the same authority architecture they instantiate. Recursive governability becomes architectural commitment, not best practice.

Action-feedback evolution is governed through proposal-and-acceptance machinery for substrate-change cycles driven by accumulated action evidence. Humans or LLMs operating under human direction propose DNA changes from action patterns; humans hold authority over the proposal-and-acceptance machinery and the orchestration rules under which proposals are constructed; verification regimes and reversion paths are themselves substrate content. The closest published neighbor is human-in-the-loop ML literature (Mosqueira-Rey et al. 2023; Wu et al. 2022), where humans appear in autonomously-running loops as reviewers of specific outputs. The actor question in HITL is bounded — who approves which decisions — but the substrate question is not: the loop’s per-step machinery is antecedent infrastructure designed once and operated thereafter, not substrate content under continuous governance. CKS commits both questions: the loop’s per-step machinery is itself substrate content humans govern.

Across the three sub-passages, the central differentiation is against mechanism-agnostic AI-governance frameworks. The pattern dominating the published landscape — EU AI Act Article 9’s continuous iterative process throughout the lifecycle, NIST AI RMF’s four functions integrated across the lifecycle, ISO/IEC 42001’s intentionally broad and sector-agnostic principles, and frontier-AI safety case frameworks proposing argument categories applicable across mechanisms — commits to unified governance with mechanism-agnostic principles as theoretical commitment. The framing treats governance as policy-overlay above mechanism: same principles apply regardless of which mechanism produced a capability shift. The closest published organizational neighbor is Anthropic’s own Responsible Scaling Policy, which commits to “Capability Thresholds: Specific AI abilities that, if reached, would require stronger safeguards than our current baseline” — capability levels (AI Safety Levels) trigger uniform governance responses regardless of which mechanism produced the capability shift. CKS commits to mechanism-specific governance machinery: distinct mechanisms produce distinct governance work, and shape is co-determined with mechanism rather than overlaid on it. The two commitments operate at different scales — organizational governance for Anthropic, Self-architecture governance for CKS — and are composable rather than mutually exclusive. CKS Selves deployed within an Anthropic-governed organization would inherit organizational RSP commitments at the deployment scope while operating mechanism-specific governance at the Self-architecture scope. A reviewer with DevOps/MLOps background may register that the three governance shapes correspond to recognizable existing patterns: verification-substrate machinery for instinct evolution shares structural shape with canary/blue-green/shadow-launch patterns; authority architecture for DNA evolution shares structural shape with GitOps PRs and CODEOWNERS; proposal-and-acceptance machinery for action-feedback shares structural shape with human-in-the-loop ML approval gates. The architectural contribution is not the patterns considered individually — those are recognizable engineering practice — but their commitment as architectural primitives co-determined with mechanism shape under one unified authority architecture, with the orchestration rules under which the patterns operate themselves substrate content humans govern.

The DevOps/MLOps mapping is structural correspondence, not collapse: each pattern in the engineering literature is a deployment hygiene practice, while CKS commits each to formal architectural primitive that mechanism-agnostic governance frameworks cannot specify.

8.3 The instinct/reasoning boundary as governed substrate content

The second component states a commitment less crowded by adjacent literature than Components 1 or 3 and produces a deliberate gap-honest finding rather than dense closest-neighbor engagement. The boundary between what lives in the instinct layer (the LLM, fast-pattern responses) and what lives in the reasoning layer (CKS substrate, governed deliberation) is itself substrate content the orchestration substrate governs. As the LLM’s instinct capability sharpens through upstream upgrades, what previously required explicit reasoning may now be handled by instinct alone — but the boundary’s relocation is governed: humans hold authority over the orchestration rules under which the boundary’s location is determined, and high-stakes decisions can be architecturally pinned to the reasoning layer regardless of how capable instinct becomes. This makes Paper 1’s substrate/LLM division at the governance boundary not a fixed architectural property but a continuously-tunable substrate-content commitment under the same authority architecture that governs DNA.

The closest published architectural neighbor is Cognitive Core (Seck 2026), which architects demand-driven delegation between System 1 and System 2 within architecture-time-defined bounds: the typed primitive vocabulary is design-time architecture rather than runtime substrate content, with the paper’s framing positioning configuration as the governed layer. CKS commits to the inverse: the parameter space defining the boundary between fast pattern-matching and explicit deliberation is itself governed substrate content. The architectural difference sits at whether the design-time-vs-runtime boundary lies at parameter-space-content (configuration) or at parameter-space-existence (substrate); adjacent governance territories — feature-flag governance, interpretability-as-governance, model-vs-policy decisions in ML deployment — stop at the configuration-vs-substrate-content barrier in the same way. Adjacent in a different direction, self-modifying AI architectures (Schmidhuber’s Gödel Machine, the broader RSI discourse) and meta-learning frameworks treat the governance boundary itself as modifiable artifact, but the modification authority sits with the system rather than with humans operating runtime authority over the boundary’s substrate content (Claim 1 §4.4 and §7.5 treat this differentiation in detail). The honest finding for Component 2 is that the boundary-as-governed-substrate-content commitment under runtime human authority is narrow rather than entirely empty territory in the surveyed literature: the territory is bounded by Cognitive Core on one side (architecture-time configuration) and by self-modifying AI on the other (substrate-modifying-itself); CKS sits between these, with humans holding runtime authority over substrate the system does not modify on its own. The synthesis lands the finding as deliverable rather than force-fit more distant neighbors into a closer-than-they-are framing.

8.4 Distinct death-type governance processes

The third component states what governance commits to over the death types Claim 3 establishes and Claim 4 develops causal arrows for. Before specifying the per-type machinery, one terminology reconciliation belongs here. Claim 3 names two death types — functional obsolescence and lineage supersession — distinguished by whether a cell’s substrate is genuinely deleted or retired-with-archival; Claim 4 names the evolution-driven case where a substrate’s function is now handled elsewhere as capability supersession. Claim 5’s reconciliation: functional obsolescence and lineage supersession are the two primary death types, with capability supersession the succession-variant where evolution drove the substrate’s function elsewhere — its source substrate either retired-with-archival as standard lineage supersession or retained-as-fallback as architectural option preserved. The two-type framing aligns with mature software-engineering practice (PEP 387, COPE, API lifecycle, Microsoft Purview), where retirement is one primitive with type-specific parameters; what Claim 5 adds is per-cause governance rather than unified-with-parameters processing. The retirement term bears bounding: CKS retirement is status change from active to non-preferred with archival or retention patterns varying by type — not the irreversibility connotation of organizational retirement, not the voluntary-decision framing of HR contexts, not the unified-EOL framing of software product lifecycles.

Functional obsolescence is governed through deletion-decision machinery. A cell whose function is no longer needed is removed under a governed retirement decision; substrate resources are released; the deletion is irreversible by architectural commitment. The governance shape responds to that irreversibility: decision processes are bounded

by impossibility of reversion, with humans holding authority over what gets deleted under what criteria. The closest computational neighbors carry the deletion-shape — Microsoft Purview Priority Cleanup with multi-tier approval architecture parameterized by retention status; capacity-management decision processes; software systems where resource pressure drives genuine deletion.

Lineage supersession is governed through retirement-with-archival decisions. After a superior offspring emerges from mating or DNA evolution, parent cells are retired but their substrate content remains addressable per Paper 1 §6’s linear-cost commitment — archived cells are not truly dead, since their substrate is reactivatable through governed processes. The governance shape responds to that reactivability: decision processes can preserve reversion paths the deletion-shape cannot, and the per-step machinery accepts retention-as-fallback when the retired substrate’s function is now handled elsewhere through evolution (the capability-supersession sub-case). The closest neighbors carry the archival-shape: software deprecation governance with PEP 387’s steering-council exception authority; academic retraction-and-archival under COPE guidelines; regulated-industry archival processes; API deprecation with version-specific lifecycle parameters. All commit to unified retirement governance with case-specific parameters.

The differentiation across Component 3 is against this unified-with-parameters pattern. The neighbors do real architectural work CKS inherits without redoing — multi-tier approval architectures, steering-council exception authority, version-specific lifecycle parameters. What CKS commits to that the neighbors do not is per-cause governance with type-appropriate decision processes. Resource-pressure deletion is irreversible and bounded by deletion-irreversibility; evolutionary-supersession archival is reactivatable given substrate-addressability; capability-migration retirement-or-retention preserves fallback options. The objection that the two death types could be re-described as two retention modes of one retirement primitive — a parameter-tuning rather than an architectural distinction — has surface plausibility, since irreversibility-versus-reversibility is parameterizable in principle. The architectural defense is that the parameter binds different governance shapes when the cause differs: the irreversibility of resource-pressure deletion bounds the decision process to a shape that cannot allow exception-after-the-fact, while the reversibility of evolutionary-supersession archival allows decision processes to preserve reversion paths the deletion shape cannot. Unifying these as one parameterized retirement primitive is acceptable for software-engineering accounting; for safety-critical AI deployment, where the difference between irreversible deletion and reactivatable archival shapes what the governance is responsible for, the two-type architectural distinction holds the operational consequence the parameter framing collapses.

8.5 Reviewer-objection responses

Six anticipated objections cluster into four response paragraphs through two paired-response treatments.

Why mechanism-specific at all — governance is governance, and three governance shapes is just complicated governance. The two related objections turn on the same architectural distinction. The mechanism-agnostic pattern across major AI-governance frameworks (EU AI Act, NIST RMF, ISO/IEC 42001, Anthropic RSP) is theoretical commitment to governance-as-policy-overlay above mechanism — the framework documents themselves describe their principles as intentionally broad and applicable throughout the lifecycle. CKS commits to governance-shape-co-determined-with-mechanism-shape: verification machinery for upstream-arriving capability shifts cannot work for proposed substrate changes (different actor, timing, substrate target), and authority architecture for proposed substrate changes cannot work for action-feedback approval (different evidence basis, different reversibility). Mismatched governance produces governance failures; the three shapes serve architecturally distinct work the unified frameworks cannot perform.

This is just human-in-the-loop with extra vocabulary. Canonical HITL taxonomies (Mosqueira-Rey et al. 2023; Wu et al. 2022) bound human “control” to influence over what the model learns or approval of specific outputs, with loop machinery antecedent to the humans-in-loop engagement. More recent extensions like HAIG (Engin 2025) introduce three-dimension framings — Decision Authority Distribution, Process Autonomy, Accountability Configuration — that gesture toward substrate-level analysis but commit to designing oversight mechanisms rather than to humans holding ongoing authority to modify loop machinery. CKS differs in kind: humans hold authority over substrate including the orchestration rules under which LLMs propose changes, with the loop’s per-step machinery itself substrate humans govern. The actor and substrate questions are architecturally separable; HITL bounds the actor question while CKS commits both. The substantive distinction is substrate locus, not vocabulary.

The boundary as governed substrate content is just configuration; three death-type governance processes is over-engineering. Configuration in modern systems (LaunchDarkly, Statsig, OPA) commits to bounded parameter spaces designed at deployment time; the parameter space itself is design-time choice. CKS commits to the parameter space itself being governed substrate — humans hold authority over what the boundary even consists of, not just where within a fixed space it sits. Unified retirement governance (PEP 387, COPE, API lifecycle, Microsoft Purview) commits to one process with case-specific parameters; CKS commits to per-cause governance because resource-pressure deletion is irreversible while evolutionary-supersession archival is reactivatable. Both parsimony framings would lose architectural commitments that matter for governance under continuous capability evolution.

Authority-vs-labor extension to evolution is just delegation with extra steps. The parsimony framing reduces governance to delegate selection plus oversight scope, with delegates operating autonomously within scope. The framing breaks down for AI Selves where the substrate within which the delegate operates is itself evolving. Constitutional AI architectures illustrate the contrast directly: governed content lives in a constitution baked into model weights through training, with the constitution functioning as antecedent infrastructure — authoritative specification fixed before runtime, modifiable only through retraining. The actor (the AI critiquing and revising according to the constitution) is delegated work within antecedent infrastructure. CKS commits to runtime substrate humans modify through governed processes without retraining: the orchestration rules under which substrate changes get proposed and authorized are themselves substrate content under the same authority architecture they instantiate, not antecedent infrastructure designed once and operated thereafter. The actor question (who does the labor) and the substrate question (what is the locus of governed content, and is that locus itself governed) are architecturally separable; delegation answers only the actor question; Constitutional AI bounds the substrate question to training-time artifact; CKS commits both questions and treats the substrate as continuously governable. Recursive governability — the architectural commitment that the operating environment is itself governed substrate, not pre-designed scaffolding — is what the actor-vs-substrate-question separation rests on architecturally and what makes the separation more than vocabulary.

8.6 Paper 1 inheritance and downstream placement

Claim 5 inherits Paper 1’s six architectural commitments. Paper 1 §3.3’s authority-vs-labor distinction is the load-bearing inheritance for the entire claim. Paper 1 §4’s substrate/LLM division at the governance boundary holds across all three governance shapes; Paper 1 §5’s two-level conflict-handling extends to evolutionary dynamics with conflicts (between proposal authority’s intent and authorization authority’s review, for instance) preserved as first-class substrate state; Paper 1 §6’s linear-cost commitment underwrites both lineage supersession’s reactivatable archive and capability supersession’s potential retention as fallback. Inheritance from prior claims is direct: Claim 1’s instinct/reasoning separation is what Component 2 commits to as governed boundary; Claim 2’s three-level structure is where governance shapes apply concurrently and its DNA/action layer distinction is what DNA evolution and action-feedback evolution operate on; Claim 3’s lifecycle primitives are what Component 3 specifies governance machinery for, with mating-pattern selection governed under the DNA-evolution authority architecture and per-pattern shapes for the three mating patterns research territory for downstream work; Claim 4’s three evolution mechanisms are what Component 1’s per-mechanism shapes operate on, and Claim 4’s causal arrows from evolution mechanisms to death types are what Component 3’s per-type machinery instantiates.

The actor-neutrality from Paper 1 §3.3 holds throughout. Human-governed does not mean human-authored: humans own authority over orchestration rules, acceptance criteria, and verification regimes; the labor is available to humans directly or to LLMs operating under human direction. As §9 develops, the multi-shaped governance Claim 5 specifies extends through every aspect-domain of the enterprise brain Self, with institutional knowledge living as Self-level substrate content under multi-shaped governance per aspect-domain and organizational authority structures composing with the architecture’s commitments. The composability framing the Anthropic-RSP-vs-CKS contrast surfaces — organizational governance at deployment scope, Self-architecture governance at internal scope — anticipates that development. Claim 5 makes the governance shapes explicit; §9 develops how those shapes compose with the human organizational structures within which CKS Selves are deployed.

9. Claim 6 — The enterprise brain Self as architecturally coherent design pattern extending Paper 1’s cell concept

9.1 Claim statement and the vision-claim framing

Claim 6 states Paper 2’s vision claim: the enterprise brain Self is architecturally coherent as a design pattern extending Paper 1’s (Li, 2026) cell concept — a functioning unit composed of cells through aspects under unified human governance, sharing the same lifecycle and evolution as any individual cell, with substrate-shared topology supporting coherent cross-aspect coordination where integration-glued enterprise architectures take a different structural approach. The claim is offered as architectural vision pending debate and observation in real-world enterprise deployment; empirical validation is downstream work.

Two architectural moves carry the claim. *Wholeness* commits the Self as functioning unit with its own lifecycle, evolution, and governance, not aggregation of cells with cell-level commitments. *Topology consequence* commits the substrate-shared aspect arrangement as producing the cross-aspect coherence at structural-pattern level. Three supporting commitments operate the moves: bidirectional traceability and propagation across all composition levels and aspect-domains; hierarchical reorganization as architectural primitive; institutional knowledge as Self-level substrate content under multi-shaped governance. §9.2 develops the wholeness move, §9.3 the topology consequence move, and §9.4 the three supporting commitments as the primitives the wholeness commitment is testable against.

The term *architecturally coherent as a design pattern* throughout names the architectural-commitment sense: the architecture’s commitments compose into a consistent design pattern at enterprise scope, with realizing the pattern in a particular deployment requiring engineering and governance work but not architectural change. This distinguishes the claim from goal-setting realism, engineering-feasibility implementability, or aspirational plausibility — none of which assert that the architecture itself composes into a coherent design pattern at the scope claimed.

Paper 2’s Claim 6 follows the structural model Paper 1 §9 establishes for vision-claim writing — claim statement extending an established commitment, calibrated humility about the empirical claim, concrete rendering downstream — there through the proof-of-concept, here through the §10 thought experiment. The acknowledgment is disciplinary parallel rather than authorizing parallel: the same shape of voice register, with the architectural content itself being Paper 2’s own.

Humans hold authority over the brain’s substrate content, over the orchestration rules under which cells and aspects operate, and over the governance shapes per aspect-domain; the brain functions as architectural unit under unified human governance, with the labor of operating the brain available to humans directly or to LLMs operating under human direction.

9.2 The wholeness move

The wholeness commitment names architectural wholeness — the Self functions as a coherent unit rather than a loose collection of parts because its architectural commitments produce unit-level lifecycle, governance, and evolution properties. The term *wholeness* in the CKS sense is engineering-technical, not philosophical, and naming the adjacent connotations CKS does not adopt at the opening lets the architectural development that follows use *wholeness* freely. Philosophical holism (Smuts; the wholes-vs-parts debate; Koestler’s holon framing) commits to claims about wholes irreducible to parts; the holon concept has been operationalized in holonic manufacturing systems (Van Bruseel et al. 1998) and the PROSA reference architecture as engineering instantiations of multi-level composition with relational properties, which CKS sits adjacent to at the multi-level structural-pattern level even as the wholeness commitment itself stays narrower than the holonic-architectures position. Gestalt psychology means perceptual wholes — the perceptual experience of a unified figure rather than its component parts. Organizational-integrity rhetoric imports holistic-management vocabulary into descriptions of how groups function. CKS’s wholeness commitment is narrower than any of these: architectural wholeness as engineering invariant supported by specific architectural primitives. The closest engineering-tradition anchor is the system-design coherence framing where “components should fit together, seamlessly working alongside each other”; Christopher Alexander’s structurally-defined wholeness sits in this lineage as architectural pattern with the metaphysics bounded out. The bound bears precision: what CKS borrows from Alexander is the architectural-pattern shape — the commitment to how all parts relate as a property of the architecture rather than as a property added after the fact — and what CKS does not borrow is Alexan-

der’s metaphysical grounding in *The Nature of Order* of wholeness as property of nature. The borrowing is at the architectural-pattern level (how parts relate within a designed system) rather than at the metaphysical-foundation level (whether wholeness is a property of natural systems generally). A reviewer pressing the selective-uptake objection — that Alexander’s pattern theory is metaphysically grounded and that bounding the metaphysics out leaves only folk software-engineering wisdom — encounters the response that the system-design coherence framing CKS actually anchors on is engineering-tradition independent of Alexander’s metaphysical foundations, with Alexander’s structurally-defined wholeness named as adjacent lineage rather than as load-bearing authority for CKS’s architectural commitment.

The wholeness move’s load-bearing differentiation is supported-by-specific-primitives versus asserted-as-philosophical-or-perceptual-position. Three architectural primitives carry the commitment: bidirectional traceability across composition levels and aspect-domains, hierarchical reorganization as architectural primitive, and institutional knowledge as Self-level substrate content. Without these, the architecture is non-conformant to the wholeness commitment; with these, the brain functions as unit. The wholeness commitment is testable against the primitives, and §9.4 develops them in turn.

The architectural-commitment shape CKS commits to has lineage at adjacent scopes. Modular-monolith advocacy commits to wholeness at single-application scope — one deployment, one process space, one coherent unit, with internal modular boundaries enforced strictly enough that the system-as-unit framing holds at the deployment level. Consensus protocols (Paxos, Raft) and Vogels’ eventually-consistent framing commit to wholeness as state-machine coherence at distributed-data-store scope, where the architectural goal is that the system appears as one rather than as a collection of collaborating components. CKS commits the same architectural commitment-shape at Self scope: the enterprise brain functions as unit composed of cells through aspects under unified human governance. The shared shape is unit-level architectural coherence; the differentiating scope is enterprise brain Self with substrate-shared topology and relational aspects.

Humans hold authority over the brain’s substrate content and over the orchestration rules that sustain the unit-level commitments; the brain functions as architectural unit under unified human governance, with the labor of operating the brain available to humans directly or to LLMs operating under human direction.

9.3 The topology consequence move

The topology consequence commitment names what follows from the substrate-shared aspect arrangement: cross-aspect coordination as first-class architectural capability rather than as inter-system message under integration glue. The term *topology* in CKS names the architectural-shape sense — Couchbase’s framing of “Topology: The Architecture of Distributed Systems” anchors the sense at first use. CKS does not import network-topology graph structures, microservices-deployment-topology shapes, or mathematical-topology continuous-deformation invariants; topology here is the architectural arrangement of substrates, aspects, and cells that governs how context, provenance, and governance flow across the enterprise brain.

The load-bearing differentiation is substrate-shared-topology versus integration-glued-topology at structural-pattern level. The integration-glue pattern Hohpe & Woolf’s *Enterprise Integration Patterns* canonized describes the structural pattern traditional enterprise architectures share: each major application speaks its own language, with integration hubs or message buses mediating via point-to-point mappings, Canonical Data Model patterns, and asynchronous messaging. The Canonical Data Model in this lineage is explicitly framed as integration pattern — a translation layer over local schemas — rather than as substrate commitment. The result is application-centric architecture with integration fabric, where each application retains its own local structures and governance, and the fabric focuses on data translation and routing rather than on global decision provenance or multi-aspect governance. Brittleness, data inconsistency, and exponential cross-system coordination cost are the named architectural consequences. The pattern is recognizable across the major enterprise software families at the structural-pattern level; the topology-consequence claim engages the pattern they share rather than each system individually.

CKS commits to substrate-shared topology in a lineage already recognized at adjacent scopes. Data warehousing (Inmon, Kimball) and the data lakehouse (Databricks, Apache Iceberg, Delta Lake) commit to substrate-sharing through one canonical analytical store; the post-microservices retrospective 2023–26 — including the Amazon Prime Video monolith-rebuild case crystallizing the substrate-sharing trade-off in popular discourse — explicitly positions

substrate-sharing as architectural axis with named trade-offs. The integration-glue characterization above engages traditional ERP/CRM-with-ESB enterprise architecture at the Hohpe & Woolf canonical-pattern level; modern enterprise architecture has substantially evolved toward substrate-sharing in specific dimensions, and the substrate-shared-topology commitment shares structural shape with several modern strands. Snowflake’s data-sharing primitive and Databricks Unity Catalog commit to substrate-sharing for analytical and governance content across operational boundaries; dbt with semantic-layer / metric-layer commits to substrate-sharing for business-logic definitions; event-streaming architectures (Kafka with Debezium and materialized views) commit to substrate-sharing for state-change events. These modern strands occupy specific substrate-sharing dimensions (analytical content; semantic definitions; event streams); CKS’s commitment is to substrate-shared topology covering decision provenance and multi-aspect governance under unified human governance — territory the modern strands do not jointly occupy at the structural-pattern level the wholeness and topology-consequence moves require. The integration-glue characterization holds at the structural-pattern level (the pattern traditional architectures share, and the pattern modern strands selectively move beyond in specific dimensions); CKS extends the substrate-shared-topology direction to the enterprise brain Self with relational aspects and multi-shaped governance per aspect-domain. The substrate-shared topology supports cross-aspect coordination as first-class architectural capability; humans hold authority at architectural level, and the labor of operating the architecture is available to humans directly or to LLMs operating under human direction.

9.4 Three supporting commitments

The three supporting commitments operate the architectural moves and are the primitives the wholeness commitment is testable against. Each is developed in turn with closest-neighbor differentiation and its own actor-neutrality phrasing; together they are what makes the wholeness move more than rhetoric and the topology-consequence move more than promise.

Bidirectional traceability and propagation. The first commitment is that traceability and propagation hold as first-class architectural property across all composition levels and aspect-domains — Self-level decisions trace to cell-level information across any aspect, and cell-level changes propagate to Self-level state coherently. The W3C PROV data model and successors (PROV-AGENT for AI agent activity) establish the architectural pattern of connected provenance graphs supporting both backward and forward queries by design; modern data-lineage platforms distinguish upstream lineage from downstream impact analysis as feature. The load-bearing differentiation: PROV is bidirectional in intent but actual completeness is instrumentation-dependent, and modern lineage tools are best-effort and partial. CKS commits substrate-shared topology making bidirectionality first-class testable invariant, in the disciplinary tradition where systems-engineering standards (automotive SPICE) treat bidirectional traceability matrices as obligation rather than nice-to-have. The architectural-pattern-level test for the bidirectionality invariant: any addressable substrate content reachable by forward query (cell—aspect—Self) is reachable by backward query (Self—aspect—cell) using the same substrate-retrieval primitives, with no separate instrumentation required for the reverse direction. The substrate-shared topology supports both query directions as architectural property; deployment-specific implementation of the retrieval primitives is engineering work the design pattern leaves to deployments. The substrate makes traceability first-class as architectural property; humans hold authority over what flows through the topology and over the bidirectionality commitment, and the labor of operating traceability is available to humans directly or to LLMs operating under human direction.

Hierarchical reorganization as architectural primitive. The second commitment is that aspects are organically grown through cell composition rather than predetermined as enterprise-software categories — and that reorganization restructures aspects as enterprise needs evolve, operating as substrate-edit primitive under the same machinery as Claim 3’s lifecycle primitives and Claim 4’s evolution mechanisms under Claim 5’s multi-shaped governance. *Organically grown* here names a structural-shape consequence of the relational role membership Claim 2 establishes carrying through to enterprise scale; it does not name autonomous-organizational-dynamics or collective-intelligence emergence, and the bound is placed deliberately so the architectural-shape sense holds throughout. The closest architectural-commitment-shaped neighbor is Team Topologies (Skelton & Pais), which treats team boundaries and interaction modes as first-class design variables and commits to Continuous Adaptation as principle. The differentiation is design-pattern-framework-with-vocabulary-requiring-change-management-labor versus architectural-primitive-with-explicit-machinery — the framework prescribes patterns and vocabulary while assuming change-management labor; CKS commits reorganization as substrate-edit operation directly enacted by the architecture

under governance. The architectural-pattern-level operational machinery: relational role membership (Claim 2) renders aspect composition as substrate-edit operation over which-cells-participate-in-which-aspects role bindings, rather than as code-change-and-redeploy. Reorganization restructures the role-binding substrate under the same authority architecture as DNA evolution; deployment-specific implementation of the role-binding substrate is engineering work the design pattern leaves to deployments. Reorganization happens under governance as substrate-edit operation; humans hold authority over reorganization decisions, and the labor of executing reorganization is available to humans directly or to LLMs operating under human direction.

Institutional knowledge as Self-level substrate content. The third commitment is that institutional knowledge — coherent organizational memory not reducible to per-cell or per-aspect content — lives as Self-level substrate content under multi-shaped governance. Enterprise knowledge graphs (Google KG, Bloomberg’s enterprise KG, Clinical KG) are the closest substrate-shaped neighbor; Talisman’s framing of knowledge graph as “a purposeful arrangement of components assembled in relationship to one another” anchors the architectural-arrangement sense. The differentiation lives on three axes the EKG cluster does not commit to simultaneously: scope of control extends from data-and-semantics to full decision-and-behavior provenance; governance shape extends from data-quality-and-access-control to multi-shaped per aspect-domain; runtime coupling extends from read-dominant to mandatory routing through substrate. The architectural-pattern-level operational machinery for mandatory routing: the orchestration substrate (Paper 1 §4) gates LLM access patterns to route through Self-level institutional-knowledge substrate for institutional queries, rather than permitting direct LLM-to-source access that bypasses institutional context. Mandatory routing is enforced as architectural property of the substrate-mediator commitment; deployment-specific implementation of routing-enforcement is engineering work the design pattern leaves to deployments. Knowledge-management canon (Nonaka; Davenport; Wenger) treats institutional knowledge as application content and social practice — KM systems as one class of apps alongside ERP and CRM. CKS commits institutional knowledge as the Self of the AI estate, with the same architectural shape — substrate plus aspect-domain composition plus governance — repeating across cell, aspect, Self, and organizational levels; Constitutional-AI architectures at corporate-deployment scope compose with the commitment as deployment-substrate-content under multi-shaped governance per aspect-domain. The brain holds knowledge as Self-level substrate under governance; humans hold authority over institutional-knowledge substrate content, and the labor of populating and updating is available to humans directly or to LLMs operating under human direction.

9.5 Reviewer-objection responses

Seven anticipated objections cluster into three response paragraphs through two paired-response treatments.

Discipline-level objections. Three related objections challenge the discipline of the claim. *Vision-mongering* attacks vision-claim writing as a genre. *Wholeness is just holism dressed up as architecture* imports philosophical holism as the wrong category for the wholeness move. *Topology consequence engages a straw-man* attacks the foil engagement as unfair characterization of named systems. The unifying response is that the discipline operates through specific architectural primitives, not through aspirational rhetoric: wholeness via three named primitives that the architecture either supports or does not, topology consequence at structural-pattern level with traditional architectures characterized at the integration-glued pattern they share rather than at any individual named system, and vision via the design-pattern-as-vision-pending-validation tradition where patterns begin as structured hypotheses requiring downstream validation and explicit pending-debate qualifiers. The wholeness move is engineering-technical wholeness with metaphysics bounded out, not philosophical assertion; the topology-consequence move engages the structural pattern traditional enterprise architectures share at the canonical-Hohpe-and-Woolf level, not specific named systems individually. The objections also surface at more specific levels of attack the genre-level response does not engage directly. The wholeness move’s specific challenge — *engineering-technical wholeness without operational specification is folk software-engineering wisdom* — is engaged at §9.4’s three supporting commitments, where each commitment is rendered with architectural-pattern-level operational machinery that supplies the test the wholeness invariant is checkable against. The topology-consequence move’s specific challenge — *the integration-glue characterization is the 2003 picture; modern enterprise architecture has substantially moved to substrate-sharing in specific dimensions* — is engaged at §9.3 through the modern-strands acknowledgment (Snowflake data-sharing, Databricks Unity Catalog, dbt semantic layer, event-streaming) sitting alongside the integration-glue characterization at structural-pattern level. The vision-claim’s specific challenge — *the comparative-cost rhetoric does load-bearing work without empirical cost comparison* — is addressed by the architectural-difference framing the vision-claim phrasing across

the manuscript carries, replacing comparative-cost rhetoric with non-comparative architectural-difference framing where integration-glued enterprise architectures take a different structural approach.

Differentiation-level objections. Three further objections — *bidirectional traceability is just provenance plus propagation*; *hierarchical reorganization is just organizational change management*; *institutional knowledge as Self-level substrate is just KM with extra vocabulary* — collapse each supporting commitment into an existing discipline. The response is per-commitment differentiation, drawing on the differentiation work §9.4 develops. Provenance and propagation exist as bidirectional capabilities in the surveyed literature; CKS commits them as testable invariants of substrate topology rather than as best-effort coverage by layered tools. Organizational change management treats reorganization as managed transformation project with humans driving and systems as levers; CKS commits reorganization as architectural primitive operating under the same governance machinery as other lifecycle primitives. KM treats institutional knowledge as application content and social practice; CKS commits it as Self-level substrate content under multi-shaped governance with architectural shape repeating across cell, aspect, Self, and organizational levels.

The empirical-grounding objection. The seventh objection — *the thought experiment is hand-waving without empirical grounding* — engages the §10 forward reference. Claim 6 is deliberately offered as architectural vision rendered through a thought experiment, not as empirically established law. This is standard for architecture and design-pattern work: patterns begin as structured hypotheses grounded in observed practice and theory, validated or revised through empirical studies and pilots. The CKS topology and Self-level substrate specify what would have to be true in the architecture for properties like cross-aspect coordination and institutional wholeness to compose into a coherent design pattern at enterprise scope; real-world enterprise deployment is the next step in testing those assumptions, and §10's healthcare thought experiment renders the architecture concretely as the design-pattern-pending-validation discipline requires.

9.6 Paper 1 inheritance, Claims 1–5 inheritance, and closing

Claim 6 inherits Paper 1's six architectural commitments at enterprise scope. Paper 1 §6's linear-cost composition is the load-bearing inheritance for the wholeness commitment's tractability — composition stays linear in cost as the unit propagates from cells through aspects and Selves to the enterprise brain. Paper 1 §3.3's authority-vs-labor distinction holds throughout: governance is an authority architecture, not a review workflow. Paper 1 §4's substrate/LLM division at the governance boundary, §5's two-level conflict-handling, and §6.3's authoring-vs-curation labor distinction extend through Claim 6 to enterprise scope unchanged. Inheritance from Paper 2's prior claims is direct as well. Claim 1's instinct/reasoning separation extends to enterprise brain scope, where the brain's instinct layer evolves following exterior LLM upgrade cycles and the reasoning layer evolves through substrate content the enterprise governs. Claim 2's three-level structure with relational role membership, the DNA/action layer distinction within every cell, and expression as governed selection extends through. Claim 3's lifecycle primitives, Claim 4's three evolution mechanisms in productive tension, and Claim 5's multi-shaped governance extend through every level of the brain.

The architectural commitments compose into a coherent design pattern carrying the vision claim. The wholeness and topology-consequence moves, supported by bidirectional traceability, hierarchical reorganization as primitive, and institutional knowledge as Self-level substrate content, render the enterprise brain Self as architecturally coherent design pattern extending Paper 1's cell concept. As §10's thought experiment will render concretely in healthcare domain continuing Paper 1's POC lineage, empirical validation through real-world enterprise deployment is downstream work; humans hold authority over what the architecture commits to.

10. The Enterprise Brain: A Thought Experiment in Healthcare

10.1 Setting and the cell-aspect-Self structure rendered

This section renders the enterprise brain Self in a hypothetical healthcare organization, continuing Paper 1's lineage where the proof-of-concept was a cell deployed in a regulated industry “with healthcare as one natural anchor” (Paper 1 §11.1). The organization is hypothetical and is not specified by size or organizational shape; what the rendering

specifies is the architectural shape produced when Paper 2's commitments compose at enterprise scope into one brain. The rendering is offered as architectural vision; empirical validation in real-world enterprise deployment is downstream work.

The cell layer is populated by many cells covering distinct work the organization performs. Paper 1's regulatory-reporting POC cell is one of them, continuing the POC lineage at enterprise brain scope; alongside it sit cells for patient intake, clinical-decision-support consultation, care coordination, quality-measure tracking, vendor management, capital planning, internal audit findings tracking, and regulatory-change impact analysis. Each cell carries its own DNA and action layers, runs under appropriate governance with humans holding authority over its substrate content and orchestration rules, and has its own lifecycle. The cell architecture Paper 1 establishes replicates across the population, with each cell doing the work cells do at the layer where the work lives.

The aspect layer is populated by kinds of work the organization performs. Four illustrative aspects render the architectural pattern: clinical-care-delivery work, regulatory-and-quality-reporting work, financial-and-operations work, and institutional-learning work. These are illustrative rather than fixed taxonomy — aspects form through cell composition based on how work flows in the organization, not through predetermined enterprise-software categories. A real organization would have aspects shaped by its own work patterns rather than by these four; what repeats across all aspects is the architectural pattern (substrate plus DNA/action layer plus governance plus lifecycle), regardless of which specific work content fills the cells composing into the aspect.

Cells participate in multiple aspects simultaneously through Claim 2's relational role membership. A clinical-data cell carrying a patient cohort's longitudinal information participates in the clinical-care-delivery aspect when clinicians draw on it for treatment decisions, in the regulatory-and-quality-reporting aspect when it informs quality-measure submissions, in the financial-and-operations aspect when it informs reimbursement workflows and resource planning, and in the institutional-learning aspect when it contributes to the organization's accumulating clinical knowledge. The same cell, four role-instances, governance properties propagating per arrangement rather than per cell. This is the substrate-shared topology in operation at structural-pattern level.

The Self layer is the enterprise brain itself: all cells, all aspects, all institutional knowledge composed into one Self under unified human governance. The Self holds institutional knowledge as Self-level substrate content — clinical philosophy, operational priorities, strategic positioning, risk appetite, accumulated decision history — and has its own lifecycle: instantiated when the organization commits to the architecture, evolving through governed mechanisms, reorganizing as needs shift. Authority over the Self's substrate content, the orchestration rules under which cells and aspects operate, and the governance shapes per aspect-domain rests with the organization's humans — clinicians, administrators, compliance officers, finance leaders, regulators with appropriate domain authority.

10.2 Architectural commitments threaded through the rendering

With the cell-aspect-Self structure in view, the six claims thread through it as architectural commitments operating in concert; the threading proceeds in numerical order, each claim rendering the same enterprise brain from a different architectural angle.

The brain's instinct layer (LLM) and reasoning layer (CKS substrate) compose into one brain under unified human governance — Claim 1's instinct/reasoning separation at enterprise scope. The elasticity property is visible in operation: as exterior LLMs sharpen at clinical text understanding, what previously required substrate reasoning collapses to instinct; when LLM upgrades introduce noise, uncertainty, or new governance requirements, additional reasoning substrates are added under governance. The instinct/reasoning boundary is itself governed substrate content, and humans hold authority over where instinct is trusted across the brain — with high-stakes clinical decisions architecturally pinned to the reasoning layer regardless of how capable instinct becomes.

The three-level structure and relational role membership the rendering above made tangible are Claim 2's structural commitment; the DNA/action distinction within every cell and expression as governed selection complete it. The DNA layer carries stabilized orchestration content — typed substrates, harness logic, conflict-handling rules, lifecycle policies — while the action layer carries recorded execution. Expression determines which DNA-layer substrates activate per cell goal: a quality-measure-tracking cell may carry the full Self's regulatory DNA with selective expression for its current measure set, or a partial slice if storage and cognitive load matter more — a per-deployment design choice under orchestration substrate humans govern.

Lifecycle primitives operate at every level — Claim 3’s commitment threading through cells, aspects, and the Self uniformly. Birth: a new clinical-decision-support cell comes into existence under governed authorization when the organization takes on a new clinical-pathway commitment, and a new aspect can be born when the organization’s work pattern shifts. Mating: two regulatory-reporting cells combine when reporting regimes consolidate, with three pattern variants (union, selective merge, lineage-preserved union) configuring the combination, and aspects can split when work bifurcates. Death has two distinct types: a vendor-management cell is retired and archived through lineage supersession when the organization divests, with archived substrate remaining addressable for audit; a cell whose function becomes obsolete is deleted through functional obsolescence.

Three evolution mechanisms operate concurrently in productive tension across all levels — Claim 4 in the brain. Instinct evolution sharpens the operational capability the brain supports as architectural unit as LLM upgrades flow in from upstream, integrated through verification substrates that determine which prior reasoning can be retired, retained as fallback, or retained as active verification on safety-critical paths. DNA evolution proceeds under governance-defined goals as clinicians, administrators, compliance officers, and finance leaders refine substrate content in their domains — orchestration rules improve, schemas tighten, policy substrate sharpens. Action-feedback evolution closes the loop from accumulated execution evidence back into governed DNA refinement under human-mediated proposal-and-acceptance machinery. Cells, aspects, and the Self evolve concurrently under unified governance, with the three mechanisms holding in productive tension at every level of the brain.

Multi-shaped governance operates across the brain with distinct governance shapes per aspect-domain — Claim 5 in operation. The clinical-care-delivery aspect runs under clinical-safety governance, where errors carry patient-safety consequences and verification regimes match. The regulatory-and-quality-reporting aspect runs under audit-trail governance, where every claim must trace to source. The financial-and-operations aspect runs under financial-reporting governance. The institutional-learning aspect runs under organizational-knowledge governance. The shapes are described by architectural function rather than by specific named regulatory frameworks; what differs across aspect-domains is the governance machinery the authority-vs-labor commitment instantiates. HIPAA propagates across aspect-domains as substrate-level cross-aspect content: anywhere patient data surfaces — in the clinical aspect when clinicians use it, in the regulatory aspect when it informs reporting, in the financial aspect when it informs reimbursement, in the learning aspect when it accumulates as institutional knowledge — the HIPAA commitment applies, governed differently in each aspect-domain per the multi-shaped commitment. The instinct/reasoning boundary is itself substrate content tunable per aspect-domain: clinical-care-delivery may pin high-stakes diagnostic decisions to the reasoning layer regardless of LLM capability, while less consequential routine work may delegate more to instinct.

What the foregoing renders concretely is Claim 6’s vision claim. *Wholeness*: the enterprise brain functions as unit serving the organization with its own lifecycle, evolution, and governance — not as aggregation of cells with cell-level commitments only. *Topology consequence*: where integration-glued enterprise architectures handle a strategic decision affecting clinical care, regulatory reporting, financial operations, and institutional learning as four operations with integration glue between them, the substrate-shared topology supports cross-aspect coordination as first-class architectural capability. Three supporting commitments operate visibly through the rendering. *Bidirectional traceability*: a clinical decision traces to cell-level information across any aspect, and substrate-level changes radiate to relevant cells and aspects under the same substrate-shared topology. *Hierarchical reorganization*: aspect composition shifts through governed substrate-edit operations under the same machinery as Claim 3’s lifecycle primitives — reorganization as architectural primitive rather than as managed transformation project running alongside the architecture. *Institutional knowledge*: the organization’s accumulated knowledge lives as Self-level substrate content under multi-shaped governance, with the architectural shape (substrate plus relational aspects plus governance plus lifecycle) repeating from cell through aspect through Self.

10.3 Honest limits and positioning

The rendering offered in this section is architectural vision rendered concretely; it is not empirical demonstration. The healthcare thought experiment specifies what an enterprise brain Self could look like when Paper 2’s commitments compose at enterprise scope, showing the architectural shape Claims 1 through 6 produce when threaded through a hypothetical organization. Empirical validation in real-world enterprise deployment is downstream work — the architecture’s commitments are testable against the primitives developed across the six claims, and whether

those primitives produce the cross-aspect coordination, institutional wholeness, and governed evolution properties the rendering describes when an actual organization deploys at scale is a question the rendering does not answer. The thought experiment is rendering at the architectural-pattern level: it specifies what the cell-aspect-Self structure looks like when the six claims thread through it, and what the supporting commitments produce in operation. It does not specify how cells are composed into aspects under deployment-scale operational machinery, how reorganization happens administratively at organizational pace, how authority architectures are implemented when thousands of cells operate concurrently, or how the multi-shaped governance is administered across an enterprise's operational footprint. These operational-specification questions are deployment-architecture work the design pattern leaves to deployments rather than thought-experiment work the rendering would over-specify by pursuing.

Several scope limits warrant explicit naming. The specific cells, aspects, and governance shapes any real organization would deploy depend on its work flows and governance commitments rather than on the four illustrative aspects rendered here; the architectural pattern repeats across organizations, but the specific shape of any organization's enterprise brain would be determined by its own work patterns, regulatory environment, and operational priorities. The rendering is not a deployment specification, and §10 does not prescribe a target architecture for any particular organization. Within Paper 2's scope, the rendering is also bounded against inter-Self dynamics — how this brain would interact with other organizations' brains across organizational boundaries is not in scope.

The vision is offered as architectural commitment pending debate and observation as systems like this get built. Where the rendering succeeds is in making concrete what the cell concept Paper 1 establishes looks like when composed through aspects into a Self at enterprise scope; where it remains pending is real-world enterprise deployment, where the architecture's commitments are tested against the operational, governance, and organizational realities of actual healthcare organizations operating at scale and over time.

11. Discussion and Implications

11.1 What the composed architecture commits to

The six claims compose into one architectural commitment, and the composition is what Paper 2's contribution amounts to as a whole. The instinct/reasoning separation Claim 1 establishes is the foundation: fast-pattern instinct (the LLM, functioning as a System-1 analogue) and deliberate reasoning (the CKS substrate, functioning as a System-2 analogue) live as two independently-evolving layers under unified human governance. Claim 2's structural commitment then makes the separation architecturally available at every composition level, with cells composing through aspects into Selves under relational role membership, distinguishable DNA and action layers within every cell, and expression as governed selection over which substrate content activates per goal. Onto that structural foundation Claim 3 places lifecycle primitives — birth, mating, death — applying uniformly at every level. Claim 4's three evolution mechanisms — instinct evolution, DNA evolution, action-feedback evolution — operate in productive tension across multiple levels and on horizontal and vertical axes. Claim 5's multi-shaped governance specifies distinct governance shapes co-determined with distinct mechanisms, the instinct/reasoning boundary as governed substrate content, and distinct death-type governance processes. And Claim 6 offers the vision claim that the enterprise extension §9 establishes is how the per-Self commitments compose into the enterprise brain Self under wholeness and topology-consequence commitments — the brain functioning as one unit composed of cells through aspects under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability.

What follows from this composition is a set of properties the architecture commits to specifically. The architecture commits to substrate handling coordination, governance, and conflict-handling outside the LLM regardless of LLM capability, for the same reasons Paper 1 established at cell scope and now extended to Self and enterprise brain scope. It commits to modularity at every level as architectural requirement at instantiation rather than as emergent property the architecture has to wait for. It commits to lifecycle primitives operating uniformly at every level under governance, with mating producing variation through three pattern variants and death producing two distinct architectural results — deletion in functional obsolescence; archival with continued substrate-addressability in lineage supersession. It commits to evolution operating through three mechanisms in productive tension under unified human

governance, with directedness alongside undirectedness and structural reorganization on operational timescales as architectural property. It commits to governance as multi-shaped, with shape co-determined with mechanism rather than overlaid as mechanism-agnostic policy, and with the boundary between instinct and reasoning itself substrate content humans tune as instinct sharpens. And it commits to the enterprise brain Self as architecturally coherent design pattern extending Paper 1's cell concept, with substrate-shared topology, hierarchical reorganization as architectural primitive, institutional knowledge as Self-level substrate content under multi-shaped governance per aspect-domain, and bidirectional traceability across all composition levels and aspect-domains. What distinguishes these architectural-consequence commitments from emergent properties beyond the architecture is testable at the architectural-pattern level: each commitment traces to specific primitives the design pattern names, and the operational test is whether the property holds when the named primitives compose under the architecture's commitments. Cross-aspect coordination as first-class capability follows from substrate-shared topology composing with relational role membership and multi-shaped governance; institutional wholeness follows from the same primitives composing across the cell-aspect-Self structure; bidirectional traceability follows from the substrate-shared topology supporting forward and backward queries with the same retrieval primitives. A property is architectural-consequence if it holds when the named primitives compose under the architecture's commitments; a property is emergent and outside what the architecture commits to if it requires conditions beyond what the primitives compose to support. The architectural-consequence commitments above are what the primitives compose to support; the disclaims that follow protect against scope creep at the boundary this test draws.

The line worth drawing here separates architectural-consequence properties from emergent-properties-beyond-architecture. The architecture commits to the topology-consequence property §9 establishes — substrate-shared topology supports cross-aspect coordination as first-class architectural capability under unified human governance. It does not commit to emergent collective intelligence in the philosophically-loaded sense: no consciousness, no autonomous goal-formation beyond what humans encode in substrate content, no integrated agency beyond architectural property. It does not commit to autonomous organizational dynamics within or beyond the architecture's commitments: hierarchical reorganization is architectural primitive operating under governance machinery, not the brain reorganizing itself, with humans holding authority over reorganization decisions and the labor of executing them available to humans directly or to LLMs operating under human direction. The architecture does commit to capabilities that follow as architectural consequence from the LLM instinct layer composed with the substrate's governed content under the wholeness and topology-consequence commitments — including the cross-aspect coordination capability §9 establishes. It does not commit to capabilities beyond what the architecture's commitments compose to support: no capability emergence outside the architecture's primitives, no behavior surfacing from the composition that humans cannot trace to substrate content under their authority. It does not commit to the LLM becoming dispensable as instinct sharpens — Claim 1's separation is governance-motivated rather than capability-motivated, which means the reasoning layer's function persists regardless of how capable instinct becomes. The disclaim is the bound this paper's vision-claim framing requires: the vision is offered as architectural commitment pending debate and observation in real-world enterprise deployment, not as prediction of properties beyond what the architecture's commitments support.

11.2 Scope and limits

The scope of Paper 2's claims warrants explicit naming. Paper 2's claims at cell scope inherit from Paper 1's proof-of-concept at cell scope: the substrate/cell two-level structure, the six architectural commitments, the outside-the-model spine, the authority-vs-labor distinction, conflict preservation, linear-cost composition, and tool-agnostic instantiation are all defended in Paper 1 with the regulated-industry POC carrying the load. The architectural commitments at aspect, Self, and enterprise brain scope are defended architecturally — through closest-neighbor differentiation, cross-claim composition, and per-claim primitive specification — and rendered concretely through §10's thought experiment in healthcare, continuing Paper 1's POC lineage at enterprise brain scope. The thought experiment is architectural vision rendered concretely rather than empirical demonstration; what it produces is a concrete picture of what an enterprise brain Self looks like when Paper 2's commitments compose at enterprise scope, not a demonstration that any particular enterprise deployment will produce that picture. Whether the architecture's commitments produce the cross-aspect coordination, institutional wholeness, and governed evolution properties §10 describes when an actual organization deploys at scale is downstream work that empirical validation in real-world enterprise deployment is positioned to address.

A second-order scope limit follows from the genre commitment. Paper 2 does not produce empirical evaluation of the per-Self commitments, the per-mechanism governance shapes, the per-pattern mating governance, the death-type decision processes, or the cross-aspect coordination property at scale. The contribution is design-pattern writing in Paper 1's genre, with the design-pattern's structured hypothesis available to be adopted, varied, composed with adjacent patterns, or argued against. Empirical work — case studies, comparative deployment studies, controlled experiments at sub-Self scope, longitudinal observation of Selves under continued evolution — would close the evaluation gap the design-pattern genre by construction does not. The scope gap between Paper 1's empirical grounding and Paper 2's architectural defense warrants explicit naming. The design-pattern-pending-validation tradition Paper 2 invokes is grounded historically in extensive observation across many deployments — Alexander's pattern language across centuries of building practice; Hohpe and Woolf's enterprise integration patterns across many enterprise deployments; the broader software-architecture pattern lineage across decades of practice. CKS's empirical grounding is one POC at cell scope plus this paper's thought experiment at enterprise scope; the wider pattern-tradition's grounding-across-many-deployments standard is not what Paper 2 carries at enterprise scope. What Paper 2 carries instead is architectural defense at the design-pattern level — closest-neighbor differentiation, cross-claim composition, per-claim primitive specification, and concrete rendering through §10's thought experiment — with empirical validation in real-world enterprise deployment positioned as the work that closes the gap the design-pattern tradition's typical grounding standard names.

11.3 Open questions for future work

Four open questions stand out as natural next steps for research extending what Paper 2 establishes. The questions are research extensions rather than foundational gaps in the architectural commitments — Paper 2 specifies the architectural primitives at the level the design pattern requires, and the open questions name research territory the primitives open rather than territory the primitives leave incomplete. First, per-mating-pattern governance shapes: Claim 3's three mating pattern variants — union, selective merge, lineage-preserved union — carry distinct governance implications Claim 5 gestures at without developing per-pattern. The union pattern's commitment to merge-time conflicts as persistent substrate state, the selective-merge pattern's commitment to curation rules as substrate content, and the lineage-preserved-union pattern's commitment to lineage pointers as substrate all suggest distinct authority architectures, distinct verification regimes, and distinct reversion paths that per-pattern governance shapes would specify. Second, per-mechanism interaction at points where mechanisms compose: Claim 4's productive tension framing commits to compositional integration rather than adversarial equilibrium, but the dynamics of how instinct evolution, DNA evolution, and action-feedback evolution interact when all three operate concurrently on the same cell or aspect — and how multi-shaped governance per Claim 5 holds across the interactions — are research territory the productive-tension framing opens rather than closes. Third, multi-Self interaction patterns: Paper 2's scope ends at the enterprise brain Self under unified human governance, and how multiple enterprise brain Selves interact across organizational boundaries — coordinate or compete, evolve under different governance regimes, share substrate content under cross-organizational authority — is a class of extension theories that would each stand on its own claim set rather than on arguments Paper 2 develops. Fourth, threshold questions for substrate tuning at scale, in two parallel forms. The boundary-tuning question: Claim 5 commits the instinct/reasoning boundary to be governed substrate content humans tune, but the criteria humans should use to tune it as instinct capability sharpens — when a previously-reasoned decision should architecturally collapse to instinct, when it should be retained as fallback, when it should be retained as active verification on safety-critical paths — is research territory the substrate-content commitment opens rather than answers. The reorganization-at-scale question: Claim 6's hierarchical reorganization commitment treats reorganization as substrate-edit operation under governance, but the criteria humans should use to authorize reorganization decisions when the enterprise brain operates with thousands of cells across dozens of aspects — when reorganization decisions can be batched at policy level under governance, when per-decision authorization is required, when authorization can be delegated to authority-architecture roles operating at sub-Self scope — is parallel research territory. Both threshold questions concern substrate-tuning criteria humans should develop as the architecture operates at organizational scale; the architecture's commitments are stated, and what is open is the operational discipline humans bring to tuning the substrate the architecture commits to.

12. Conclusion

This paper has extended the Coordination Knowledge Substrate pattern Paper 1 (Li, 2026) defends at cell scope to AI Self scope and toward the enterprise brain Self as architecturally coherent design pattern extending Paper 1's cell concept, defending six architectural commitments in turn. The first separates fast-pattern instinct from deliberate reasoning into two independently-evolving layers under unified human governance, extending Paper 1's substrate/LLM division at the governance boundary from cell scope to Self scope. The second commits structural modularity at every level — multi-level composition with relational roles, the DNA/action layer distinction within every cell, expression as governed selection — as architectural requirement rather than evolved property. The third commits lifecycle as governed primitives — birth, mating with three pattern variants, death with two distinct architectural results — operating uniformly at every level. The fourth commits evolution as three mechanisms in productive tension under unified human governance: instinct evolution, DNA evolution, and action-feedback evolution, operating across multiple levels and on horizontal and vertical axes simultaneously. The fifth commits human governance as multi-shaped, with shape co-determined with mechanism, the instinct/reasoning boundary as governed substrate content, and distinct death-type governance processes. The sixth offers the vision claim that the enterprise brain Self is architecturally coherent as a design pattern extending Paper 1's cell concept — a functioning unit composed of cells through aspects under unified human governance, with substrate-shared topology supporting cross-aspect coordination as first-class architectural capability where integration-glued enterprise architectures take a different structural approach.

The unifying architectural move across all six claims is the separation of fast-pattern instinct and deliberate reasoning into independently-evolving layers under unified human governance — Paper 1's *outside-the-model vs. inside-the-model* spine extended through Self lifecycle, evolution, and into the enterprise brain Self. Each claim either establishes structures the separation requires, develops dynamics the separation enables, or specifies governance properties that keep the separation safe. The resulting architecture answers the field-level architectural question §1.1 names: an LLM-era AI system that is governed across its lifecycle and evolution looks like a CKS-governed AI Self in which coordination, governance, conflict-handling, and the reasoning humans should govern remain architecturally separable from the model layer for the same reasons Paper 1 identified at cell scope, now extended to a complete AI Self with a lifecycle, evolutionary dynamics, and a vision of how per-Self commitments compose into the enterprise brain Self.

The commitments are calibrated explicitly. The architecture commits to what the closest-neighbor literature does not jointly occupy and to what Paper 1's proof-of-concept supports through inheritance at cell scope. It does not commit to emergent collective intelligence, autonomous organizational dynamics, capabilities beyond what the architecture's commitments compose to support, or the LLM becoming dispensable as instinct sharpens. The vision claim Claim 6 offers is calibrated to the architectural-commitment sense — the architecture's commitments compose into a coherent design pattern at enterprise scope, with the rendering in §10's healthcare thought experiment offering an architectural rendering pending the validation real-world enterprise deployment provides.

The contribution is design-pattern writing in Paper 1's genre. Paper 2 articulates a design pattern available to be adopted, varied, composed with adjacent patterns, or argued against, sitting at the intersection of dual-process AI, structural-modularity-as-architectural-commitment, biology-of-evolvability, governance-as-architecture, and enterprise AI deployment lineages. The design pattern's empirical validation through real-world enterprise deployment is positioned as downstream work; what Paper 2 establishes is the architectural commitments and their composition into one CKS-governed AI Self extending toward the enterprise brain Self as architecturally coherent design pattern extending Paper 1's cell concept.

References

This references list is partial. It includes the academic and peer-reviewed sources most directly cited across Paper 2's six claims, the contemporary industry voice sources most directly engaged in closest-neighbor differentiation, and the primary architectural lineage anchors. Additional entries — including vendor documentation, standards body publications, practitioner essays, conference workshop papers, and preprints requiring further verification — will appear in a revised version. Paper 1 (Li, 2026) is cited throughout Paper 2 as the substrate-cell foundation Paper 2

extends; the citation scheme uses “Paper 1 \$X” for section references and full citation only on first introduction.

Stage 9 reconciliation expanded the lineage-anchor coverage to include the cognitive-architectures (SOAR, ACT-R, CLARION, ICARUS) and organizational-design (Galbraith, Mintzberg, Argyris & Schön, Senge) traditions Paper 2’s body engages, the holonic manufacturing reference architecture, and Alexander’s *The Nature of Order* as the canonical anchor for §9.2’s wholeness move. Stage 9 also reconciled the cross-paper-stub entries Paper 2 v2.1 carried as “see Paper 1 references” pointers to full canonical entries inherited from Paper 1 v1.1, and resolved verification on five preprint entries whose v2.1 metadata required correction (Bai, Engin, Mohsin, Quirke, Ziabari).

Alexander, C. (1979). *The Timeless Way of Building*. New York: Oxford University Press. (Architectural-pattern lineage anchor for §9.2’s wholeness move; metaphysics of “wholeness as property of nature” bounded out per this paper’s engineering-technical commitment.)

Alexander, C. (2002). *The Nature of Order: An Essay on the Art of Building and the Nature of the Universe, Book 1 — The Phenomenon of Life*. Berkeley, CA: Center for Environmental Structure. (Canonical reference for §9.2’s wholeness concept; Volume 1 of the four-volume series 2002–2005.)

Anderson, J. R. (1993). *Rules of the Mind*. Hillsdale, NJ: Lawrence Erlbaum Associates. (Cognitive-architectures lineage anchor for §3.1 and §4.3; ACT-R as named architecture with declarative/procedural distinction.)

Argyris, C., & Schön, D. A. (1978). *Organizational Learning: A Theory of Action Perspective*. Reading, MA: Addison-Wesley. (Organizational-design lineage anchor.)

Bai, Y., Kadavath, S., Kundu, S., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL]. <https://arxiv.org/abs/2212.08073> (51 authors; first three named, “et al.” abbreviation per citation discipline for very-large-author preprints.)

Booch, G., Fabiano, F., Horesh, L., Kate, K., Lenchner, J., Linck, N., Loreggia, A., Murugesan, K., Mattei, N., Rossi, F., & Srivastava, B. (2021). Thinking Fast and Slow in AI. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(17), 15042–15046. <https://doi.org/10.1609/aaai.v35i17.17765> (Canonical AAAI 2021 paper from the SOFAI program at IBM Research; subsequent CACM 2025 and CP 2024 follow-ups extend the program.)

Bottino, F., Ferrero, C., Dosio, N., & Beneventano, P. (2026). Retrieval is not enough: Why organizational AI needs epistemic infrastructure. arXiv:2604.11759 [cs.AI]. <https://arxiv.org/abs/2604.11759> (OIDA framework; surfaced as adjacent substrate-based architecture in §3 Related Work and §4 Claim 1 inheritance; canonical entry inherited from Paper 1 v1.1.)

Chang, E. Y. (2025). MACI: Multi-agent collaborative intelligence for adaptive reasoning and temporal planning. arXiv:2501.16689 [cs.AI]. <https://arxiv.org/abs/2501.16689> (Canonical entry inherited from Paper 1 v1.1.)

Coyne, J. A., & Orr, H. A. (2004). *Speciation*. Sunderland, MA: Sinauer Associates.

Davenport, T. H., & Prusak, L. (1998). *Working Knowledge: How Organizations Manage What They Know*. Boston, MA: Harvard Business School Press. (KM canon anchor for §9.4’s institutional-knowledge commitment differentiation.)

Davidson, E. H., & Erwin, D. H. (2006). Gene regulatory networks and the evolution of animal body plans. *Science*, 311(5762), 796–800. <https://doi.org/10.1126/science.1113832>

Davidson, E. H., & Levine, M. S. (2008). Properties of developmental gene regulatory networks. *Proceedings of the National Academy of Sciences*, 105(51), 20063–20066. <https://doi.org/10.1073/pnas.0806007105>

Engin, Z. (2025). Human-AI Governance (HAIG): A Trust-Utility Approach. arXiv:2505.01651 [cs.AI]. <https://arxiv.org/abs/2505.01651> (Solo-authored; cited at §8.5 for the Decision Authority / Process Autonomy / Accountability Configuration three-dimension framing.)

Fabiano, F., Pallagani, V., Bergamaschi Ganapini, M., Horesh, L., Loreggia, A., Murugesan, K., Rossi, F., & Srivastava, B. (2024). Plan-SOFAI: A Neuro-Symbolic Planning Architecture. *NuCLear Workshop at AAAI 2024*.

Galbraith, J. R. (1973). *Designing Complex Organizations*. Reading, MA: Addison-Wesley. (Organizational-design lineage anchor.)

- Hartwell, L. H., Hopfield, J. J., Leibler, S., & Murray, A. W. (1999). From molecular to modular cell biology. *Nature*, 402(Suppl), C47–C52. <https://doi.org/10.1038/35011540>
- Hodgson, P. (2017). *Feature Toggles (aka Feature Flags)*. martinowler.com. <https://martinowler.com/articles/feature-toggles.html> (retrieved April 2026).
- Hohpe, G., & Woolf, B. (2003). *Enterprise Integration Patterns: Designing, Building, and Deploying Messaging Solutions*. Boston, MA: Addison-Wesley. (Canonical-source anchor for \$9.3's topology-consequence move differentiation against the integration-glued pattern traditional enterprise architectures share.)
- Holland, J. H. (1975). *Adaptation in Natural and Artificial Systems*. Ann Arbor, MI: University of Michigan Press.
- Inmon, W. H. (2005). *Building the Data Warehouse* (4th ed.). Indianapolis, IN: Wiley. (Substrate-sharing lineage anchor for \$3.4 and \$9.3.)
- Jablonka, E., & Lamb, M. J. (2014). *Evolution in Four Dimensions: Genetic, Epigenetic, Behavioral, and Symbolic Variation in the History of Life* (revised ed.). Cambridge, MA: MIT Press.
- Kahneman, D. (2011). *Thinking, Fast and Slow*. New York: Farrar, Straus and Giroux.
- Kerr, J. F. R., Wyllie, A. H., & Currie, A. R. (1972). Apoptosis: a basic biological phenomenon with wide-ranging implications in tissue kinetics. *British Journal of Cancer*, 26, 239–257. <https://doi.org/10.1038/bjc.1972.33>
- Kimball, R., & Ross, M. (2013). *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling* (3rd ed.). Hoboken, NJ: Wiley. (Substrate-sharing lineage anchor for \$3.4 and \$9.3.)
- Kirschner, M. W., & Gerhart, J. C. (2005). *The Plausibility of Life: Resolving Darwin's Dilemma*. New Haven, CT: Yale University Press.
- Laird, J. E., Newell, A., & Rosenbloom, P. S. (1987). SOAR: An architecture for general intelligence. *Artificial Intelligence*, 33(1), 1–64. [https://doi.org/10.1016/0004-3702\(87\)90050-6](https://doi.org/10.1016/0004-3702(87)90050-6) (Cognitive-architectures lineage anchor for \$3.1 and \$4.3; SOAR with working-memory architecture.)
- Laland, K. N., Uller, T., Feldman, M. W., Sterelny, K., Müller, G. B., Moczek, A., Jablonka, E., & Odling-Smee, J. (2015). The extended evolutionary synthesis: its structure, assumptions and predictions. *Proceedings of the Royal Society B*, 282(1813), 20151019. <https://doi.org/10.1098/rspb.2015.1019>
- Langley, P., & Choi, D. (2006). A unified cognitive architecture for physical agents. In *Proceedings of the Twenty-First National Conference on Artificial Intelligence* (pp. 1469–1474). Boston, MA: AAAI Press. (Cognitive-architectures lineage anchor for \$3.1 and \$4.3; ICARUS with knowledge-and-skill memory split.)
- Lehman, J., & Stanley, K. O. (2011). Abandoning objectives: Evolution through the search for novelty alone. *Evolutionary Computation*, 19(2), 189–223. https://doi.org/10.1162/EVCO_a_00025
- Li, W. (2026). *Coordination Outside the Model: A Human-Governed Substrate Pattern for AI Systems* [Paper 1 of the CKS theory series]. (Cited as “Paper 1” throughout Paper 2.)
- Maynard Smith, J., & Szathmáry, E. (1995). *The Major Transitions in Evolution*. Oxford: Oxford University Press.
- Miller, M. S. (2006). *Robust Composition: Towards a Unified Approach to Access Control and Concurrency Control*. PhD dissertation, Johns Hopkins University, Baltimore, MD, May 2006.
- Mintzberg, H. (1979). *The Structuring of Organizations: A Synthesis of the Research*. Englewood Cliffs, NJ: Prentice-Hall. (Organizational-design lineage anchor.)
- Mohsin, M. A., Umer, M., Bilal, A., Memon, Z., Qadir, M. I., Bhattacharya, S., Rizwan, H., Gorle, A. R., Kazmi, M. Z., Mohsin, A., Rafique, M. U., He, Z., Mehta, P., Jamshed, M. A., & Cioffi, J. M. (2025). On the Fundamental Limits of LLMs at Scale. arXiv:2511.12869 [cs.CL]. <https://arxiv.org/abs/2511.12869> (15 authors per v2 (Jan 2026) canonical author list.)
- Mosqueira-Rey, E., Hernández-Pereira, E., Alonso-Ríos, D., Bobes-Bascarán, J., & Fernández-Leal, Á. (2023). Human-in-the-loop machine learning: a state of the art. *Artificial Intelligence Review*, 56(4), 3005–3054. <https://doi.org/10.1007/s10462-022-10246-w>

- Nonaka, I., & Takeuchi, H. (1995). *The Knowledge-Creating Company: How Japanese Companies Create the Dynamics of Innovation*. New York: Oxford University Press. (KM canon anchor for §9.4's institutional-knowledge commitment differentiation.)
- Orlikowski, W. J. (1992). The duality of technology: Rethinking the concept of technology in organizations. *Organization Science*, 3(3), 398–427. <https://doi.org/10.1287/orsc.3.3.398>
- Potter, M. A., & De Jong, K. A. (2000). Cooperative coevolution: An architecture for evolving coadapted subcomponents. *Evolutionary Computation*, 8(1), 1–29. <https://doi.org/10.1162/106365600568086>
- Quirke, P., Oozeer, N., Bandi, C., Abdullah, A., Hoelscher-Obermaier, J., Phillips, J. M., Greaves, J., Neo, C., Lan, M., Barez, F., & Upadhyay, S. (2025). Beyond Monoliths: Expert Orchestration for More Capable, Democratic, and Safe Language Models. arXiv:2506.00051 [cs.CL]. <https://arxiv.org/abs/2506.00051> (Position paper; venue characterization downgraded to arXiv preprint pending NeurIPS 2025 Position Paper Track acceptance verification.)
- Schmidhuber, J. (2009). Ultimate cognition à la Gödel. *Cognitive Computation*, 1(2), 177–193. <https://doi.org/10.1007/s12559-009-9014-y> (Gödel Machine.)
- Seck, M. (2026). Cognitive Core: Architecture for demand-driven dual-process delegation. Personal communication and ongoing engagement via LinkedIn long-form content (2025–2026). (Adjacent dual-process configuration architecture engaged at architecture-time configuration; substantive engagement with the three-layer model — workflow YAML, domain YAML, case JSON — and the typed-epistemic-primitives decomposition draws from these communications. Cited at §3.2, §4.5, and §8.3.)
- Senge, P. M. (1990). *The Fifth Discipline: The Art and Practice of the Learning Organization*. New York: Doubleday/Currency. (Organizational-design lineage anchor.)
- Skelton, M., & Pais, M. (2019). *Team Topologies: Organizing Business and Technology Teams for Fast Flow*. Portland, OR: IT Revolution Press.
- Steimann, F. (2000). On the representation of roles in object-oriented and conceptual modelling. *Data & Knowledge Engineering*, 35(1), 83–106. [https://doi.org/10.1016/S0169-023X\(00\)00023-9](https://doi.org/10.1016/S0169-023X(00)00023-9)
- Stevens, A. (2025). The missing primitives for trustworthy AI agents (17-part series). *Sakura Sky Blog*, September 8 – December 1, 2025. Series index: <https://www.sakurasky.com/blog/missing-primitives-for-trustworthy-ai/> (retrieved April 2026). (Canonical entry inherited from Paper 1 v1.1; cited as practitioner-source control-plane neighbor.)
- Sun, R., Merrill, E., & Peterson, T. (2001). From implicit skills to explicit knowledge: A bottom-up model of skill learning. *Cognitive Science*, 25(2), 203–244. https://doi.org/10.1207/s15516709cog2502_2 (Cognitive-architectures lineage anchor for §3.1 and §4.3; CLARION with explicit/implicit knowledge separation.)
- Van Brussel, H., Wyls, J., Valckenaers, P., Bongaerts, L., & Peeters, P. (1998). Reference architecture for holonic manufacturing systems: PROSA. *Computers in Industry*, 37(3), 255–274. [https://doi.org/10.1016/S0166-3615\(98\)00102-X](https://doi.org/10.1016/S0166-3615(98)00102-X) (Holonic manufacturing reference architecture.)
- Vanchurin, V., Wolf, Y. I., Katsnelson, M. I., & Koonin, E. V. (2022). Toward a theory of evolution as multilevel learning. *Proceedings of the National Academy of Sciences*, 119(6), e2120037119. <https://doi.org/10.1073/pnas.2120037119>
- Vincenti, W. G. (1990). *What Engineers Know and How They Know It: Analytical Studies from Aeronautical History*. Baltimore, MD: Johns Hopkins University Press.
- Wenger, E. (1998). *Communities of Practice: Learning, Meaning, and Identity*. Cambridge: Cambridge University Press. (KM canon anchor for §9.4's institutional-knowledge commitment differentiation.)
- West-Eberhard, M. J. (2005). Developmental plasticity and the origin of species differences. *Proceedings of the National Academy of Sciences*, 102(suppl 1), 6543–6549. <https://doi.org/10.1073/pnas.0501844102>
- Wu, X., Xiao, L., Sun, Y., Zhang, J., Ma, T., & He, L. (2022). A survey of human-in-the-loop for machine learning. *Future Generation Computer Systems*, 135, 364–381. <https://doi.org/10.1016/j.future.2022.05.014>

Zahn, O., & Chana, S. (2026). Facts as first class objects: Knowledge objects for persistent LLM memory. arXiv:2603.17781 [cs.AI]. <https://arxiv.org/abs/2603.17781> (Canonical entry inherited from Paper 1 v1.1; demonstrates database-like cost model of hash-addressed persistent memory.)

Ziabari, A. S., Ghazizadeh, N., Sourati, Z., Karimi-Malekabadi, F., Piray, P., & Dehghani, M. (2025). Reasoning on a Spectrum: Aligning LLMs to System 1 and System 2 Thinking. arXiv:2502.12470 [cs.CL]. <https://arxiv.org/abs/2502.12470>

Industry voice and standards-body sources cited across the claims include: Anthropic Responsible Scaling Policy; Apache Software Foundation Project Management Committee governance documentation; Atlan, Epsilla AgentStudio, Graph Digital, and Unstructured.io agent-memory and metadata-control-plane writings; AutoGen, LangGraph, Microsoft Agent Framework, Claude Agent SDK, OpenAI Agent Builder, StackAI, CrewAI AMP, A2A “agent cards” agent-framework documentation; automotive SPICE bidirectional traceability standard; Bloomberg Enterprise Knowledge Graph, Clinical Knowledge Graph, and Google Knowledge Graph as EKG cluster references engaged in §9.4’s institutional-knowledge differentiation; Couchbase “Topology: The Architecture of Distributed Systems” framing; Databricks, Apache Iceberg, and Delta Lake lakehouse documentation as substrate-sharing-lineage anchors; Davidson and Levine on gene regulatory networks (additional citation context); EU AI Act; ISO/IEC 42001; LiteLLM, OpenRouter, Portkey, Agentgateway gateway documentation; MIT Sloan Intelligent Choice Architectures and meta-decision-rights writings; NIST AI Risk Management Framework; Open Policy Agent / Conftest / Checkov policy-as-code documentation; PEP 387 and GitHub repository archival policy documentation; Puppet, Ansible, Chef, Kubernetes configuration management documentation; Talisman knowledge-graph writings; W3C PROV provenance specification and PROV-AGENT extension; the post-microservices retrospective with the Amazon Prime Video monolith-rebuild case as its popular crystallization; the system-design coherence framing on components fitting together. Sectoral and regulatory sources, vendor enterprise AI deployment platform documentation (Azure AI Foundry, Vertex AI, AWS Bedrock with SageMaker, Salesforce Agentforce, Anthropic Enterprise, OpenAI Enterprise), and practitioner essays referenced in closest-neighbor differentiation are not enumerated here; they will appear with full citations in a revised version.